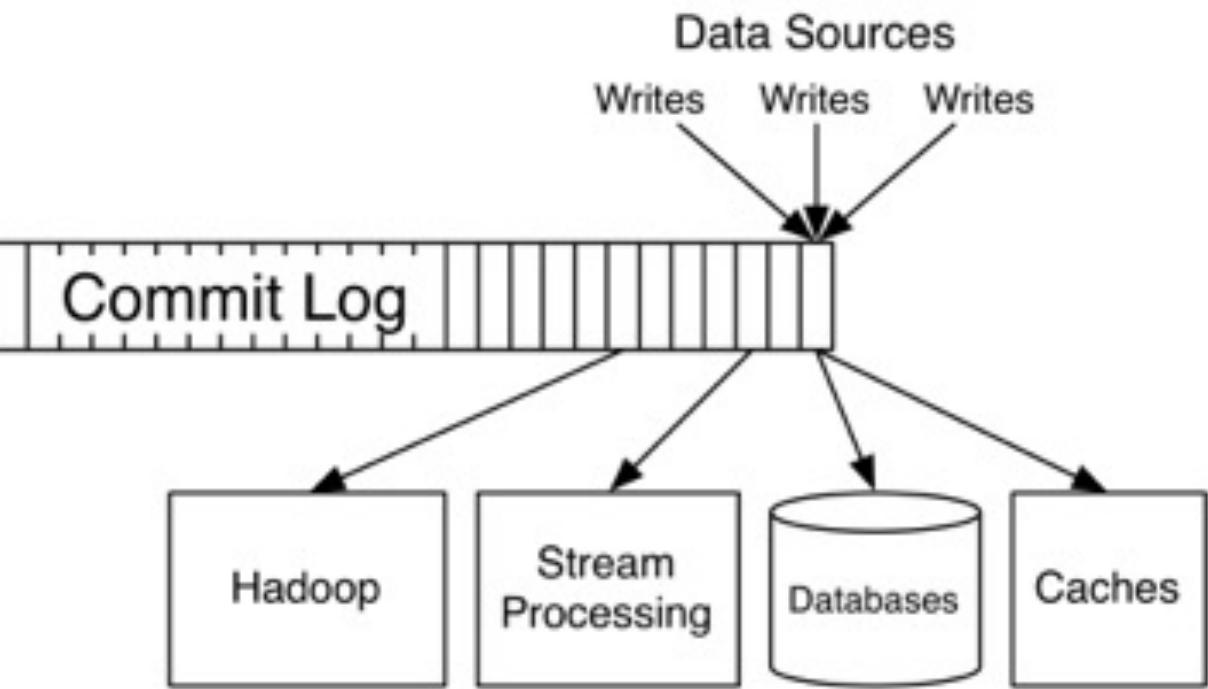




Why is My Stream Processing Job Slow?

Xavier Léauté, Software Engineer

Gwen Shapira, Principal Data Architect



Distributed
Scalable
Fault-Tolerant

Partitioned + Replicated Log
Ordering guarantees

Consumers advance independently

Exactly-once delivery
Transactional commits



What people think of Stream Monitoring

A wide-angle photograph capturing a massive colony of King penguins in their natural, cold environment. The foreground is filled with penguins standing in deep white snow, their dark bodies contrasting sharply with the bright, overexposed background. The scene is shrouded in a thick, white mist or falling snow, which obscures the sky and creates a sense of depth and density. The penguins are packed closely together, filling the frame from left to right.

What our typical experience is

Real Customer Experiences

Client Side

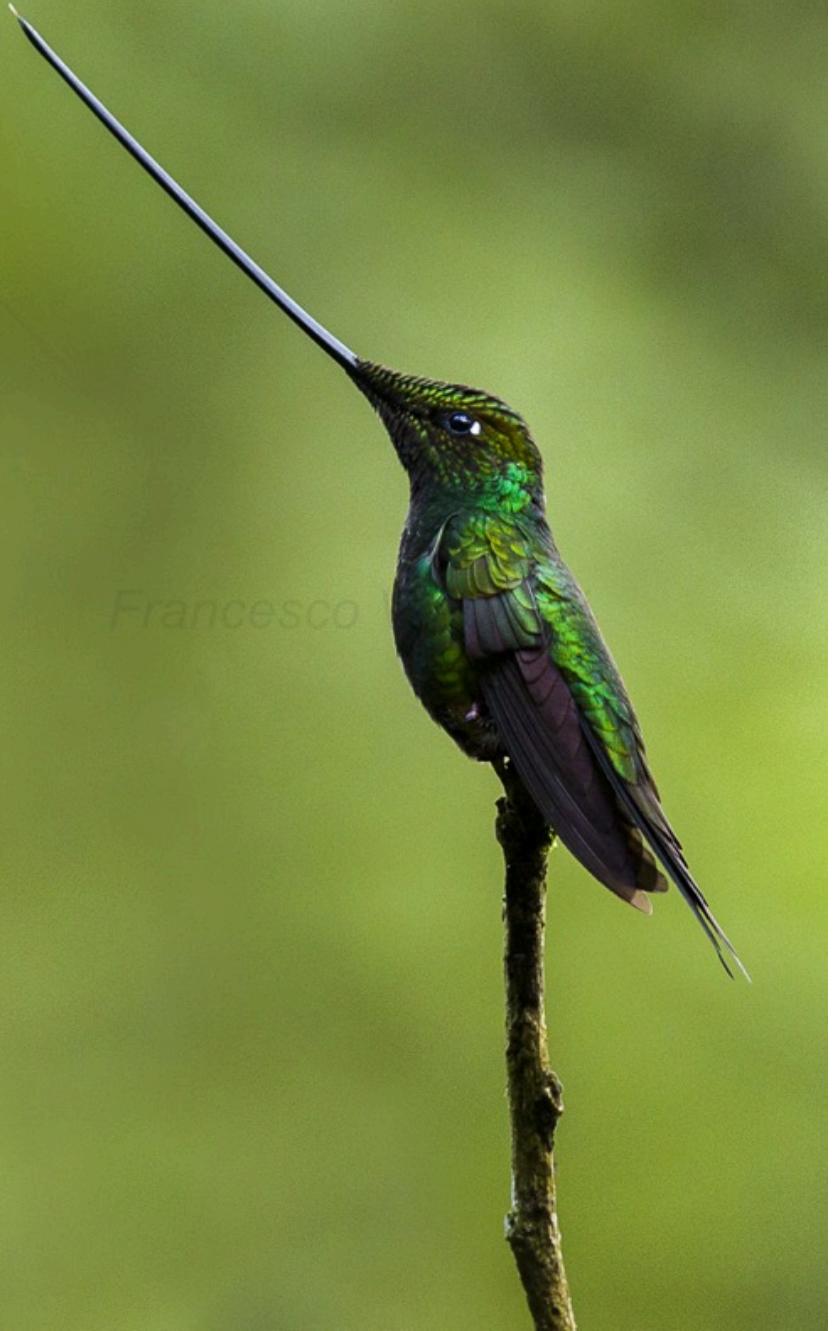
Broken Streaming Job / App

Client Side

Broken Streaming Job / App

End-to-End

Slow Replication



**Your Kafka stream job stopped
humming... now what?**



What we check

- Consumer Lag
- Partition Assignment
- Partition Skew
- Client Logs
- GC Log
- Metrics
 - Request Latencies
 - Commit Rates
 - Group Rebalancing
- Basic Tuning
 - Batch Sizes
 - Commit Rate
- Application Profiling

The Newbie - During an incident...

GC Logs? Metrics?
How do I get those?

I'll just change some configs
and reboot everything.

Consumer Lag



Bad Capacity Allocation

```
kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group fast-data-reader
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
fast-data	1	8661694	8703404	41710	myapp-1
fast-data	3	8577975	8616490	38515	myapp-2
fast-data	0	4902354	8741872	3839518	myapp-3
fast-data	2	4922614	8621757	3699143	myapp-3

Bad Capacity Allocation

```
kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group fast-data-reader
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
fast-data	1	8661694	8703404	41710	myapp-1
fast-data	3	8577975	8616490	38515	myapp-2
fast-data	0	4902354	8741872	3839518	myapp-3
fast-data	2	4922614	8621757	3699143	myapp-3

Bad Capacity Allocation

```
kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group fast-data-reader
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
fast-data	1	8661694	8703404	41710	myapp-1
fast-data	3	8577975	8616490	38515	mvapp-2
fast-data	0	4902354	8741872	3839518	myapp-3
fast-data	2	4922614	8621757	3699143	myapp-3

Watch for Partition Skew

```
kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group fast-data-reader
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
fast-data	1	8661694	8703404	41710	myapp-1
fast-data	3	8577975	8616490	38515	myapp-2
fast-data	0	4902354	8741872	3839518	myapp-3
fast-data	2	4922614	8621757	3699143	myapp-3

Watch for Partition Skew

```
kafka-consumer-groups --bootstrap-server localhost:9092 --describe --group fast-data-reader
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
fast-data	1	8661694	8703404	41710	myapp-1
fast-data	3	8577975	8616490	38515	myapp-2
fast-data	0	4902354	8741872	3839518	myapp-3
fast-data	2	4922614	8621757	3699143	myapp-3

Not all partitions are created equal

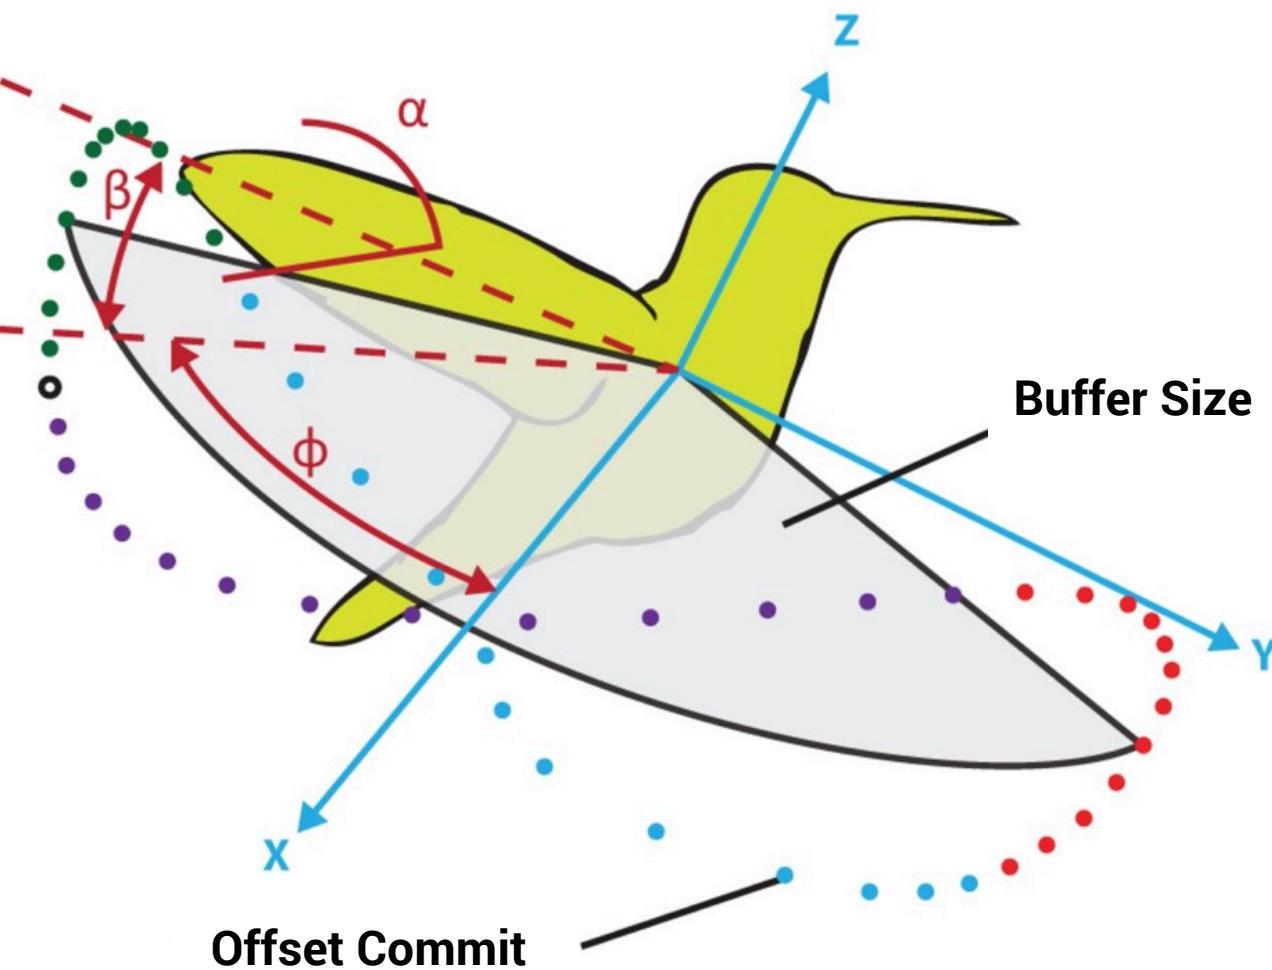


Important for
Keyed topics
Custom partitioned topics

Early warning signs
some partitions lagging
uneven CPU / Network usage

Typical cause
skewed key distribution in your data
bad joins (null keys)
imbalance across brokers

Clients have metrics too!



Start with the basics **GC / CPU / Network**

General Slowness

Consumer or Producer Side?

Global Request Latencies

Some partitions still lagging

Per Broker metrics (bad node / network)

Per Topic metrics (data / tuning)

Turn up the log level

The logs took too much space, so we deleted them.



Time for some profiling

<https://github.com/jvm-profiling-tools/async-profiler>

<https://github.com/brendangregg/FlameGraph>

```
./profiler.sh -d 30 -f flamegraph.svg <pid>
```

To impress your coworkers

<https://github.com/Netflix/flamescope>

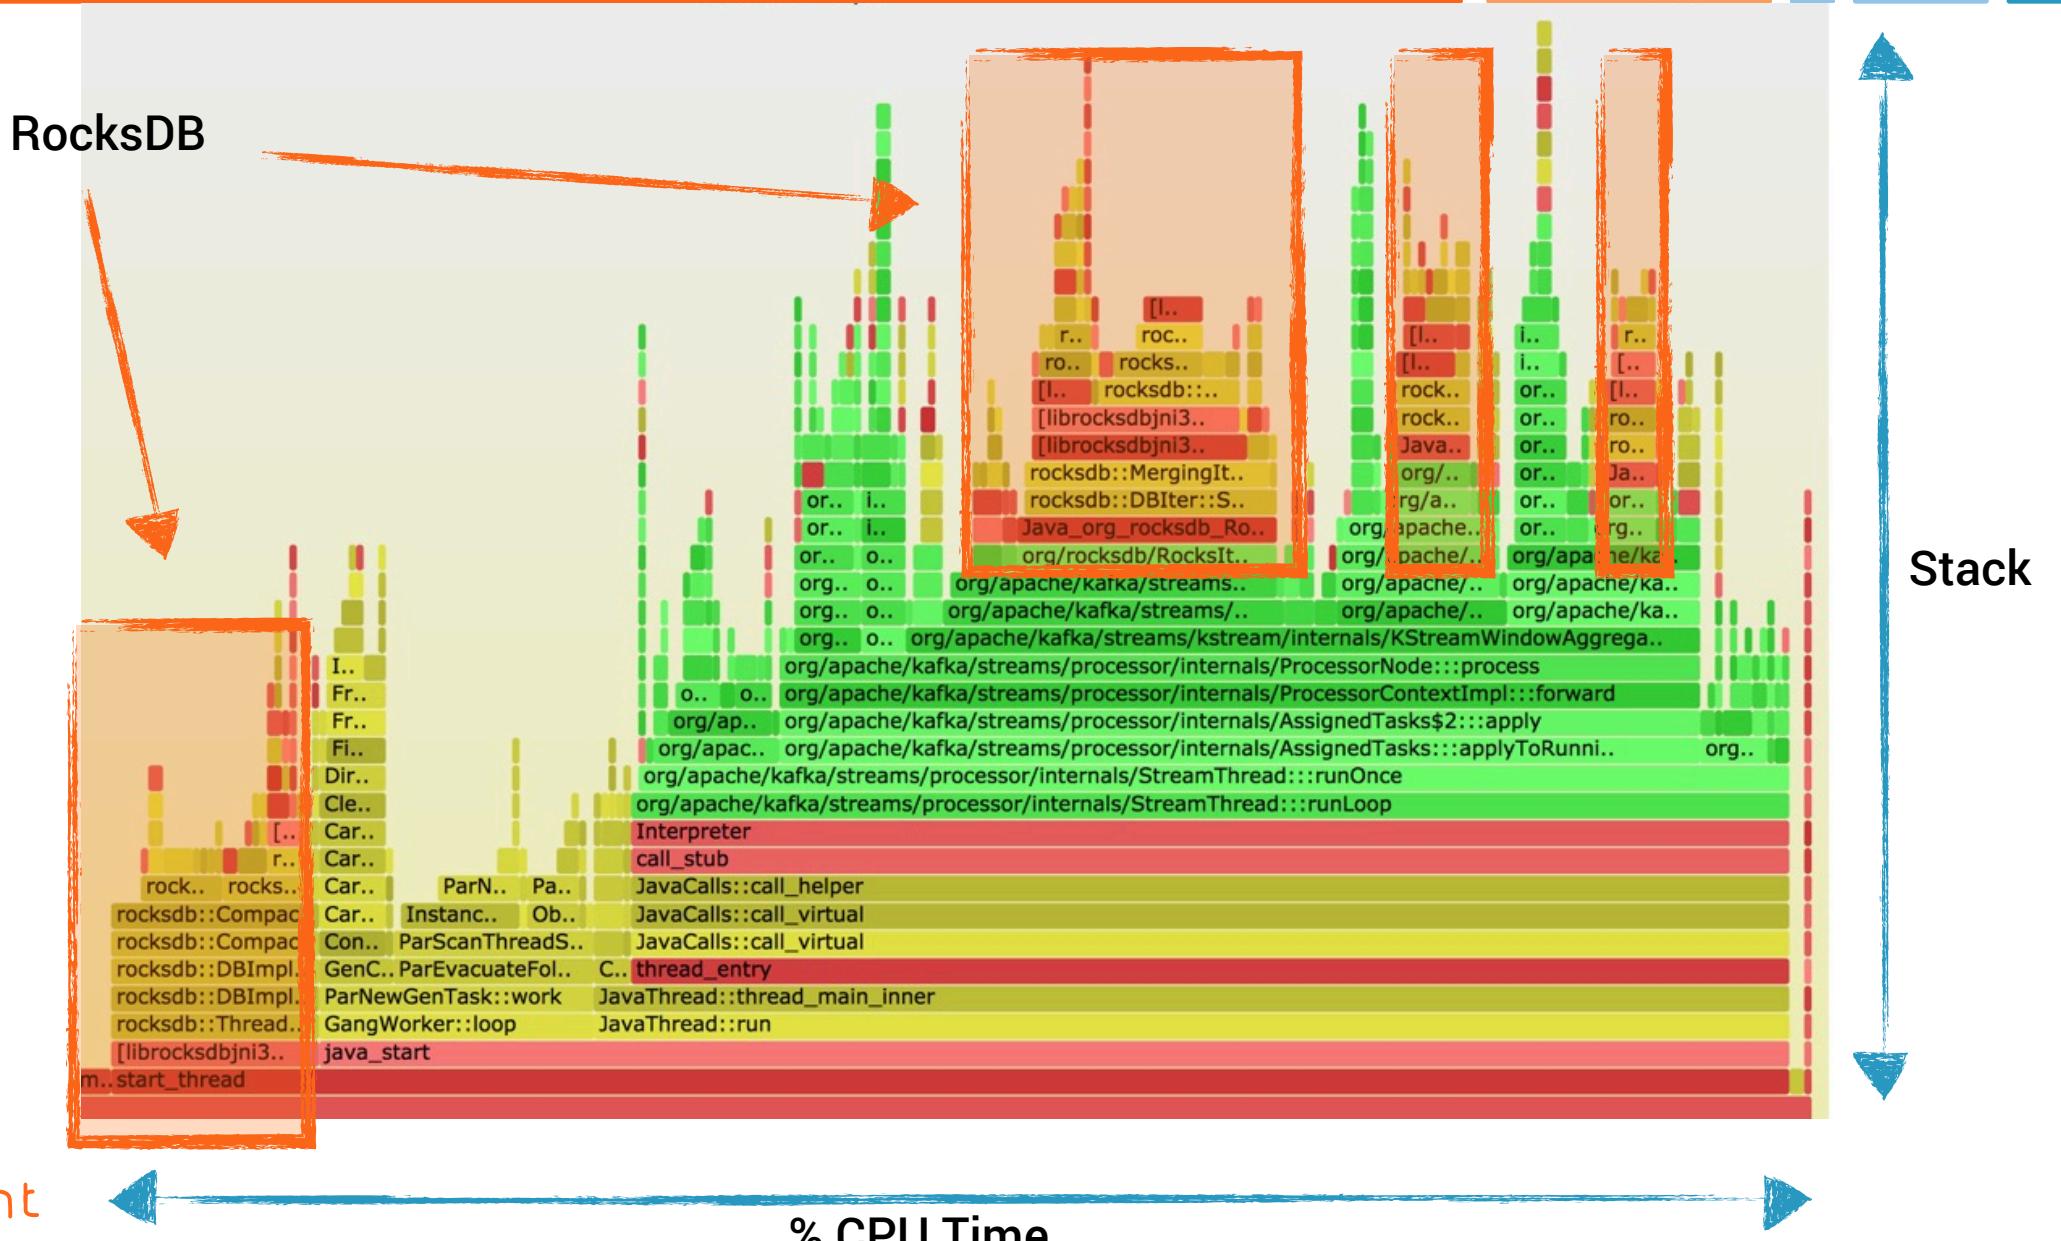
Here's where your CPU cycles went



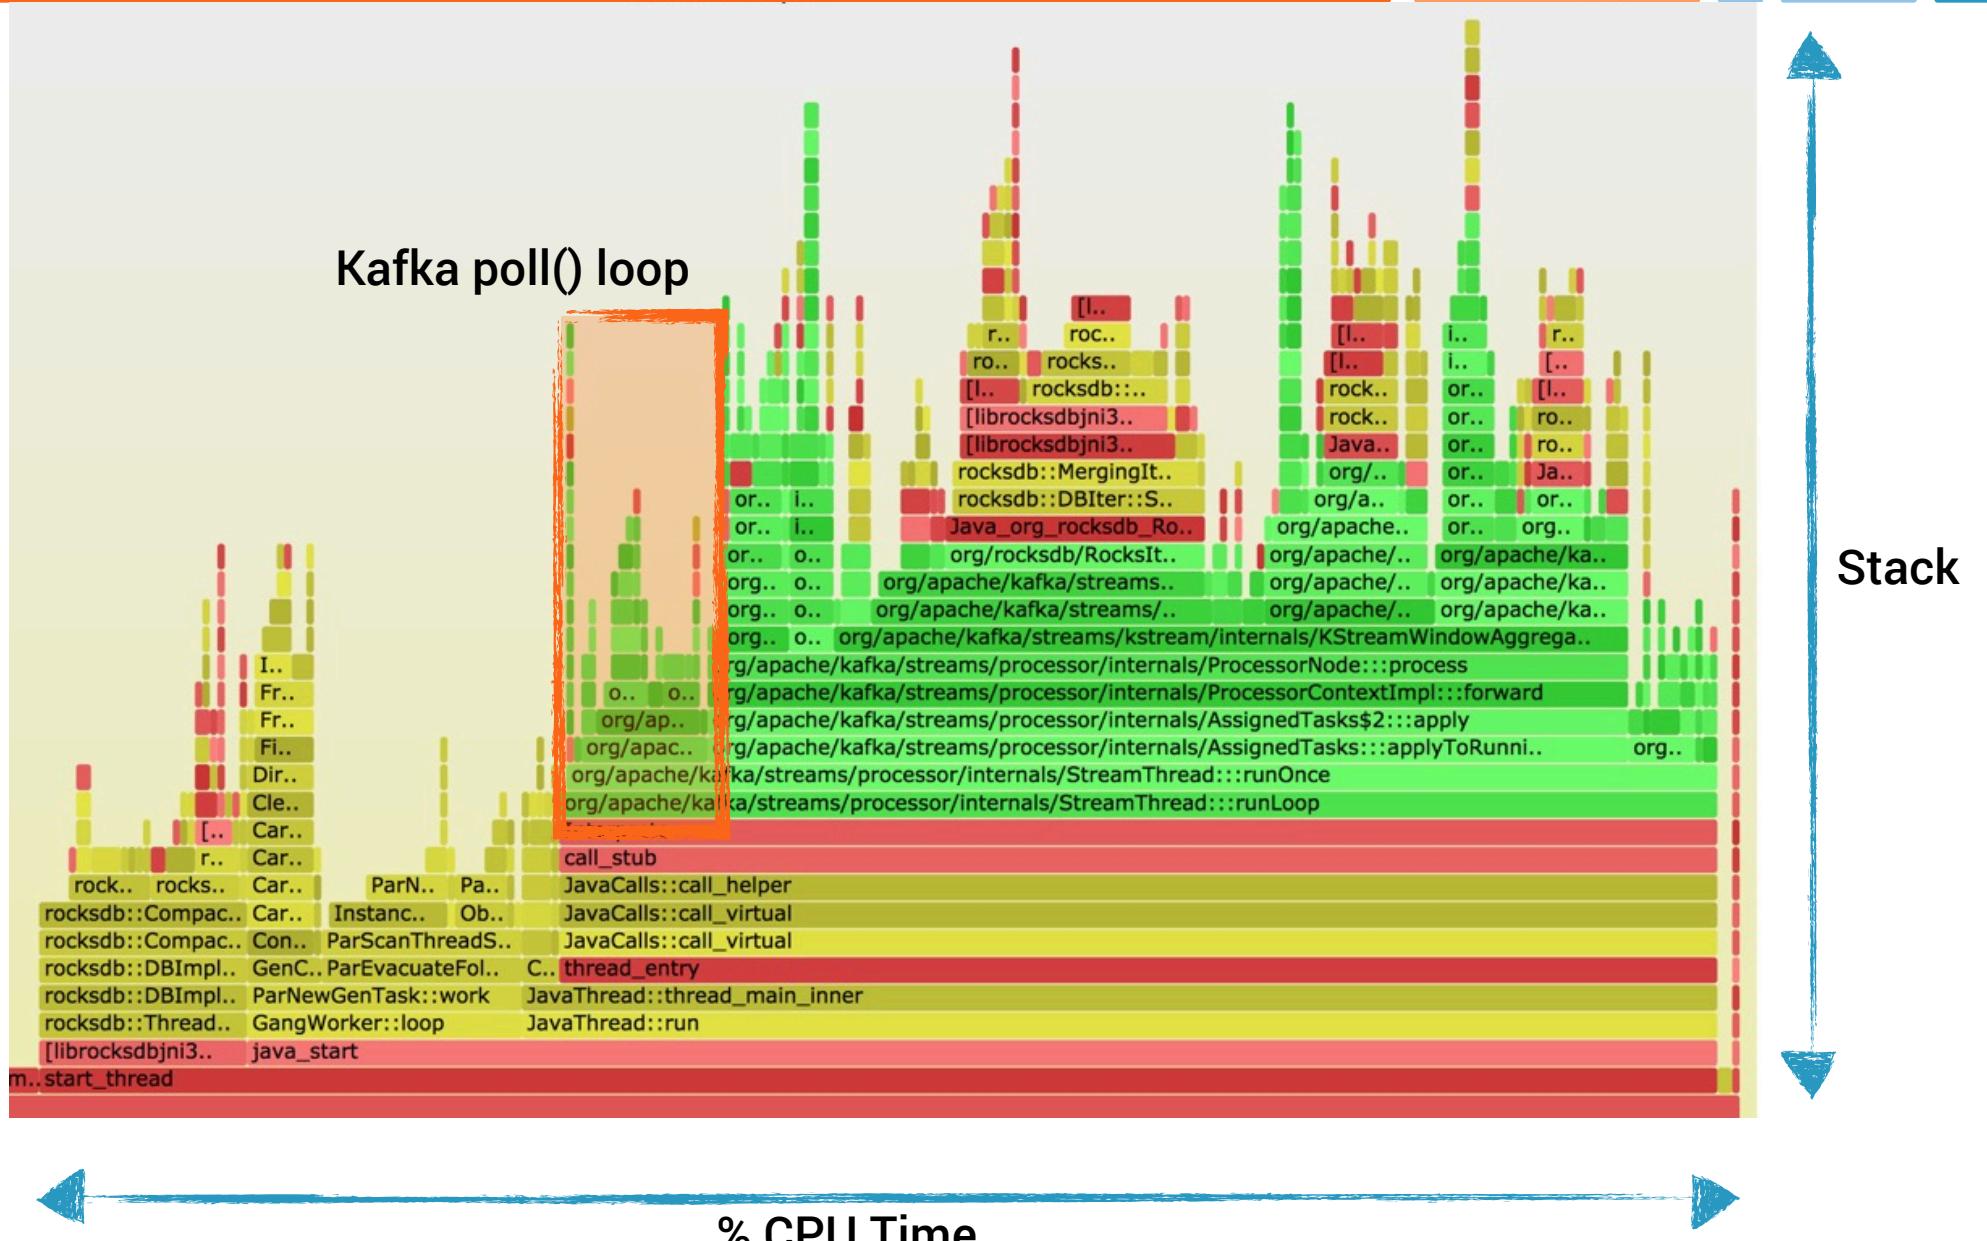
Here's where your CPU cycles went



Here's where your CPU cycles went



Here's where your CPU cycles went



Here's where your CPU cycles went

Actual Processing Time

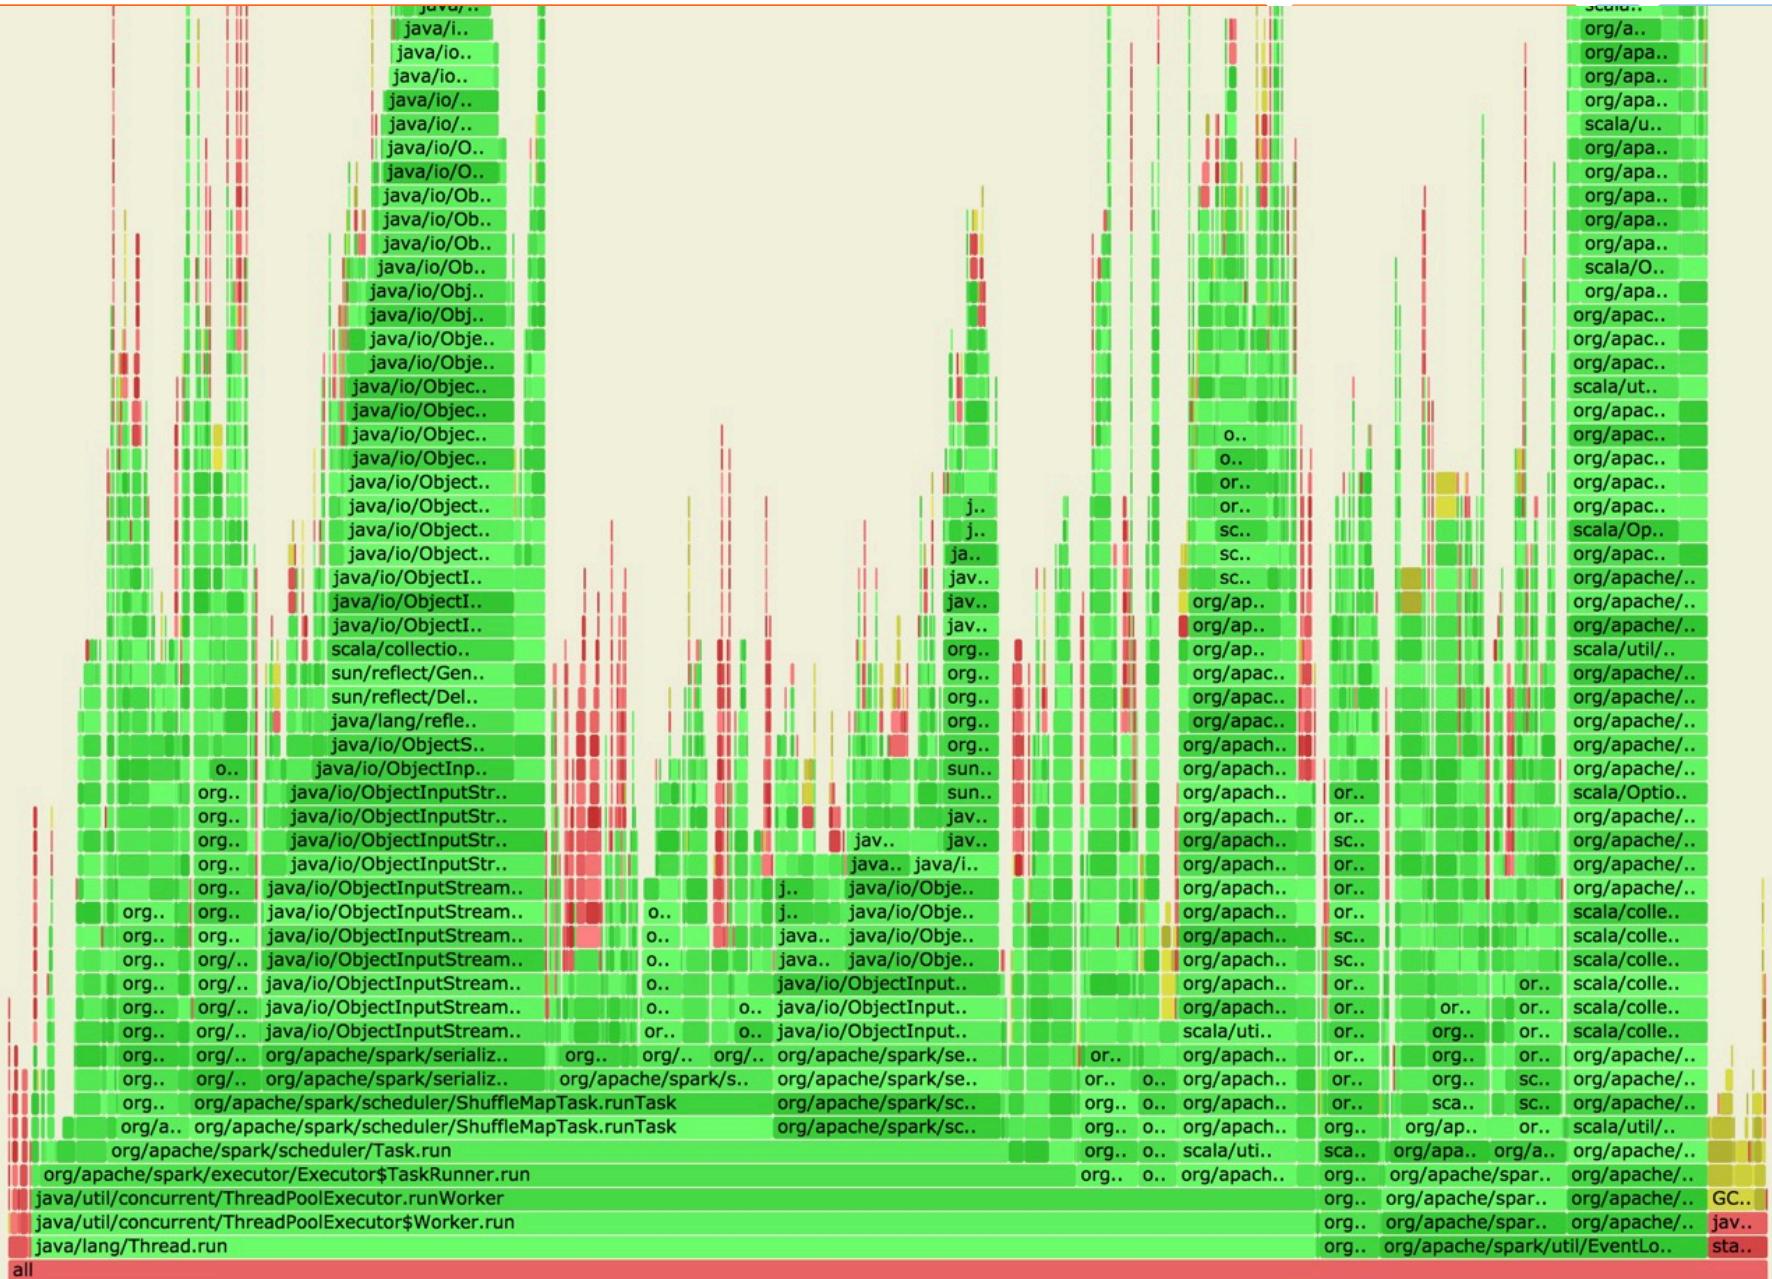
Stack



Spark Streaming Clickstream Example (using Kafka)

```
val pageViews = KafkaUtils.createDirectStream[String, String](  
    streamingContext, LocationStrategies.PreferConsistent,  
    ConsumerStrategies.Subscribe[String, String](topicsSet, kafkaParams))  
    .map { record => PageView.fromString(record.value) }  
// Return the rate of error pages (a non 200 status) in each zip code over the last 30 seconds  
val statusesPerZipCode = pageViews.window(Seconds(30), Seconds(2))  
    .map(view => (view.zipCode, view.status))  
    .groupByKey()  
val errorRatePerZipCode = statusesPerZipCode.map{  
    case(zip, statuses) =>  
        val normalCount = statuses.count(_ == 200)  
        val errorCount = statuses.size - normalCount  
        val errorRatio = errorCount.toFloat / statuses.size  
        if (errorRatio > 0.05) {  
            "%s: **%s**".format(zip, errorRatio)  
        } else {  
            "%s: %s".format(zip, errorRatio)  
        }  
}  
errorRatePerZipCode.print()
```

Spark Streaming Clickstream Example (using Kafka)



Spark Streaming Clickstream Example (using Kafka)

Scheduler
Event Loop



Spark Streaming Clickstream Example (using Kafka)

Shuffle Writes

Scheduler
Event Loop

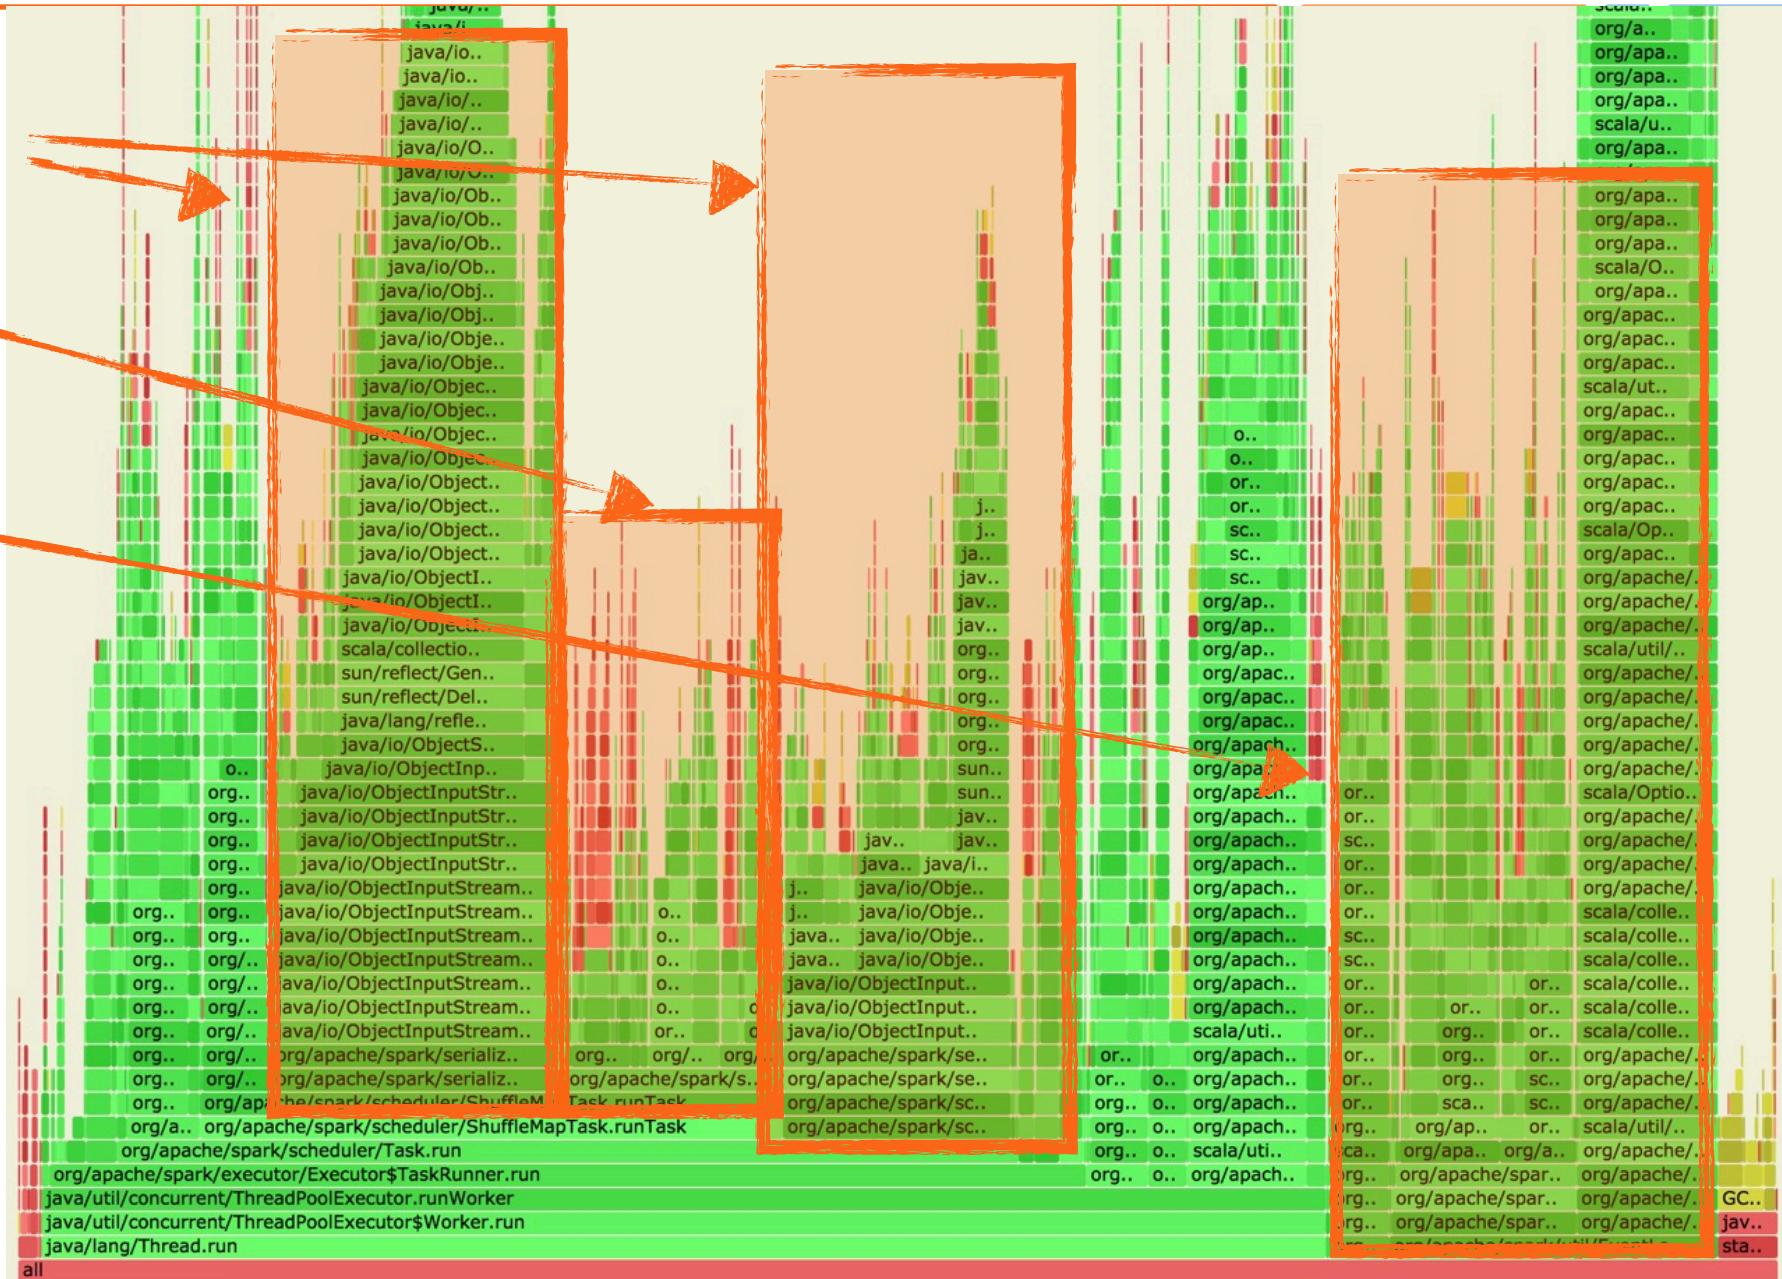


Spark Streaming Clickstream Example (using Kafka)

30% deserialization

Shuffle Writes

Scheduler
Event Loop



Spark Streaming Clickstream Example (using Kafka)

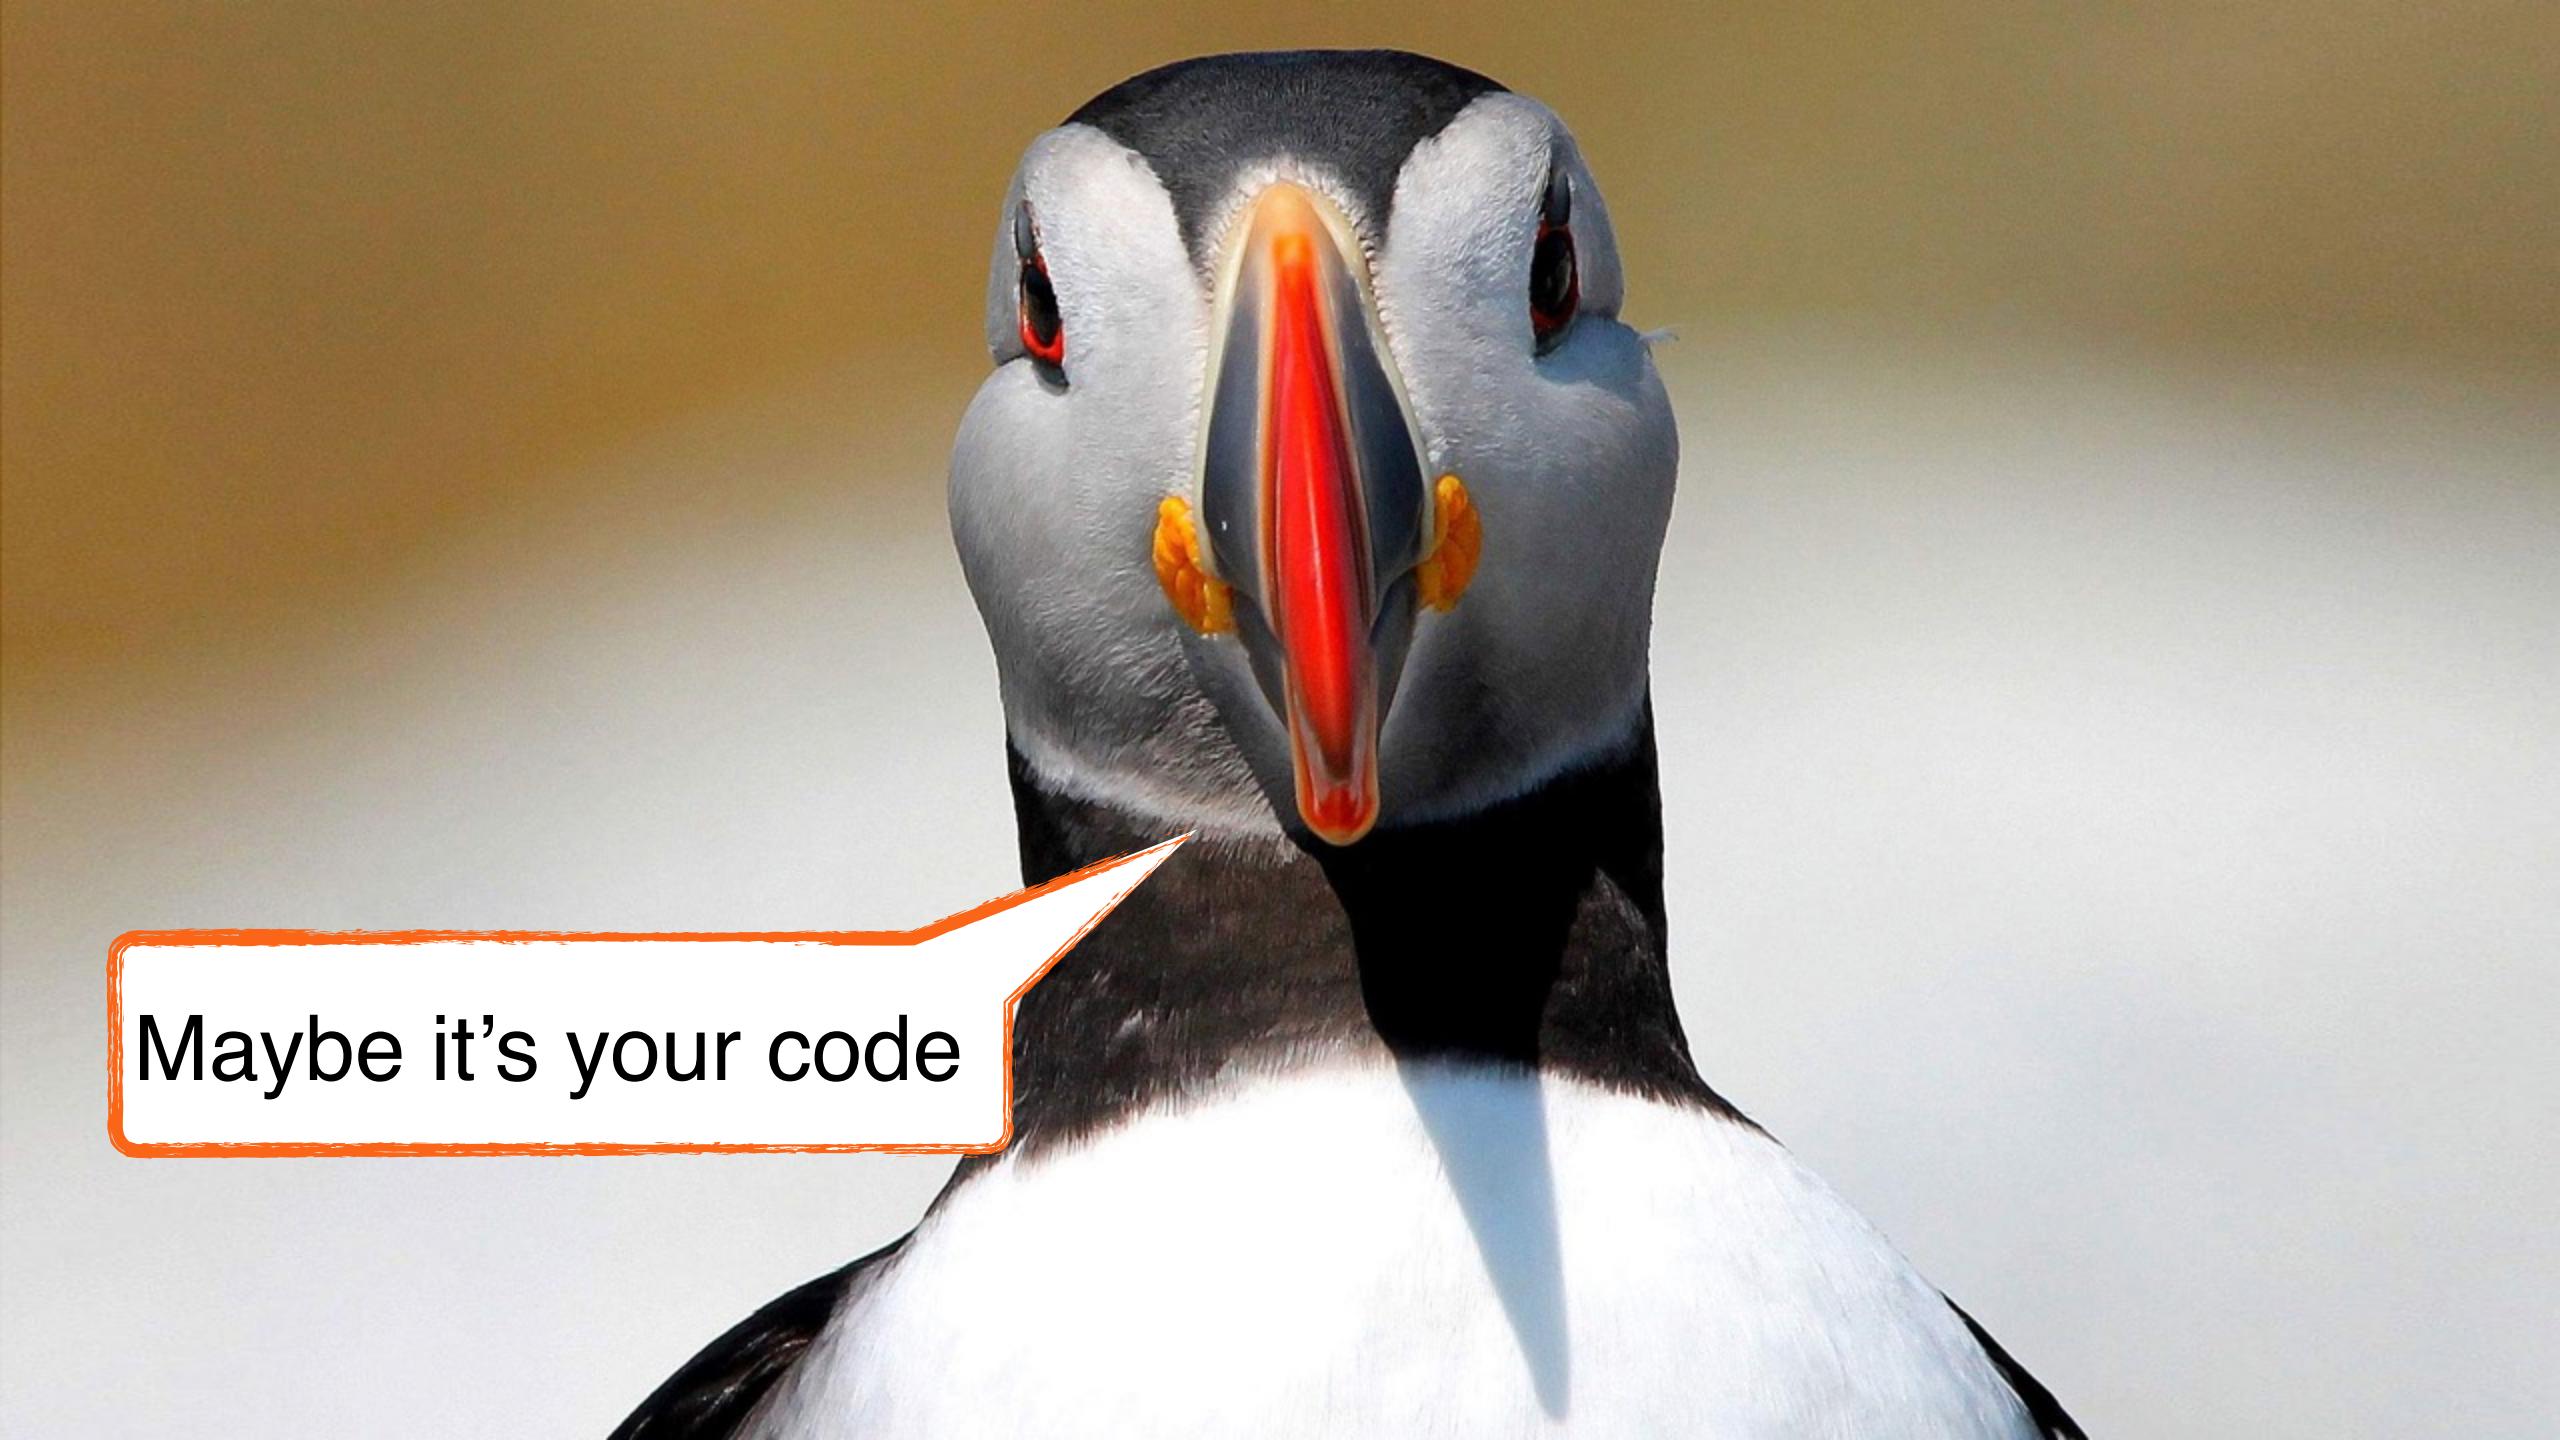
30% deserialization

Shuffle Writes

Scheduler
Event Loop

Read from Kafka
& Processing





Maybe it's your code

Let's commit, just to be safe, right?



Common beginner mistake

Commit only as needed
keep recovery short
maximize throughput

Metrics to validate
commit-rate
commit-latency-avg

Right-size your batches



Right-size your batches



Bigger Batches
increase throughput
improve compression

Right-size your batches



Bigger Batches

increase throughput

improve compression

Small enough (<< 10MB) to keep GC low

Right-size your batches



Bigger Batches

increase throughput

improve compression

Small enough (<< 10MB) to keep GC low

batch.size + **linger.ms**

Right-size your batches



Bigger Batches

increase throughput

improve compression

Small enough (<< 10MB) to keep GC low

batch.size + **linger.ms**

don't forget!

Right-size your batches



Bigger Batches

increase throughput

improve compression

Small enough (<< 10MB) to keep GC low

batch.size + **linger.ms**

don't forget!

Watch

request-rate

request-latency-avg

compression-rate

My app keeps rebalancing



Symptoms

low throughput

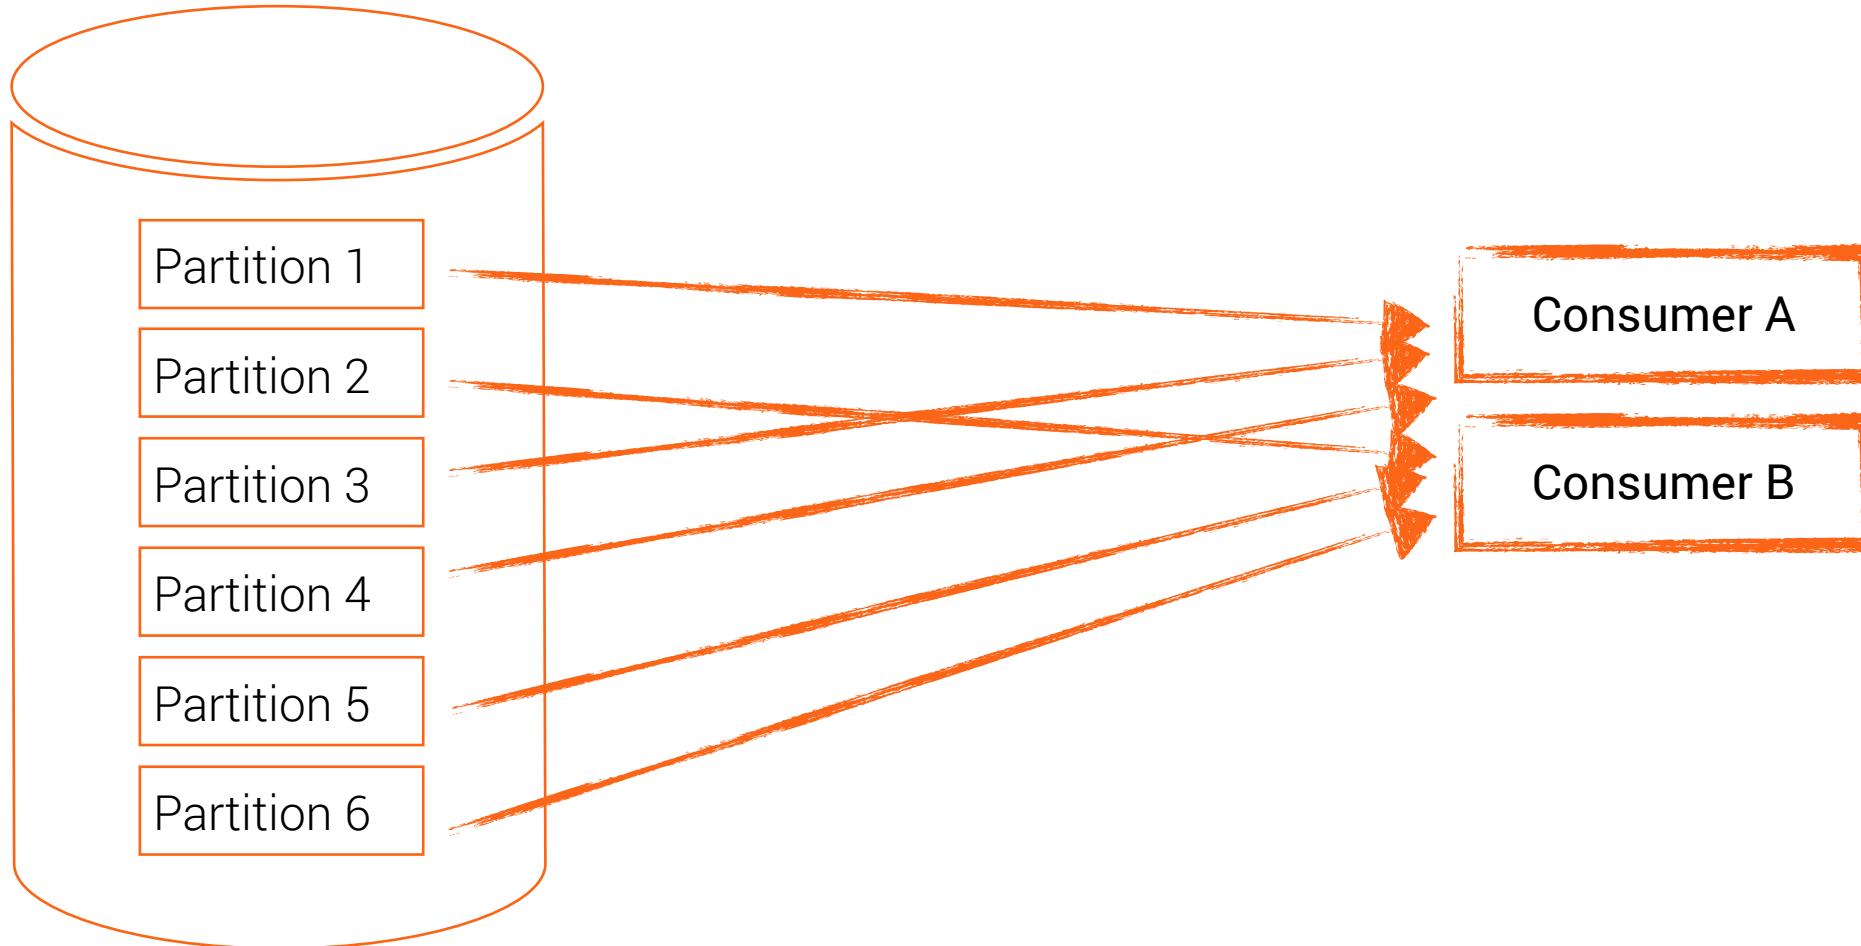
high network chatter

consumer logs galore

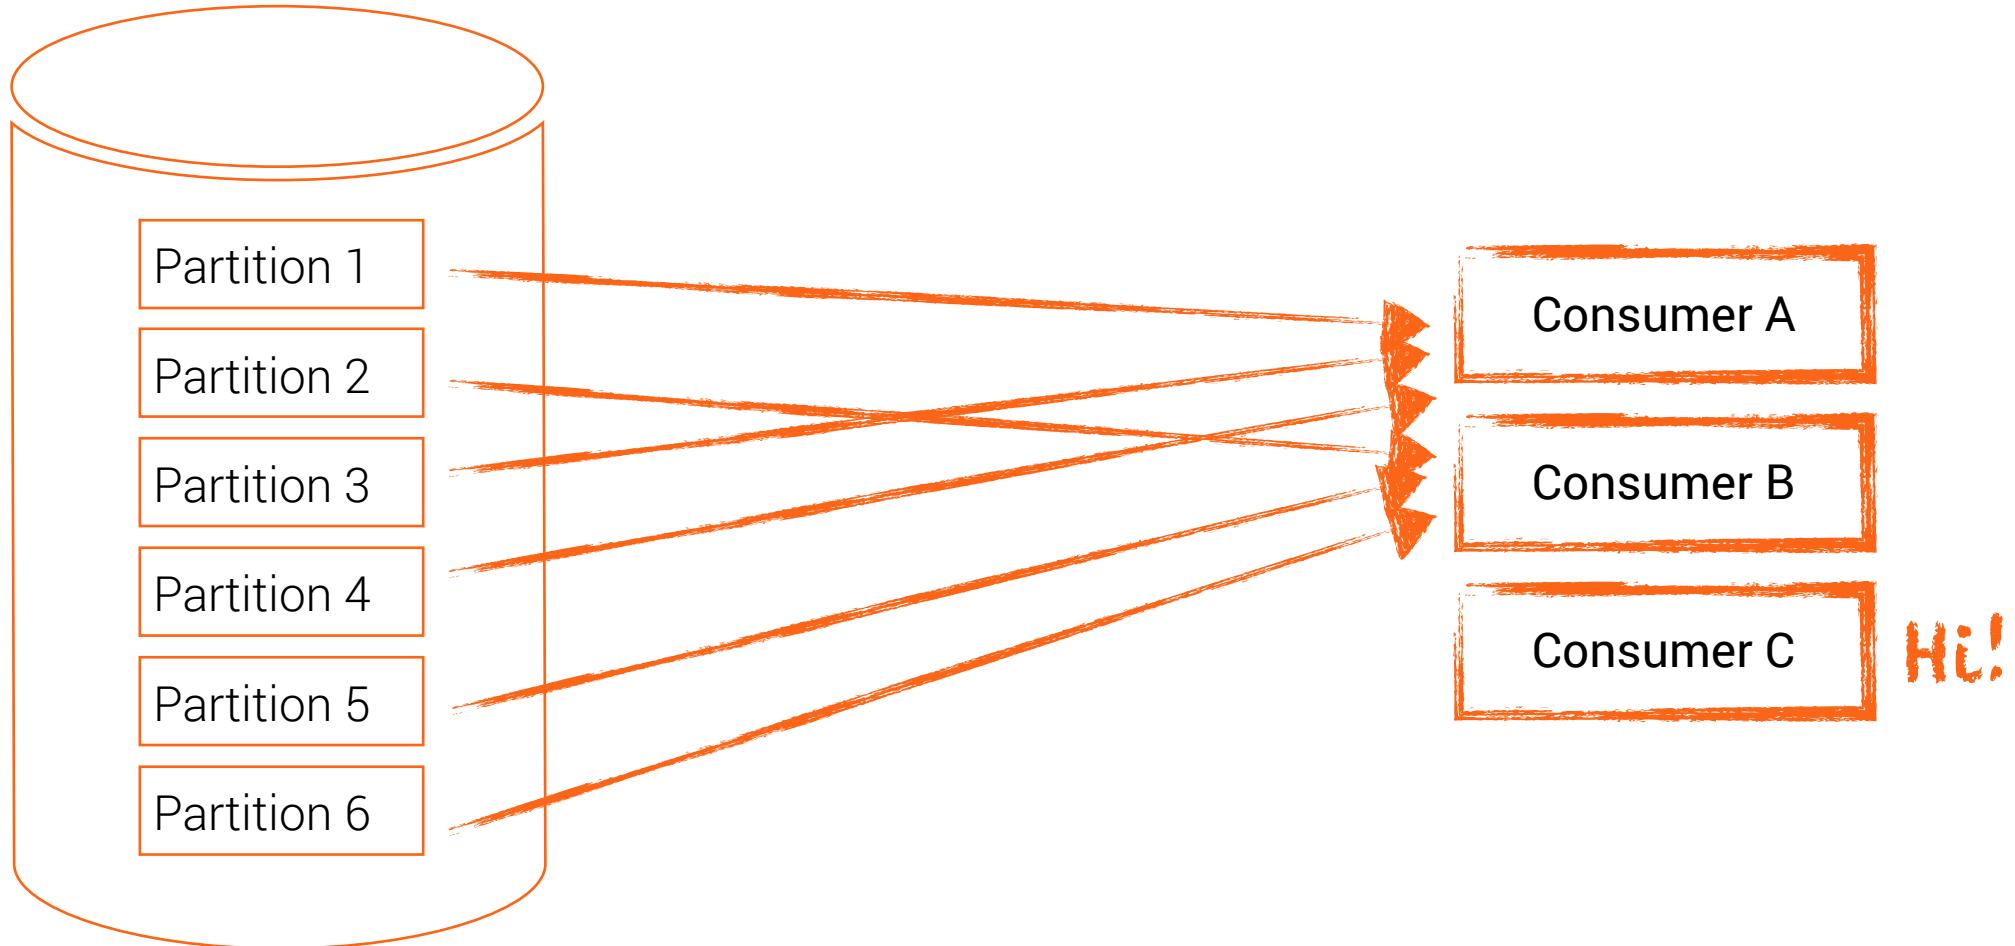
no progress

hanging

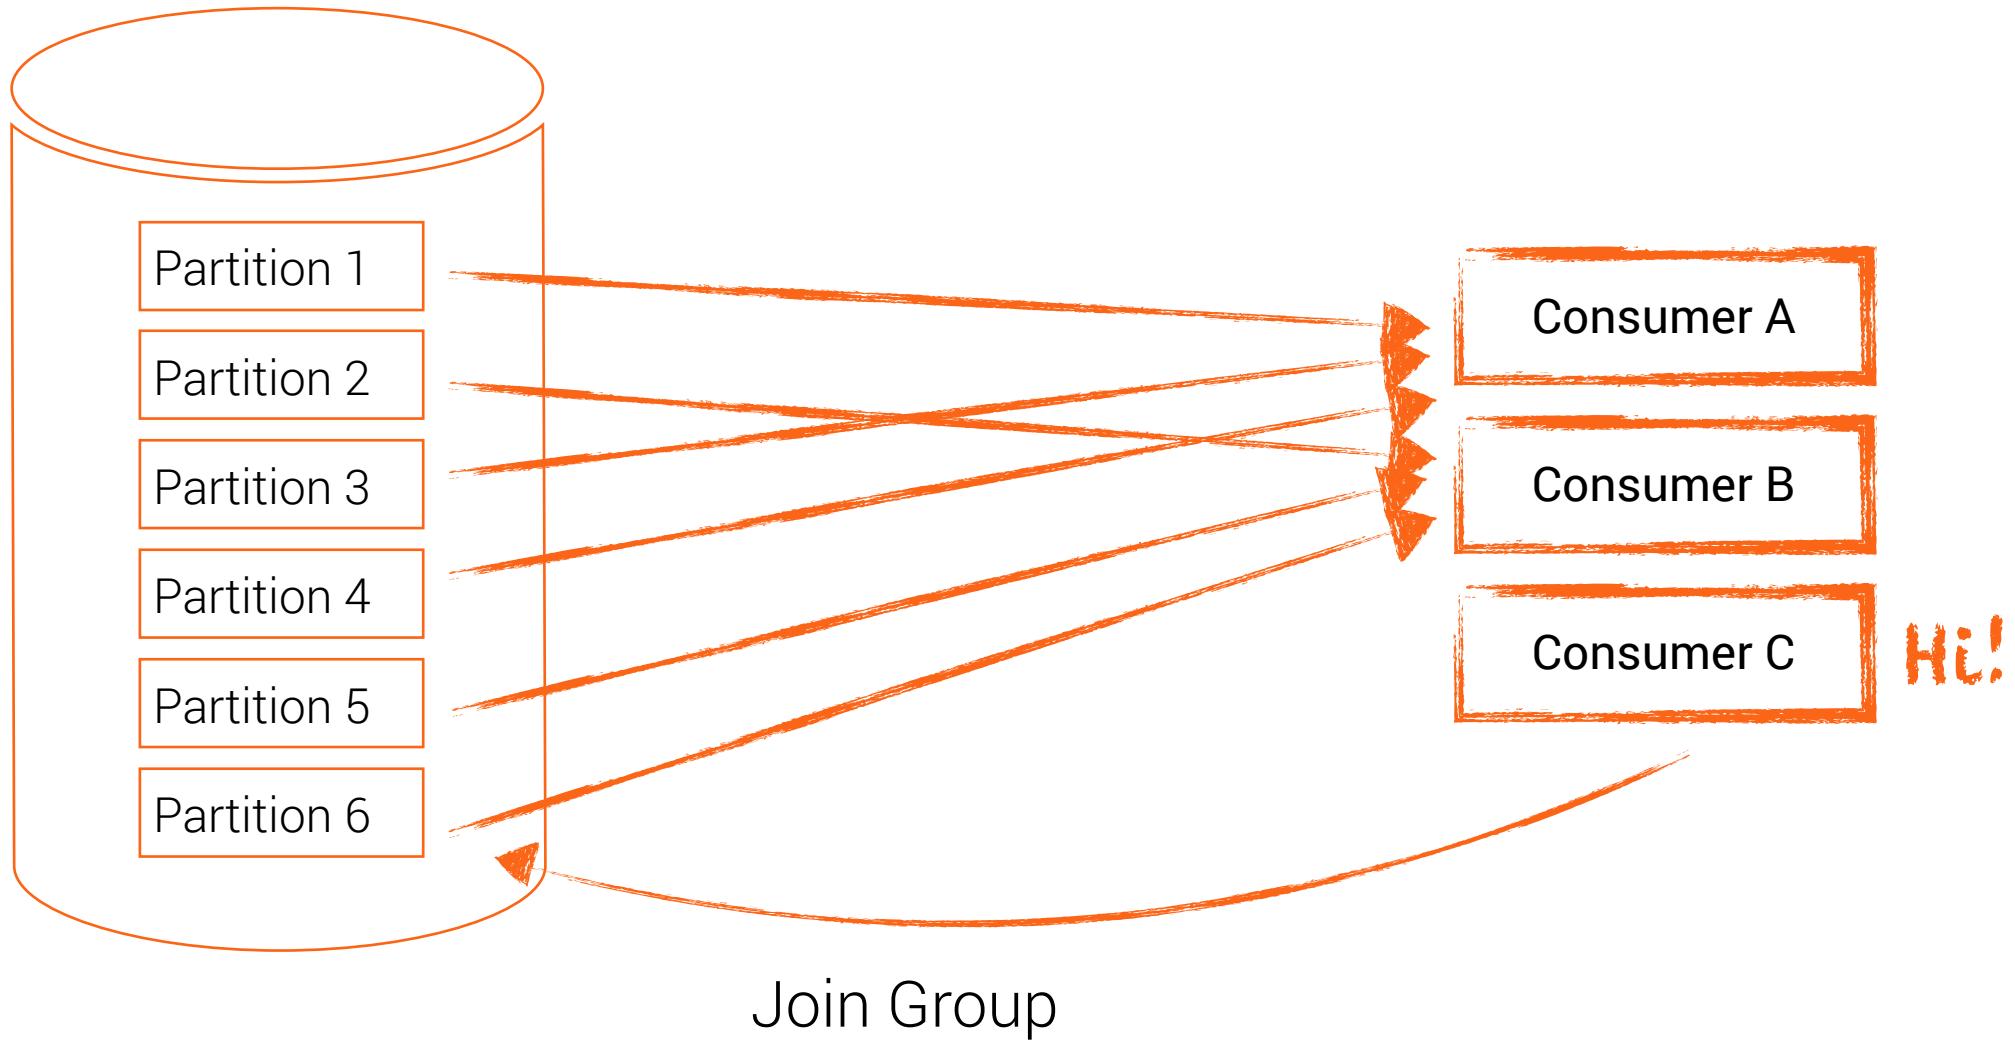
Kafka Consumer Group Rebalancing 101



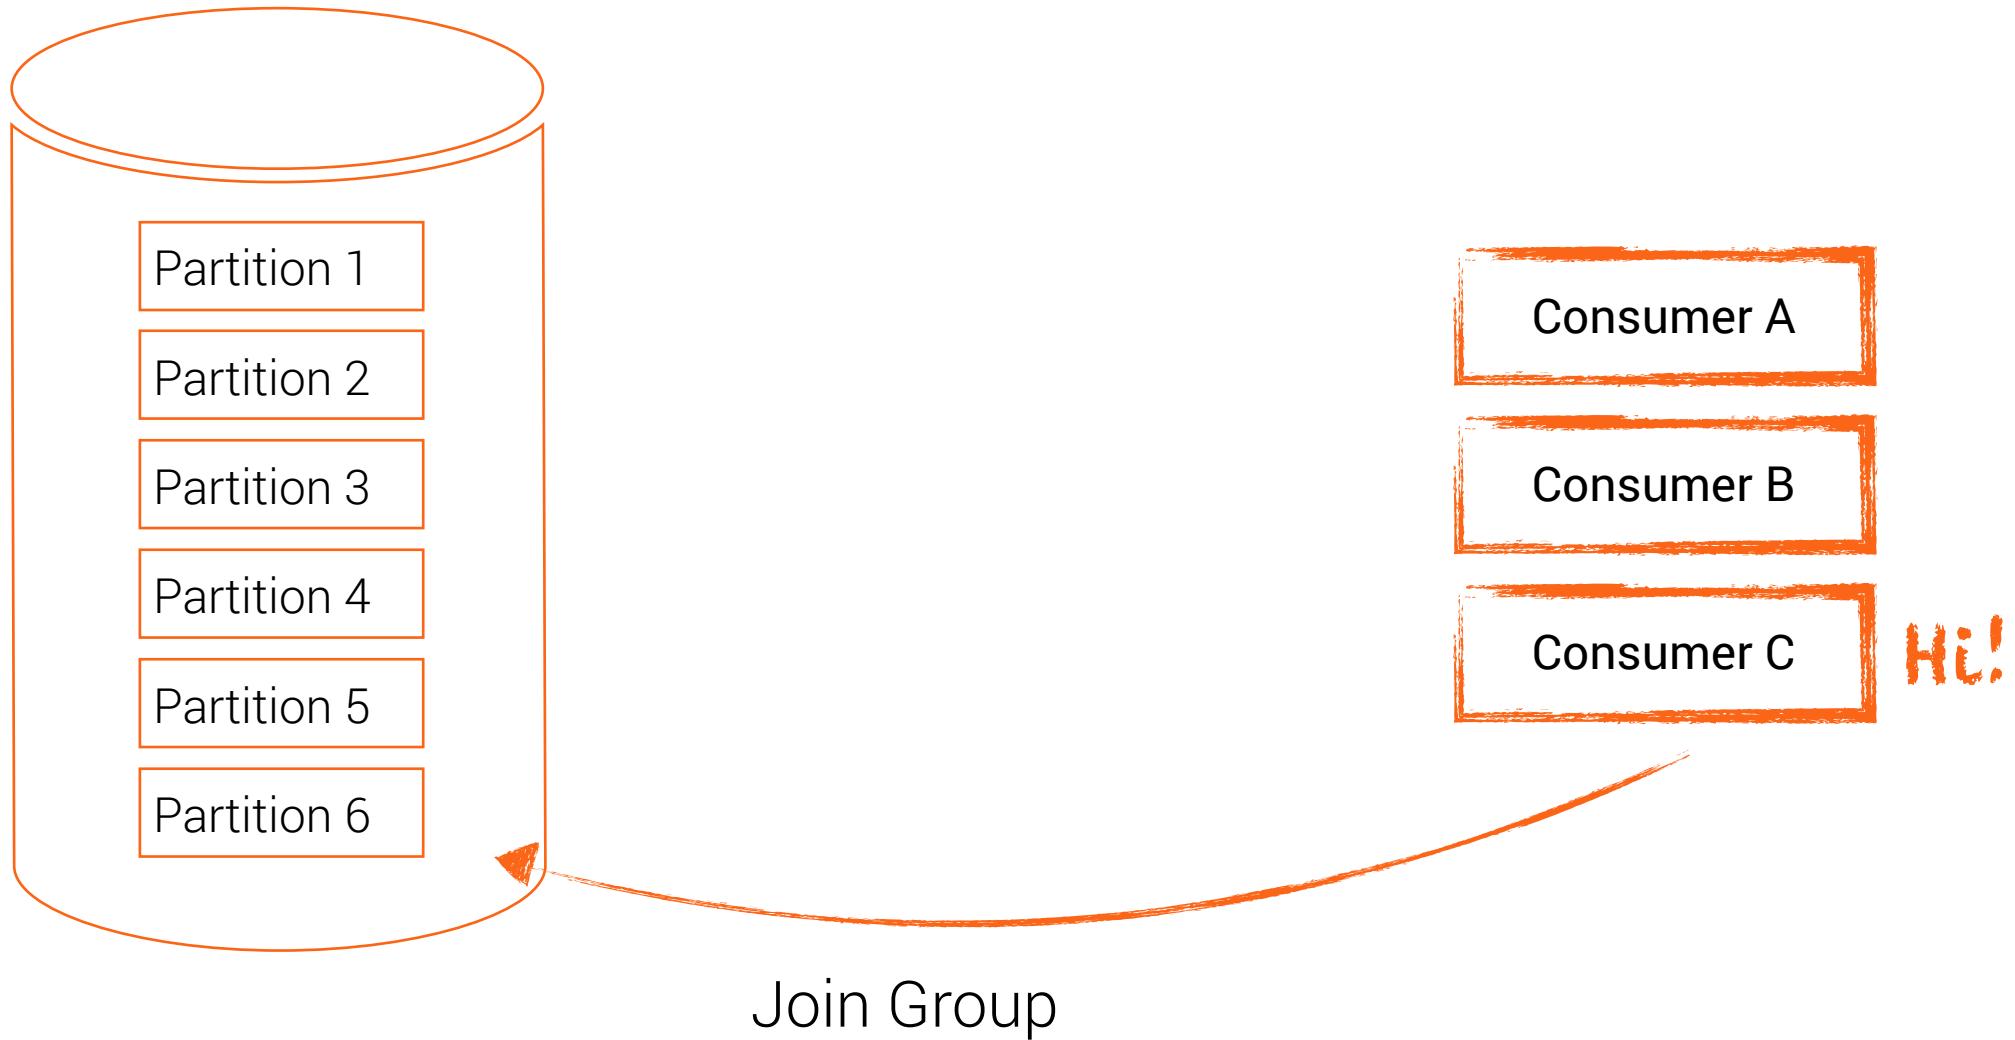
Kafka Consumer Group Rebalancing 101



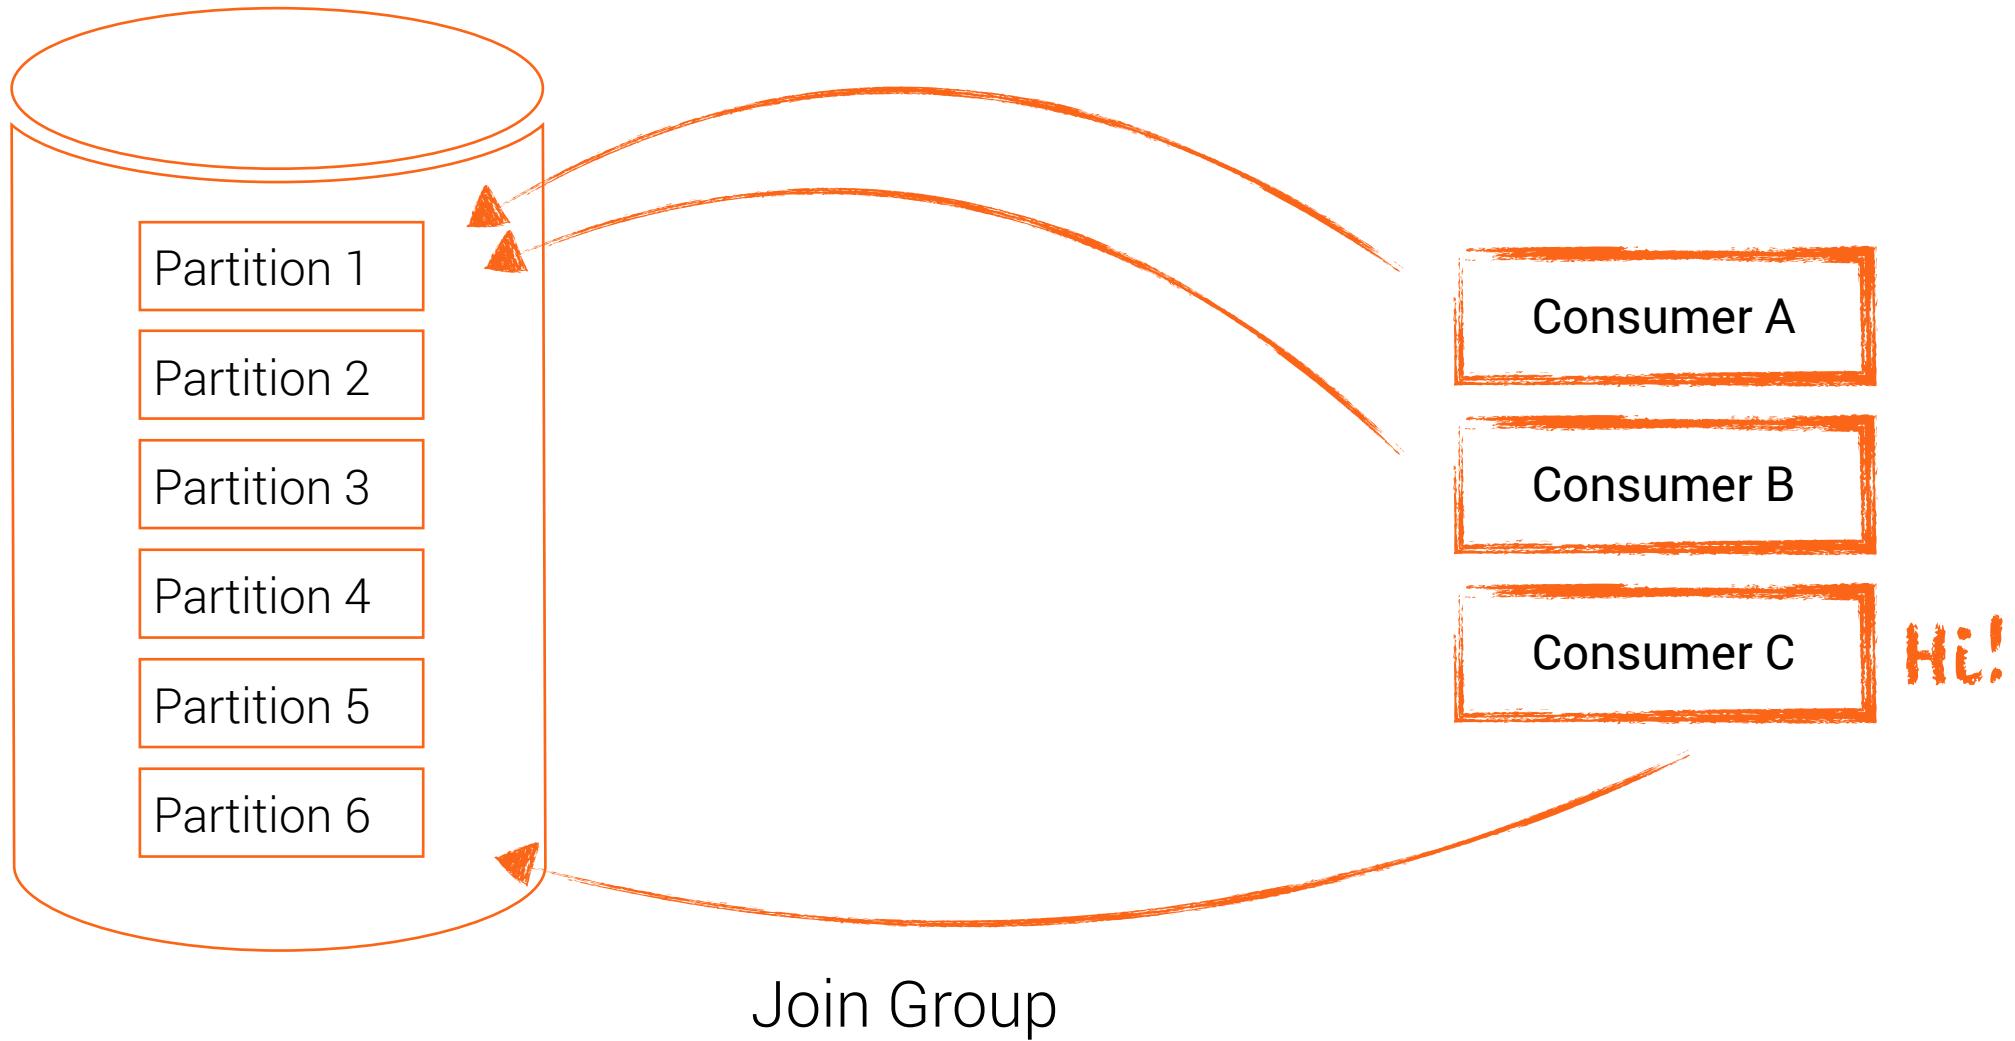
Kafka Consumer Group Rebalancing 101



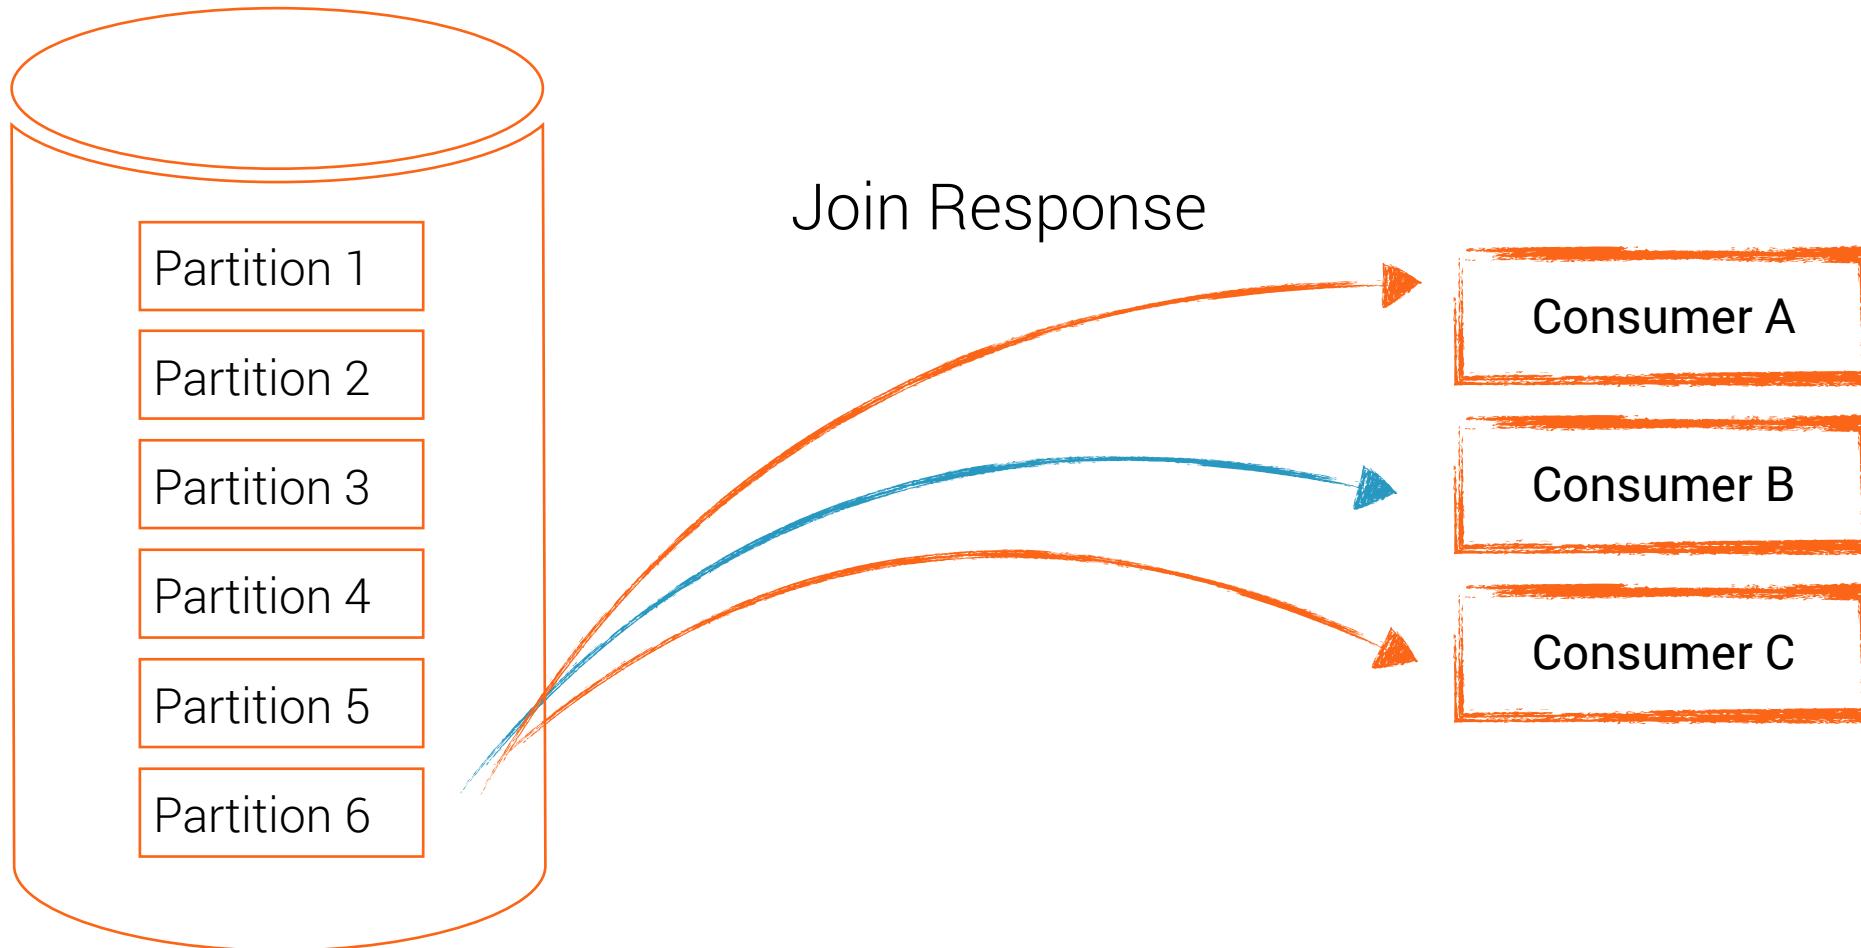
Kafka Consumer Group Rebalancing 101



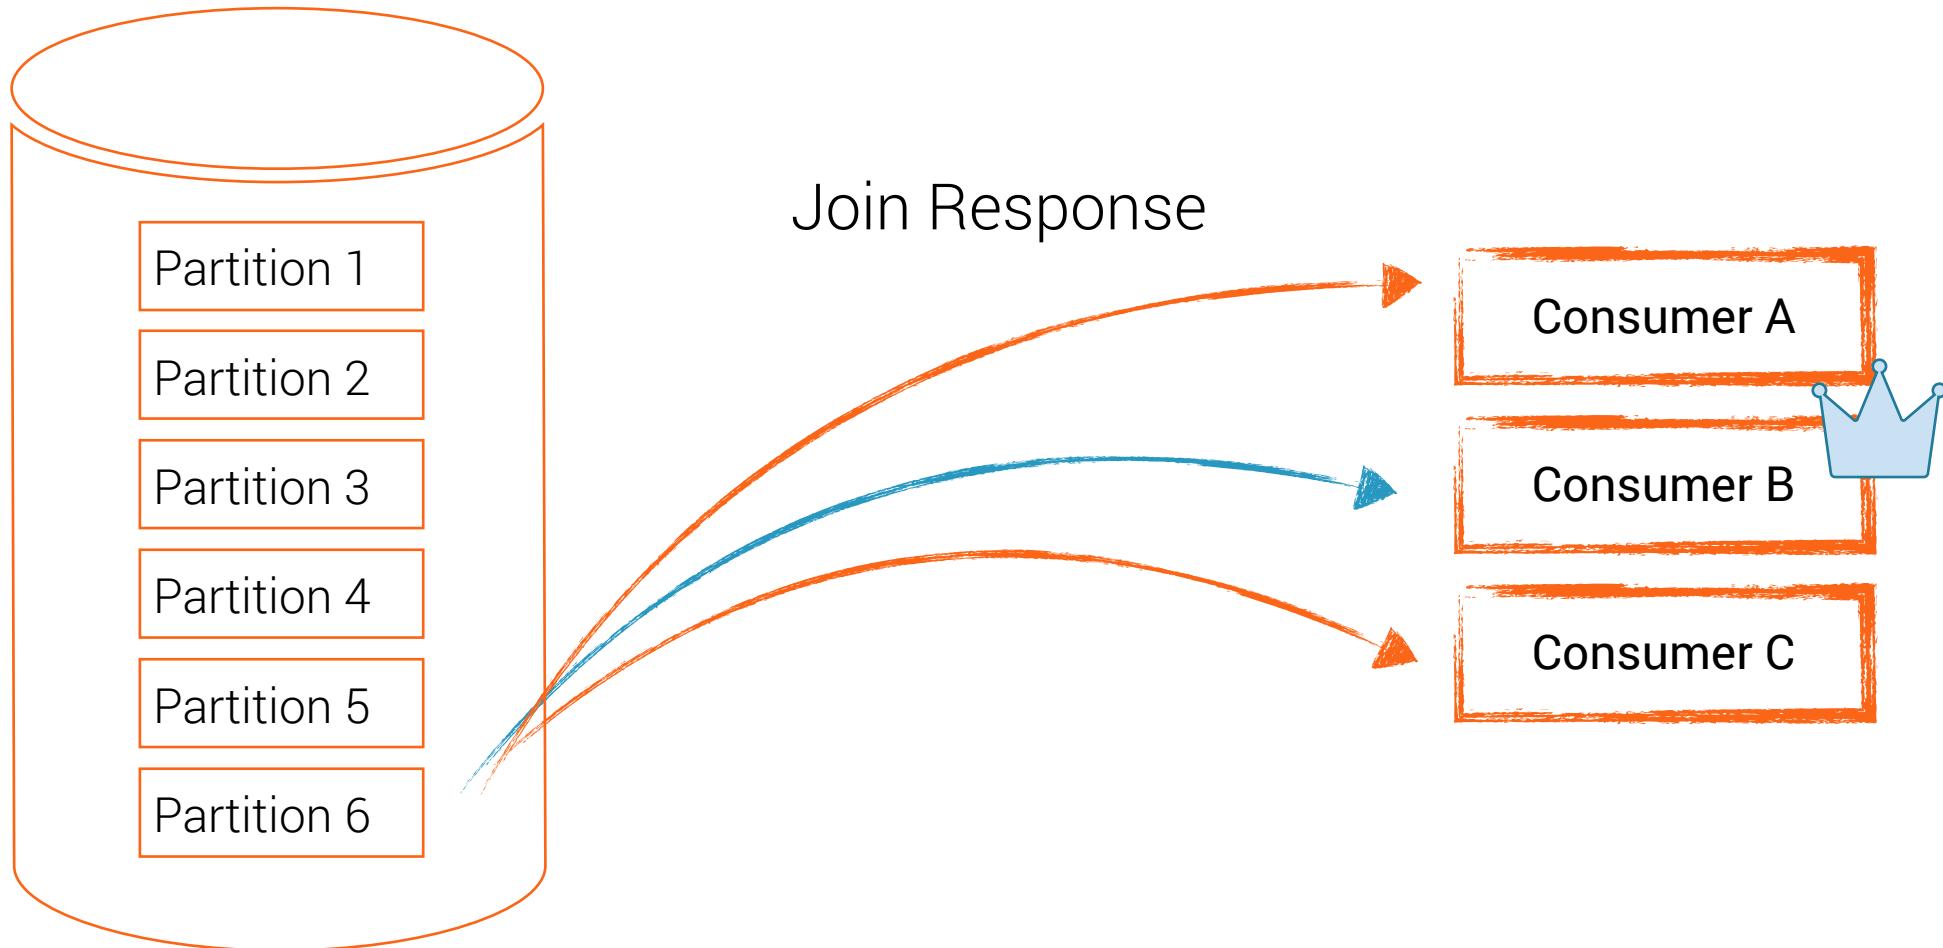
Kafka Consumer Group Rebalancing 101



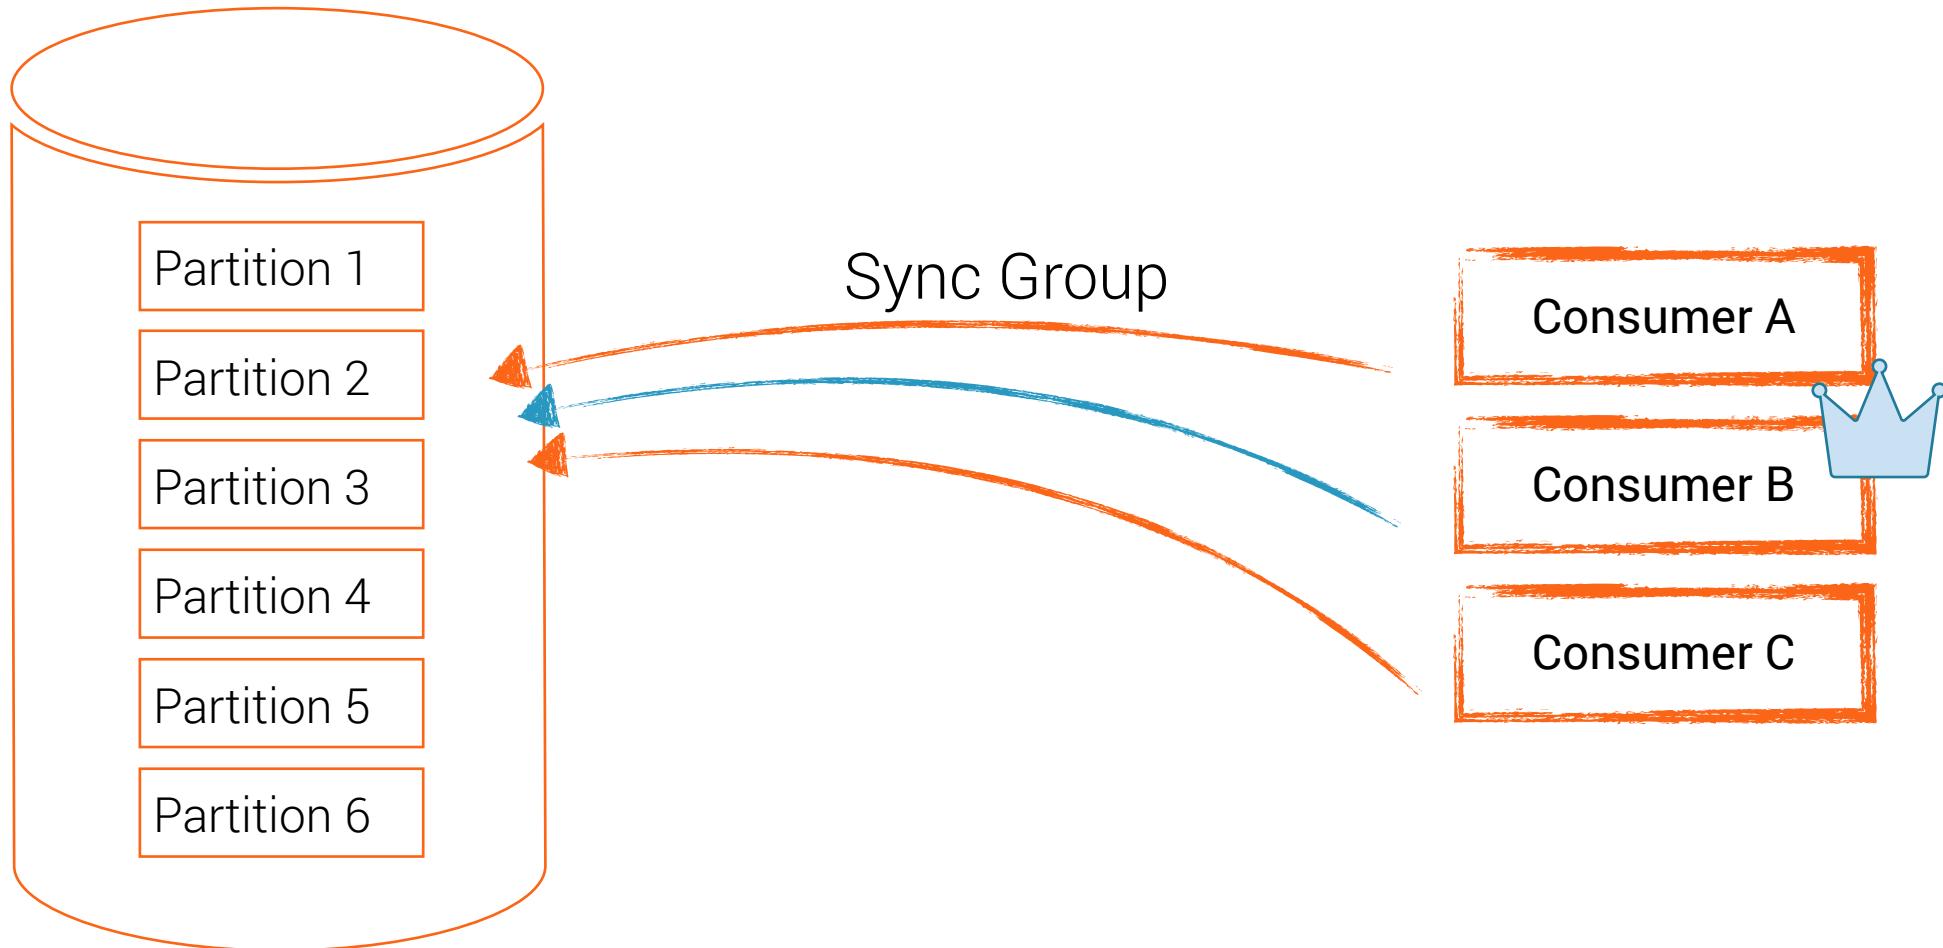
Kafka Consumer Group Rebalancing 101



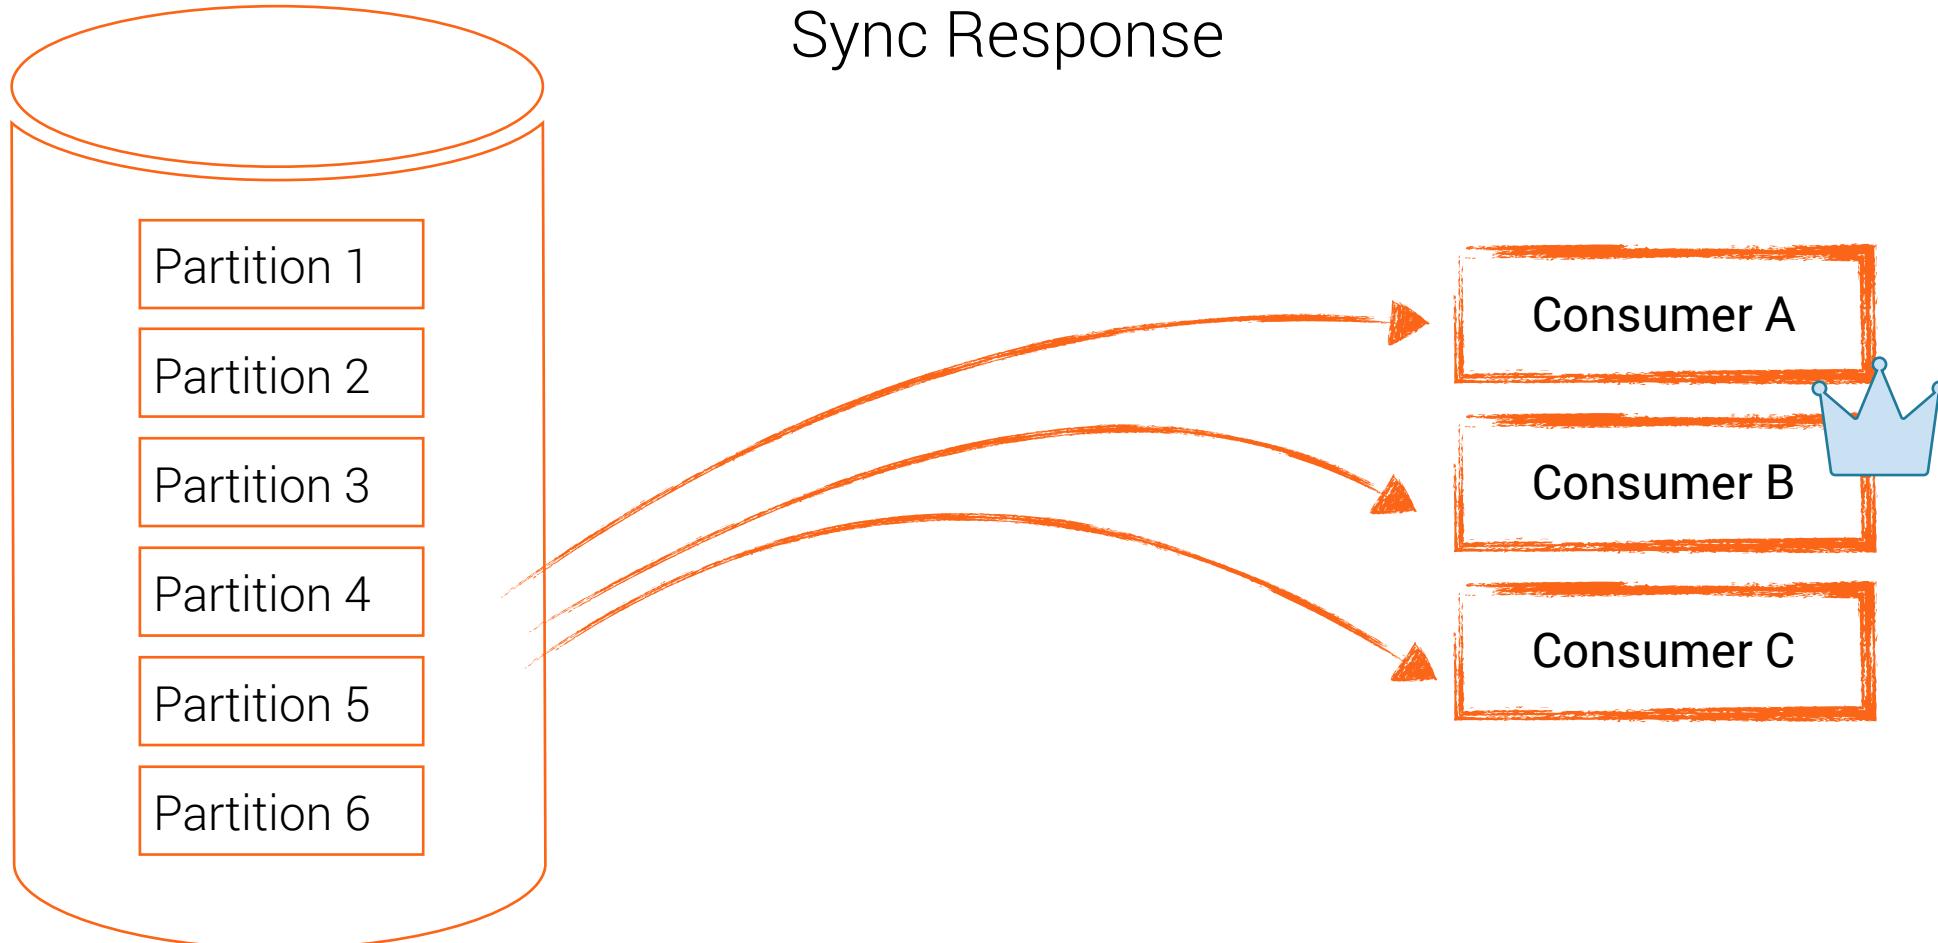
Kafka Consumer Group Rebalancing 101



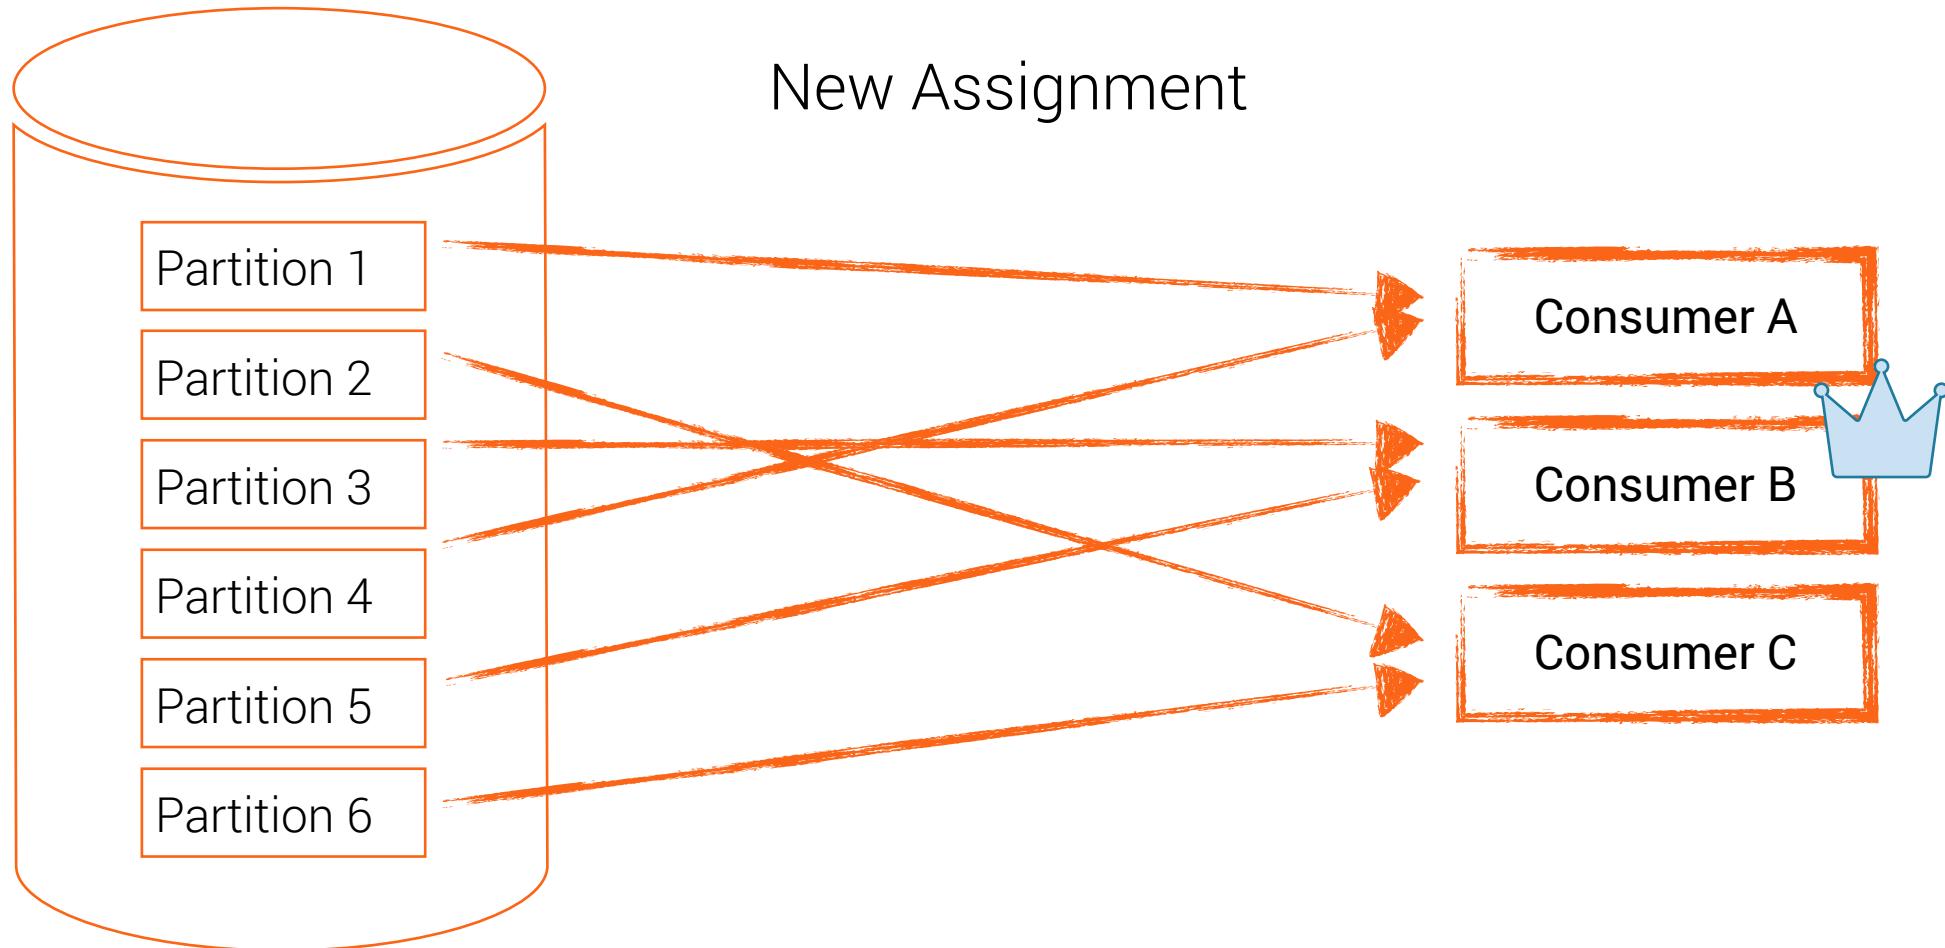
Kafka Consumer Group Rebalancing 101



Kafka Consumer Group Rebalancing 101



Kafka Consumer Group Rebalancing 101



Restoring a Happy Balance



Timing Issues

- long GC pauses (tens of seconds)
- infrequent calls to `poll()`
- timeouts too short?
- flaky network
- 1 bad machine affects the entire group

Watch

- join-rate
- sync-rate

Competent Users



- Monitor Consumer Lag
- Lookout for Partition Skew
- Commit Offsets Sparingly
- Collect Logs
- Understand how to tune Batch Sizes

Kafka Pros



- Watch Group Partition Assignment
- Monitor Client Metrics
- Understand Consumer Rebalancing
- Profile their applications
- Distinguish Client/App/Broker problems

~~Replication~~ Everything is Slow

Famous last words...

"You just consume, and produce. How hard can this be?"



Famous last words...



"We have a disaster in our main cluster. Can we fail over to secondary? We can't lose more than 7 seconds of data."

Monitor Replication Lag - In messages

MONITORING > CONSUMER LAG >
NY_PHX_Replicator

15,007 Total messages behind

+231 messages behind
2,009 produced - 1,994 read

Set an alert

ad_clicks_agg_1min
Max lag /consumer: 421

user_profile_updates
Max lag /consumer: 3,421

Consumer's location on a partition

Consumer	Topic partition	Lag		
	Topic	Partition	Messages behind	Current offset
1	ad_clicks_agg_1min	1	1,241 of 3,292,199	1829198 of 1929933
2	ad_clicks_agg_1min	2	15,231 of 56,928,323	1282381 of 1318981
3	ad_clicks_agg_1min	3	905 of 4,292,199	2838289 of 3828289
4	ad_clicks_agg_1min	4	1,131 of 5,111,918	2888889 of 8188821

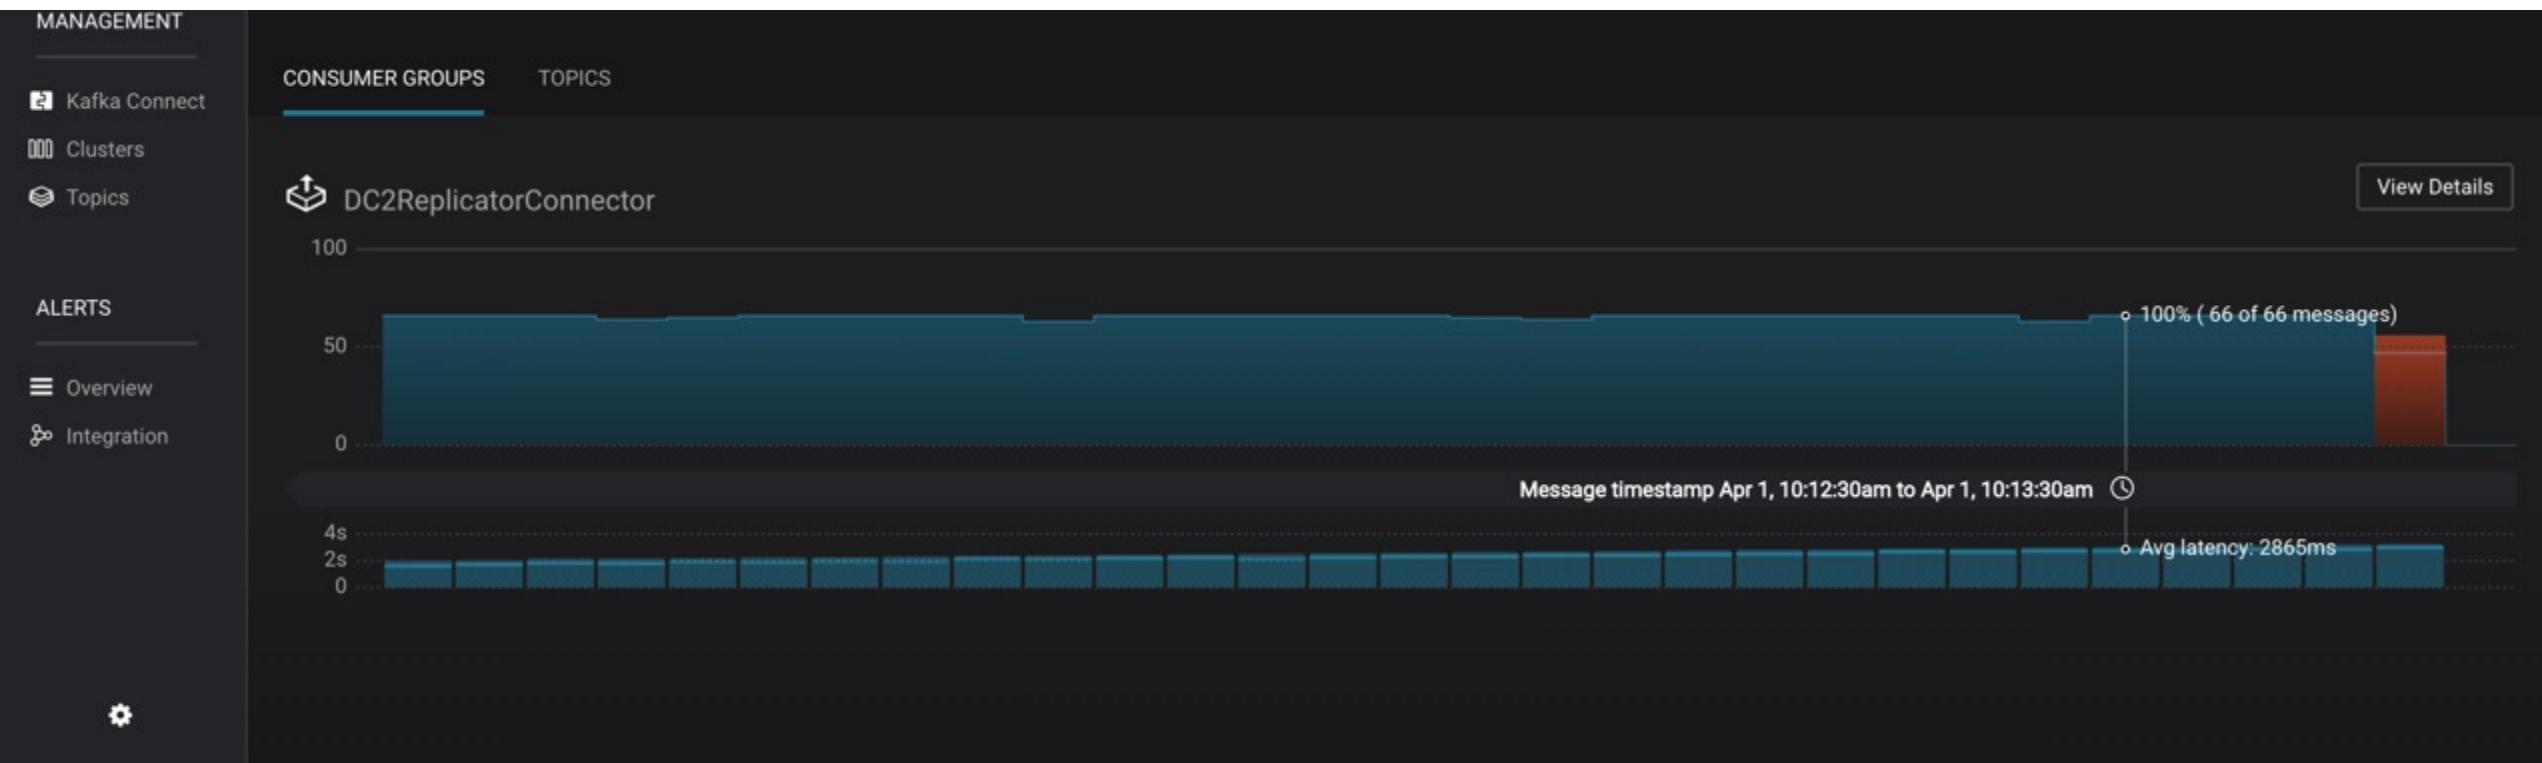
Addison Cooper

confluent

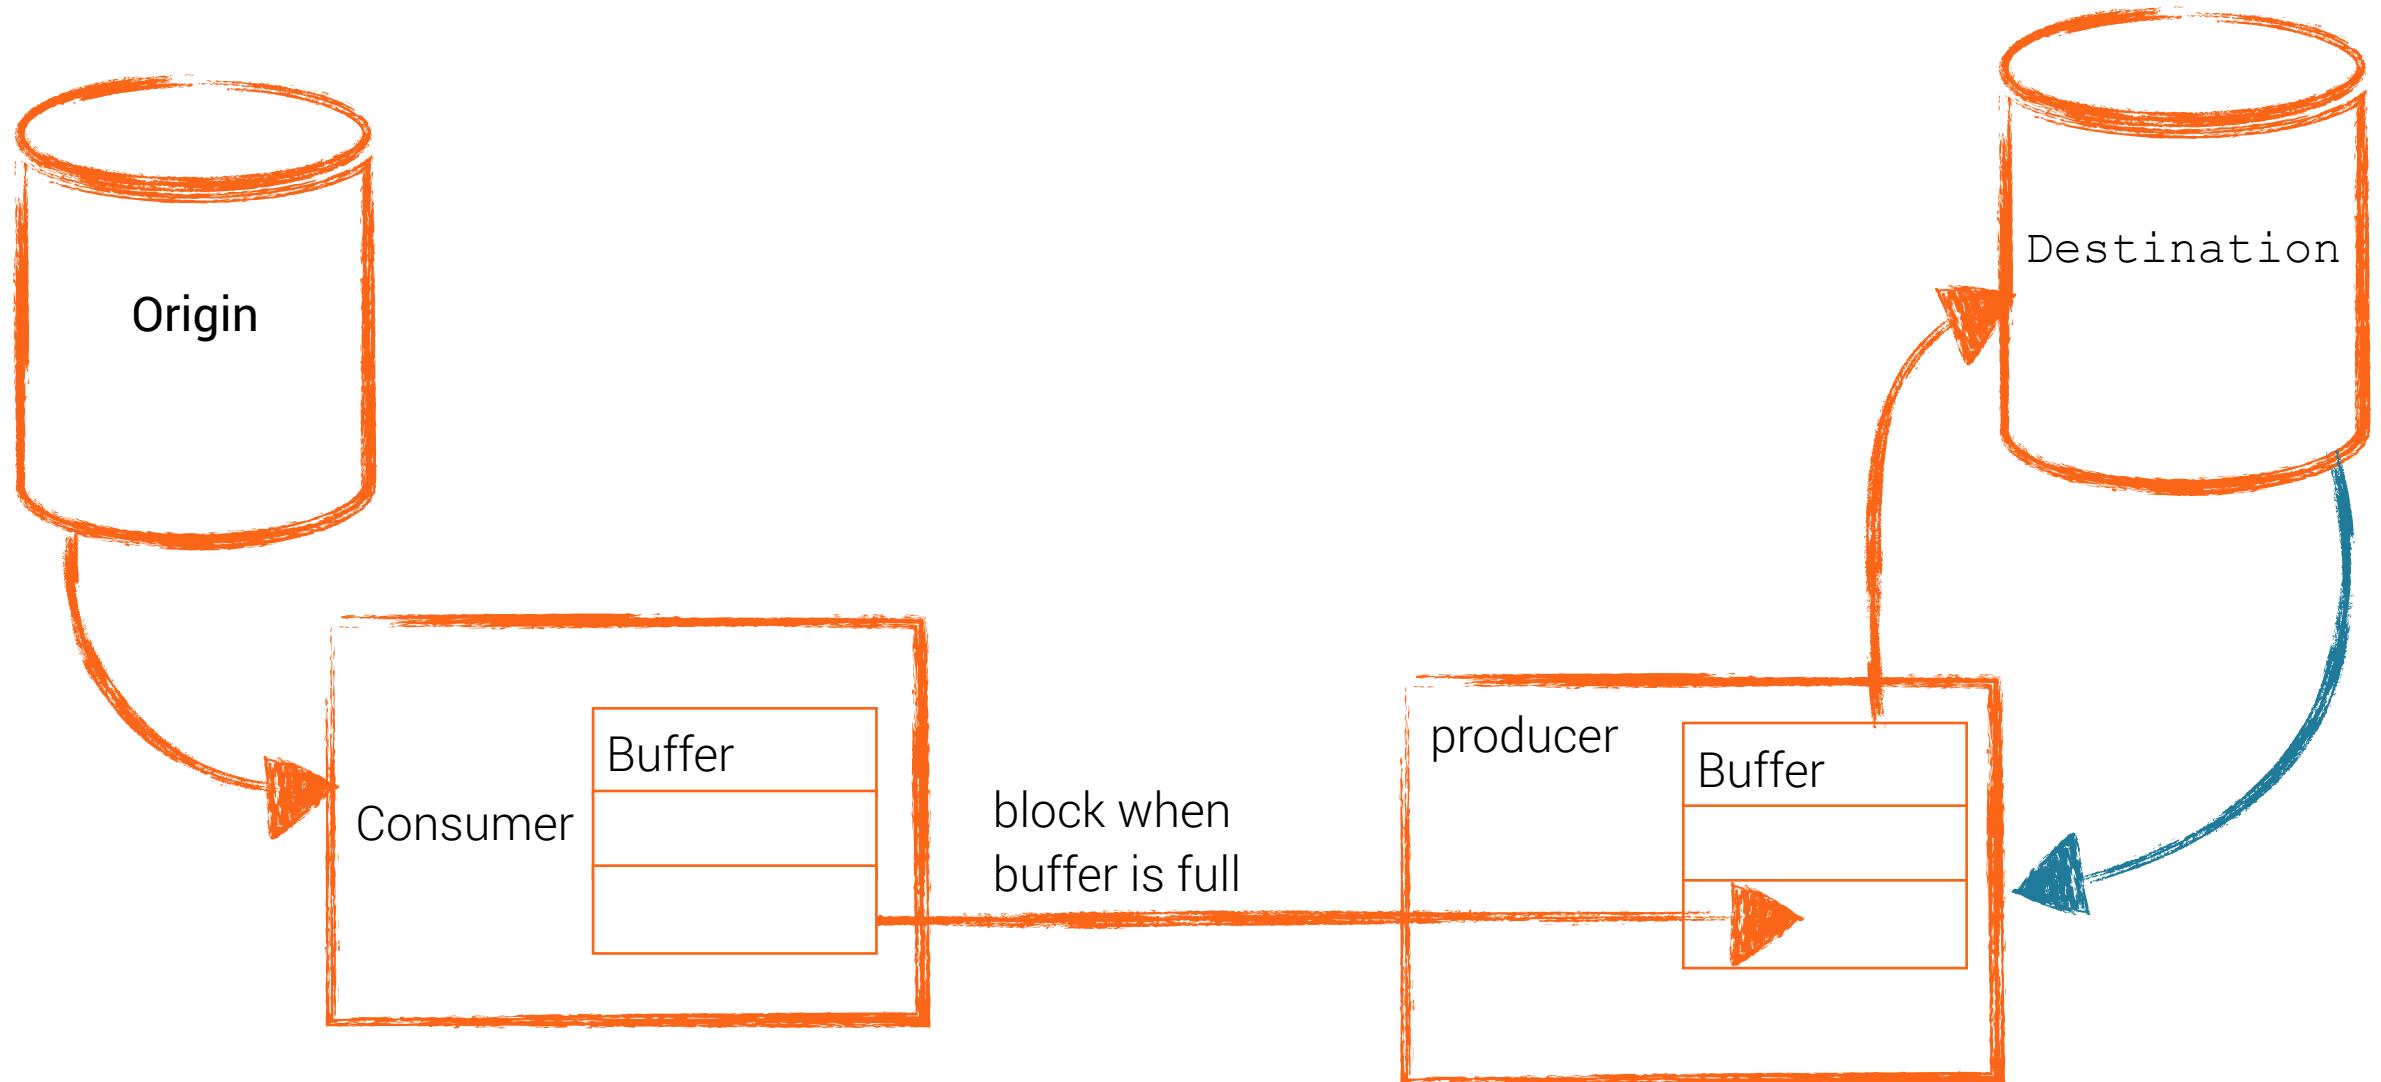
Cluster: NorthernCa

30

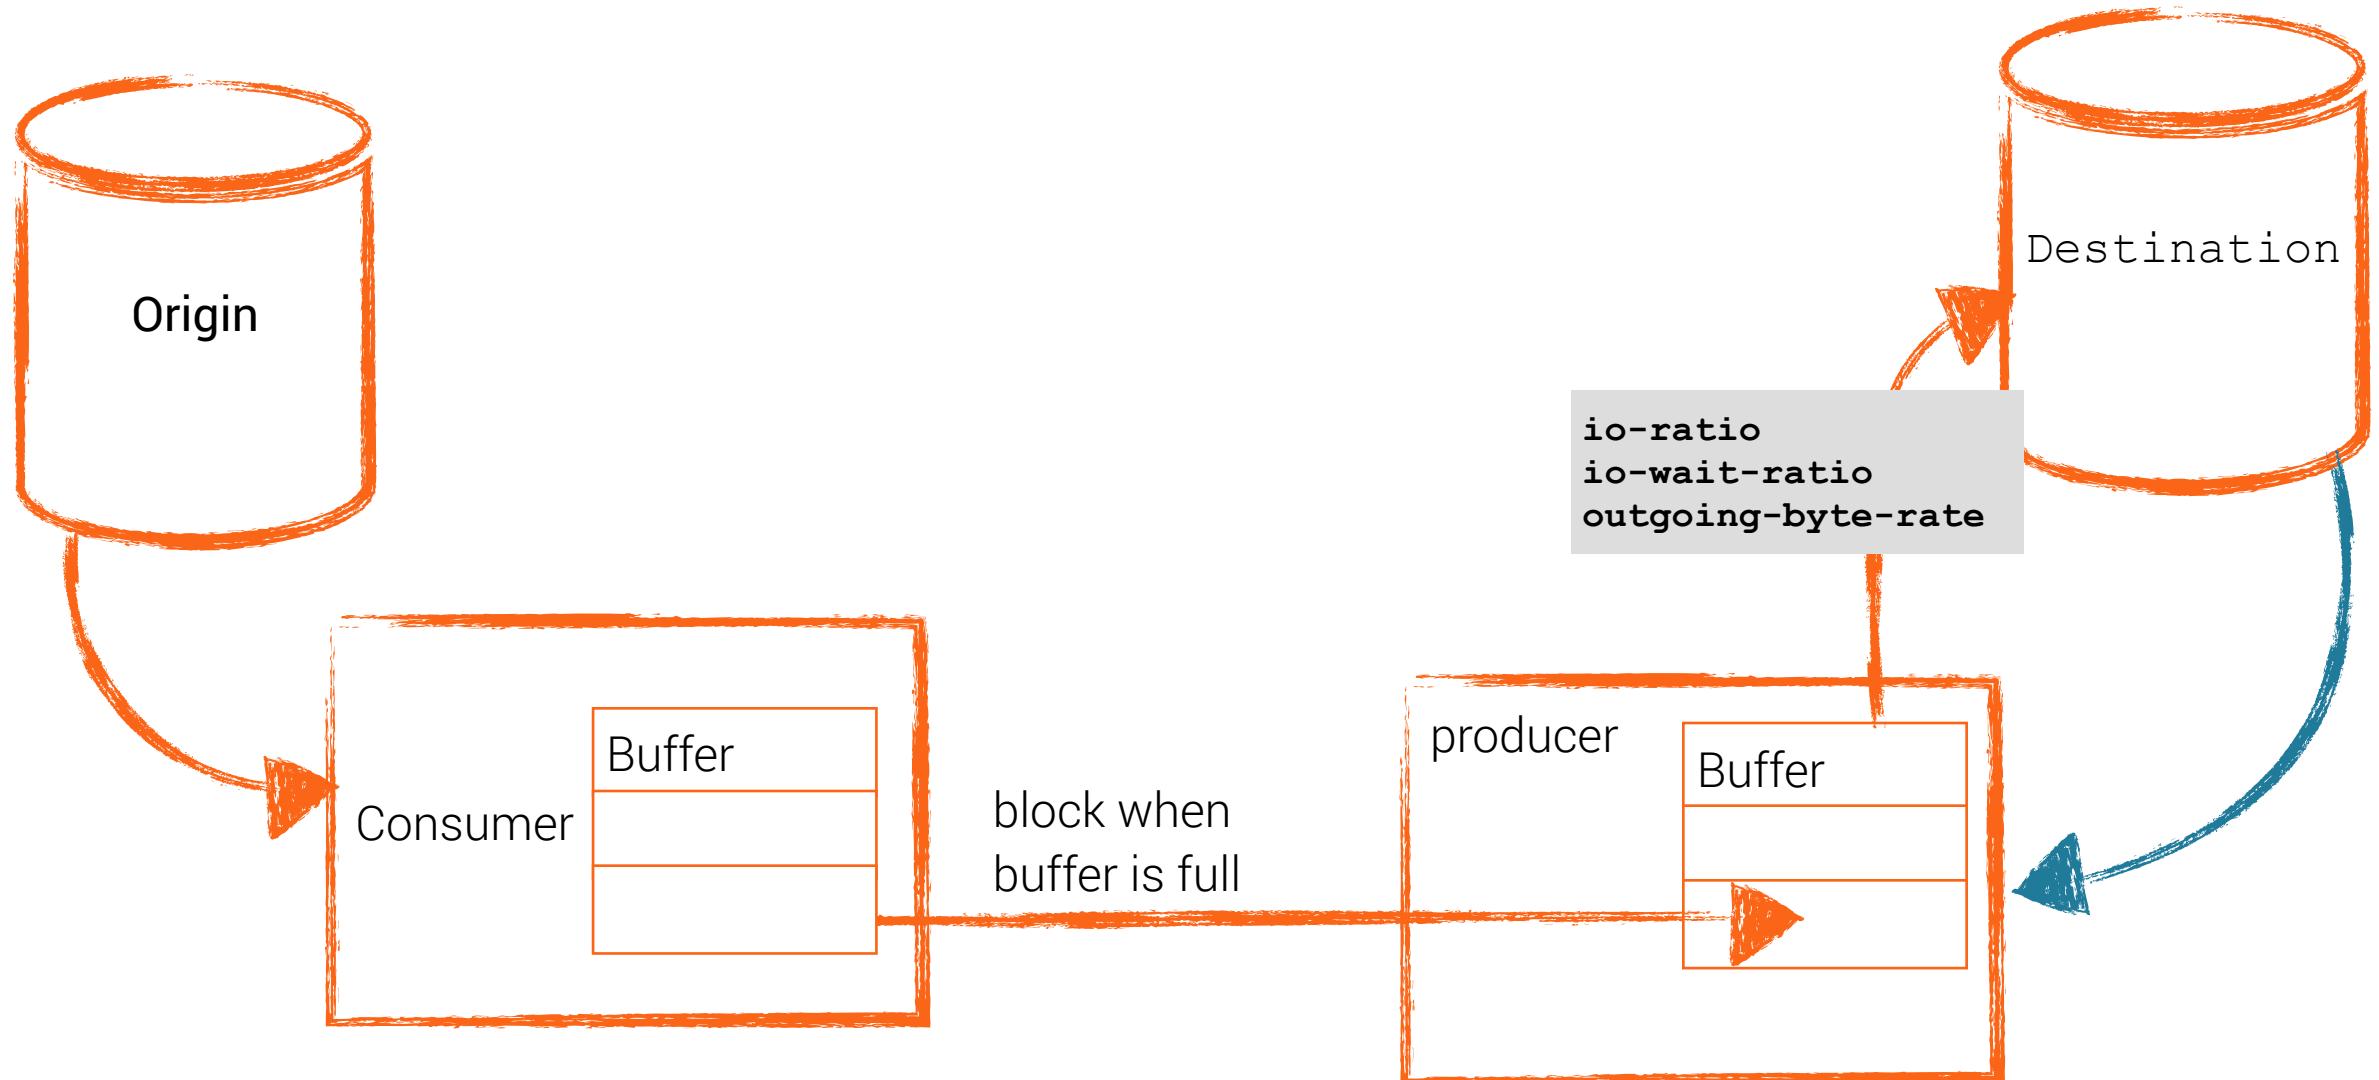
Monitor Replication Lag - or in seconds...



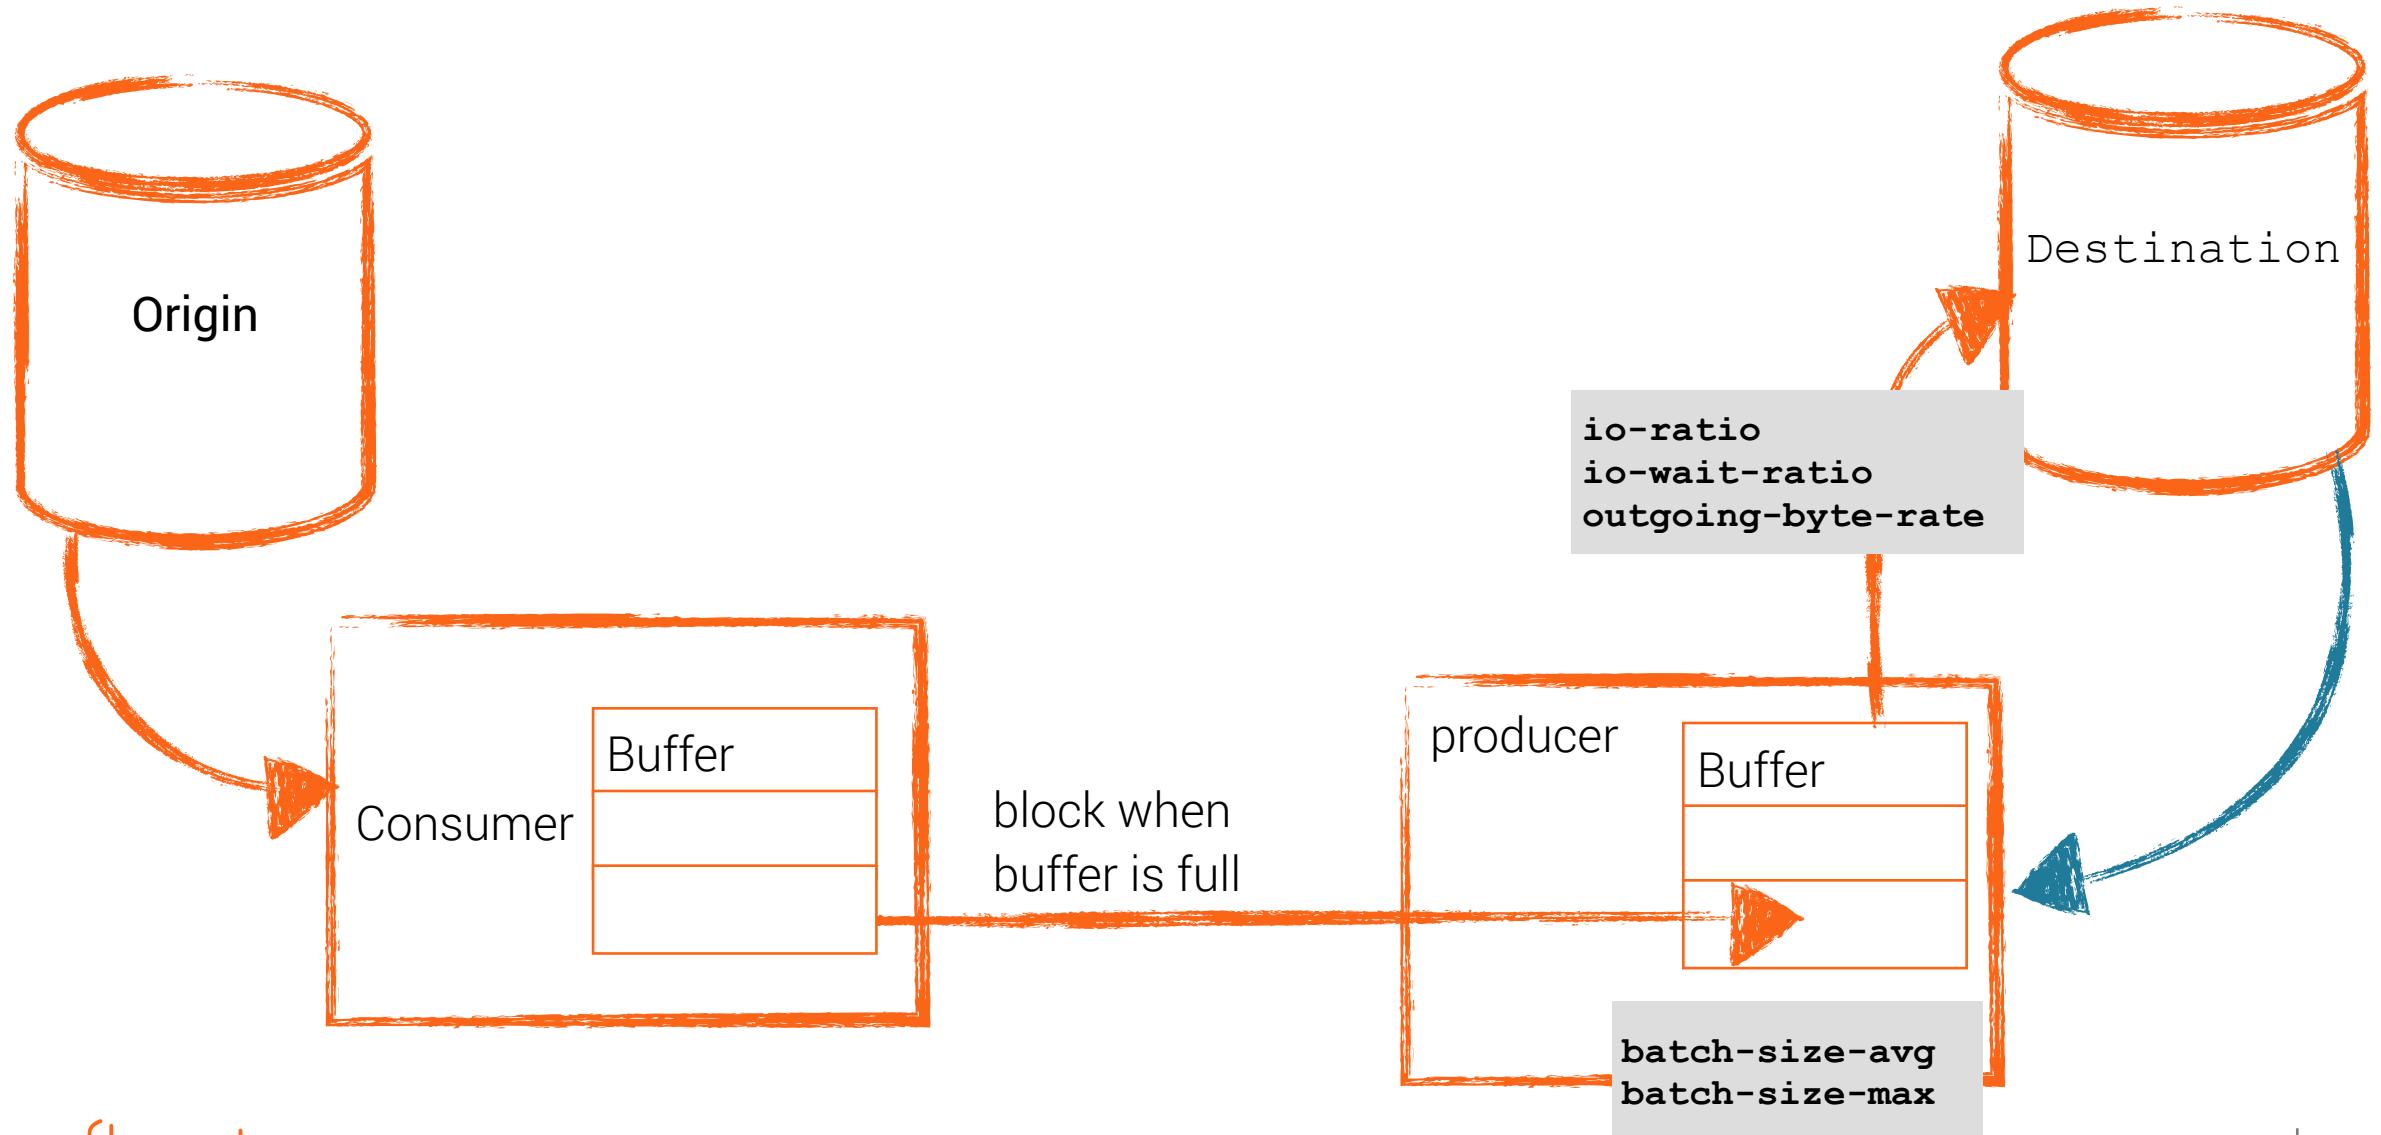
Simple and elegant design



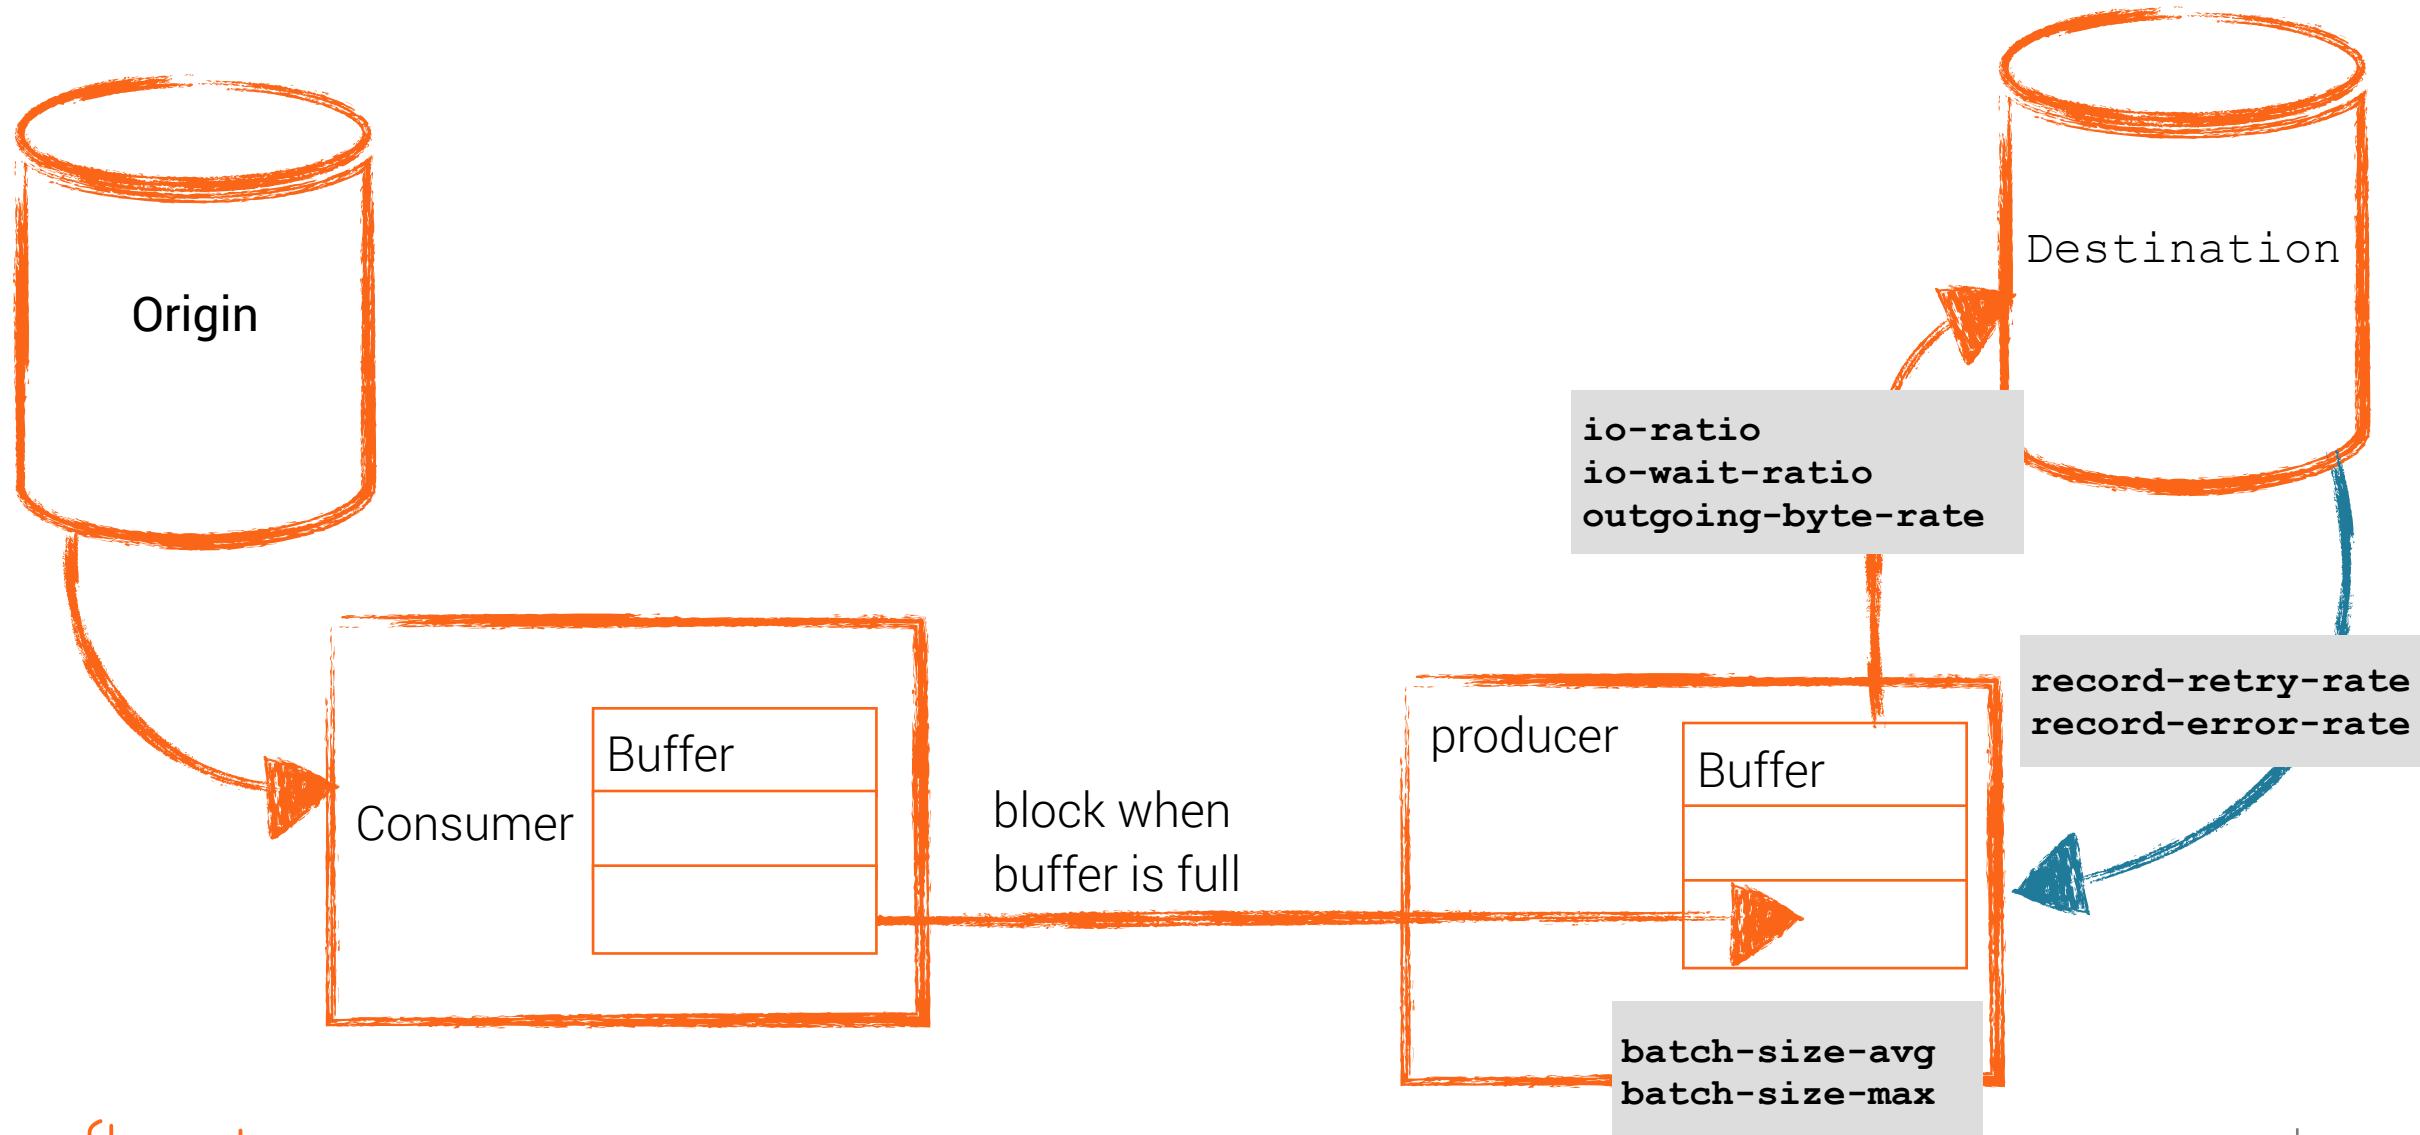
Simple and elegant design



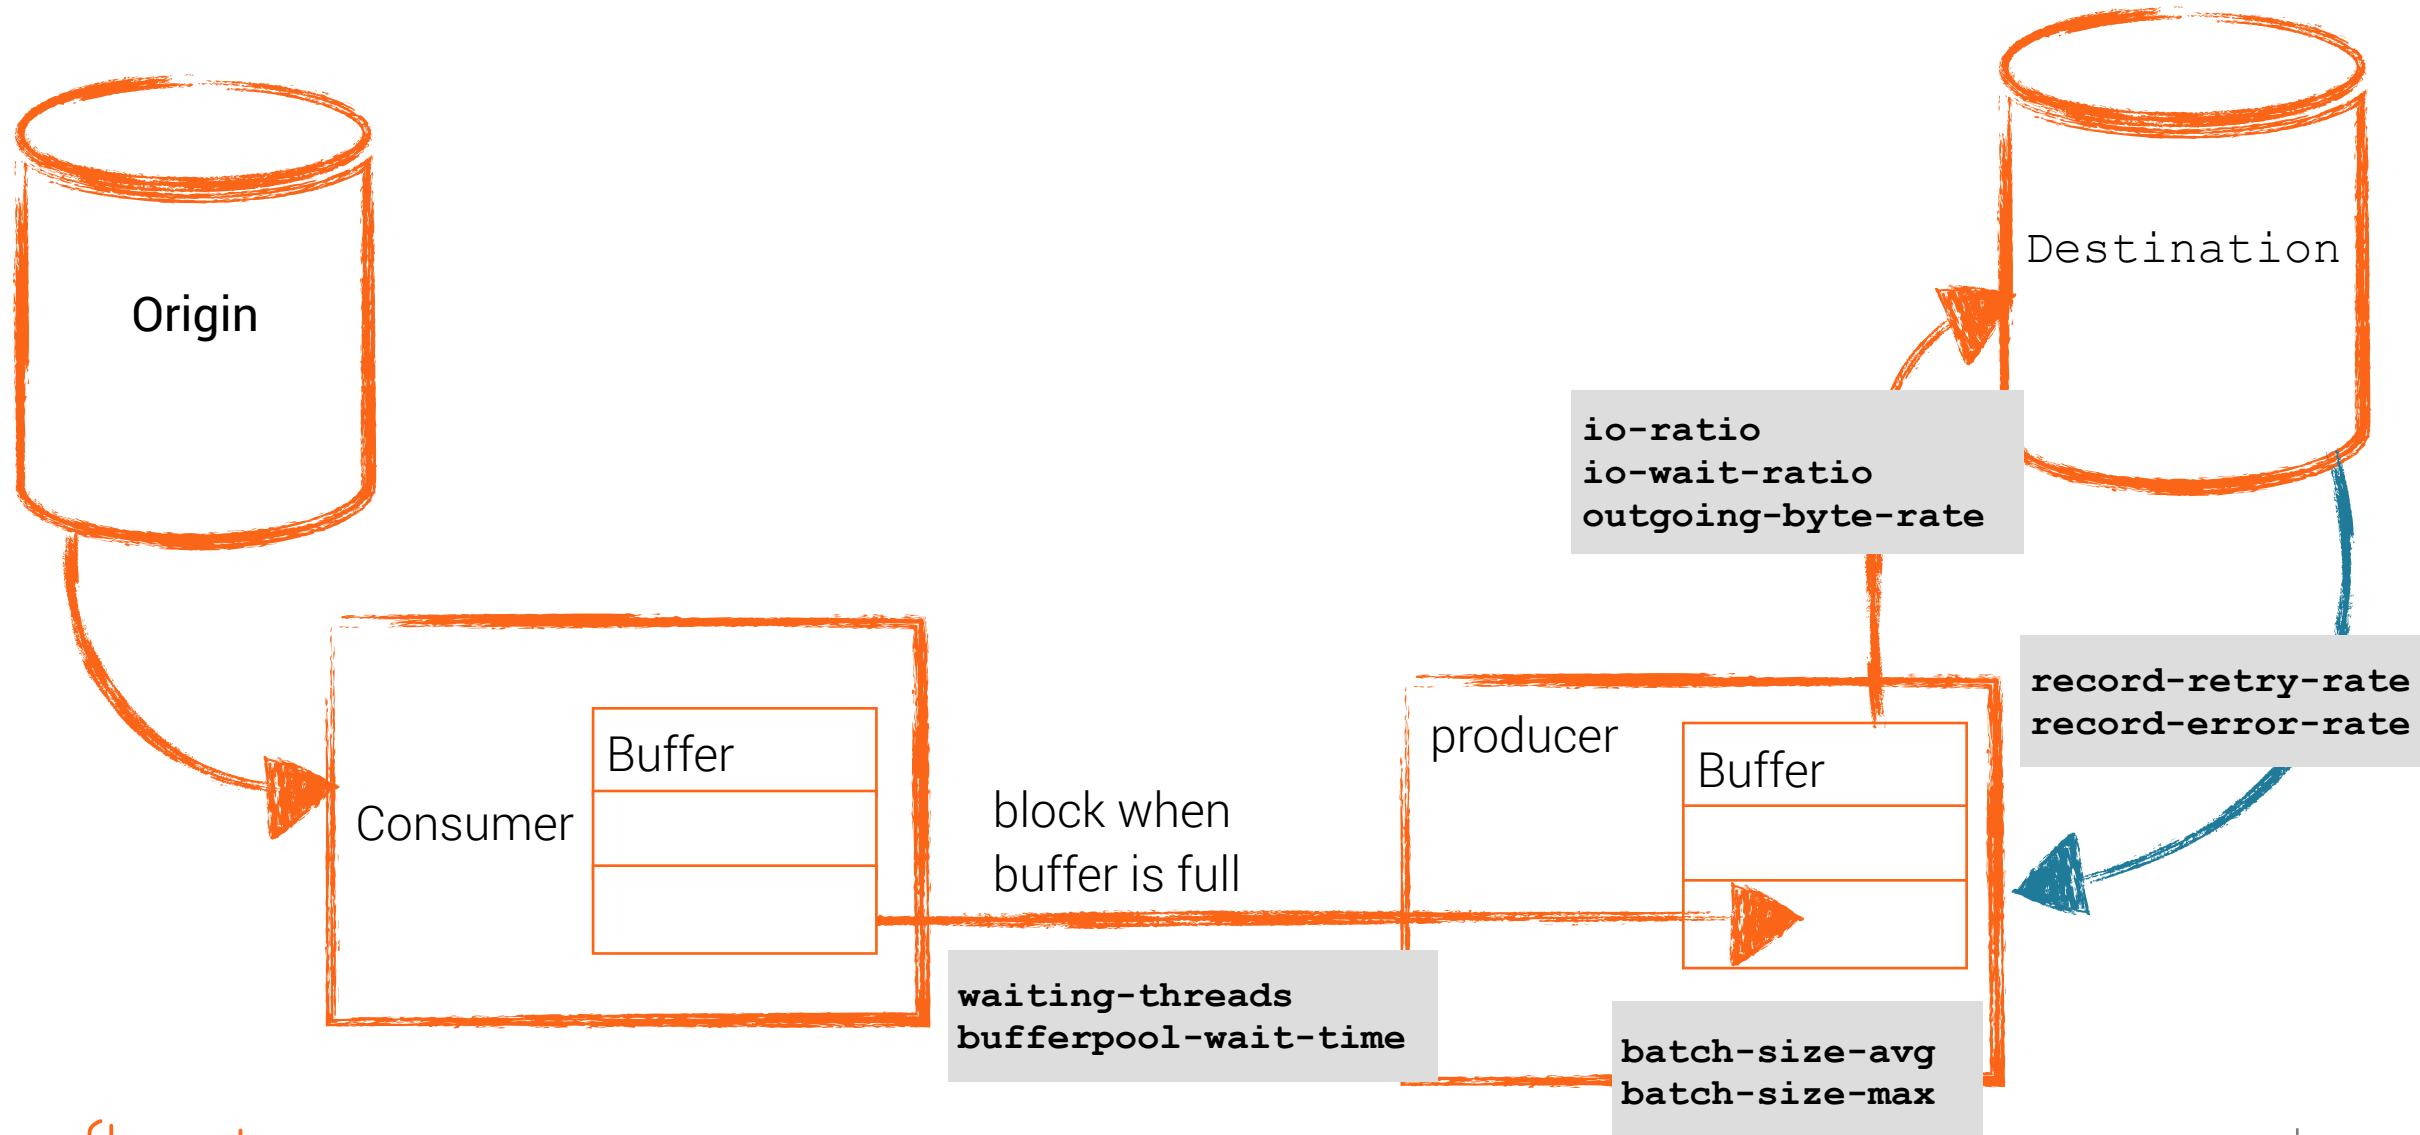
Simple and elegant design



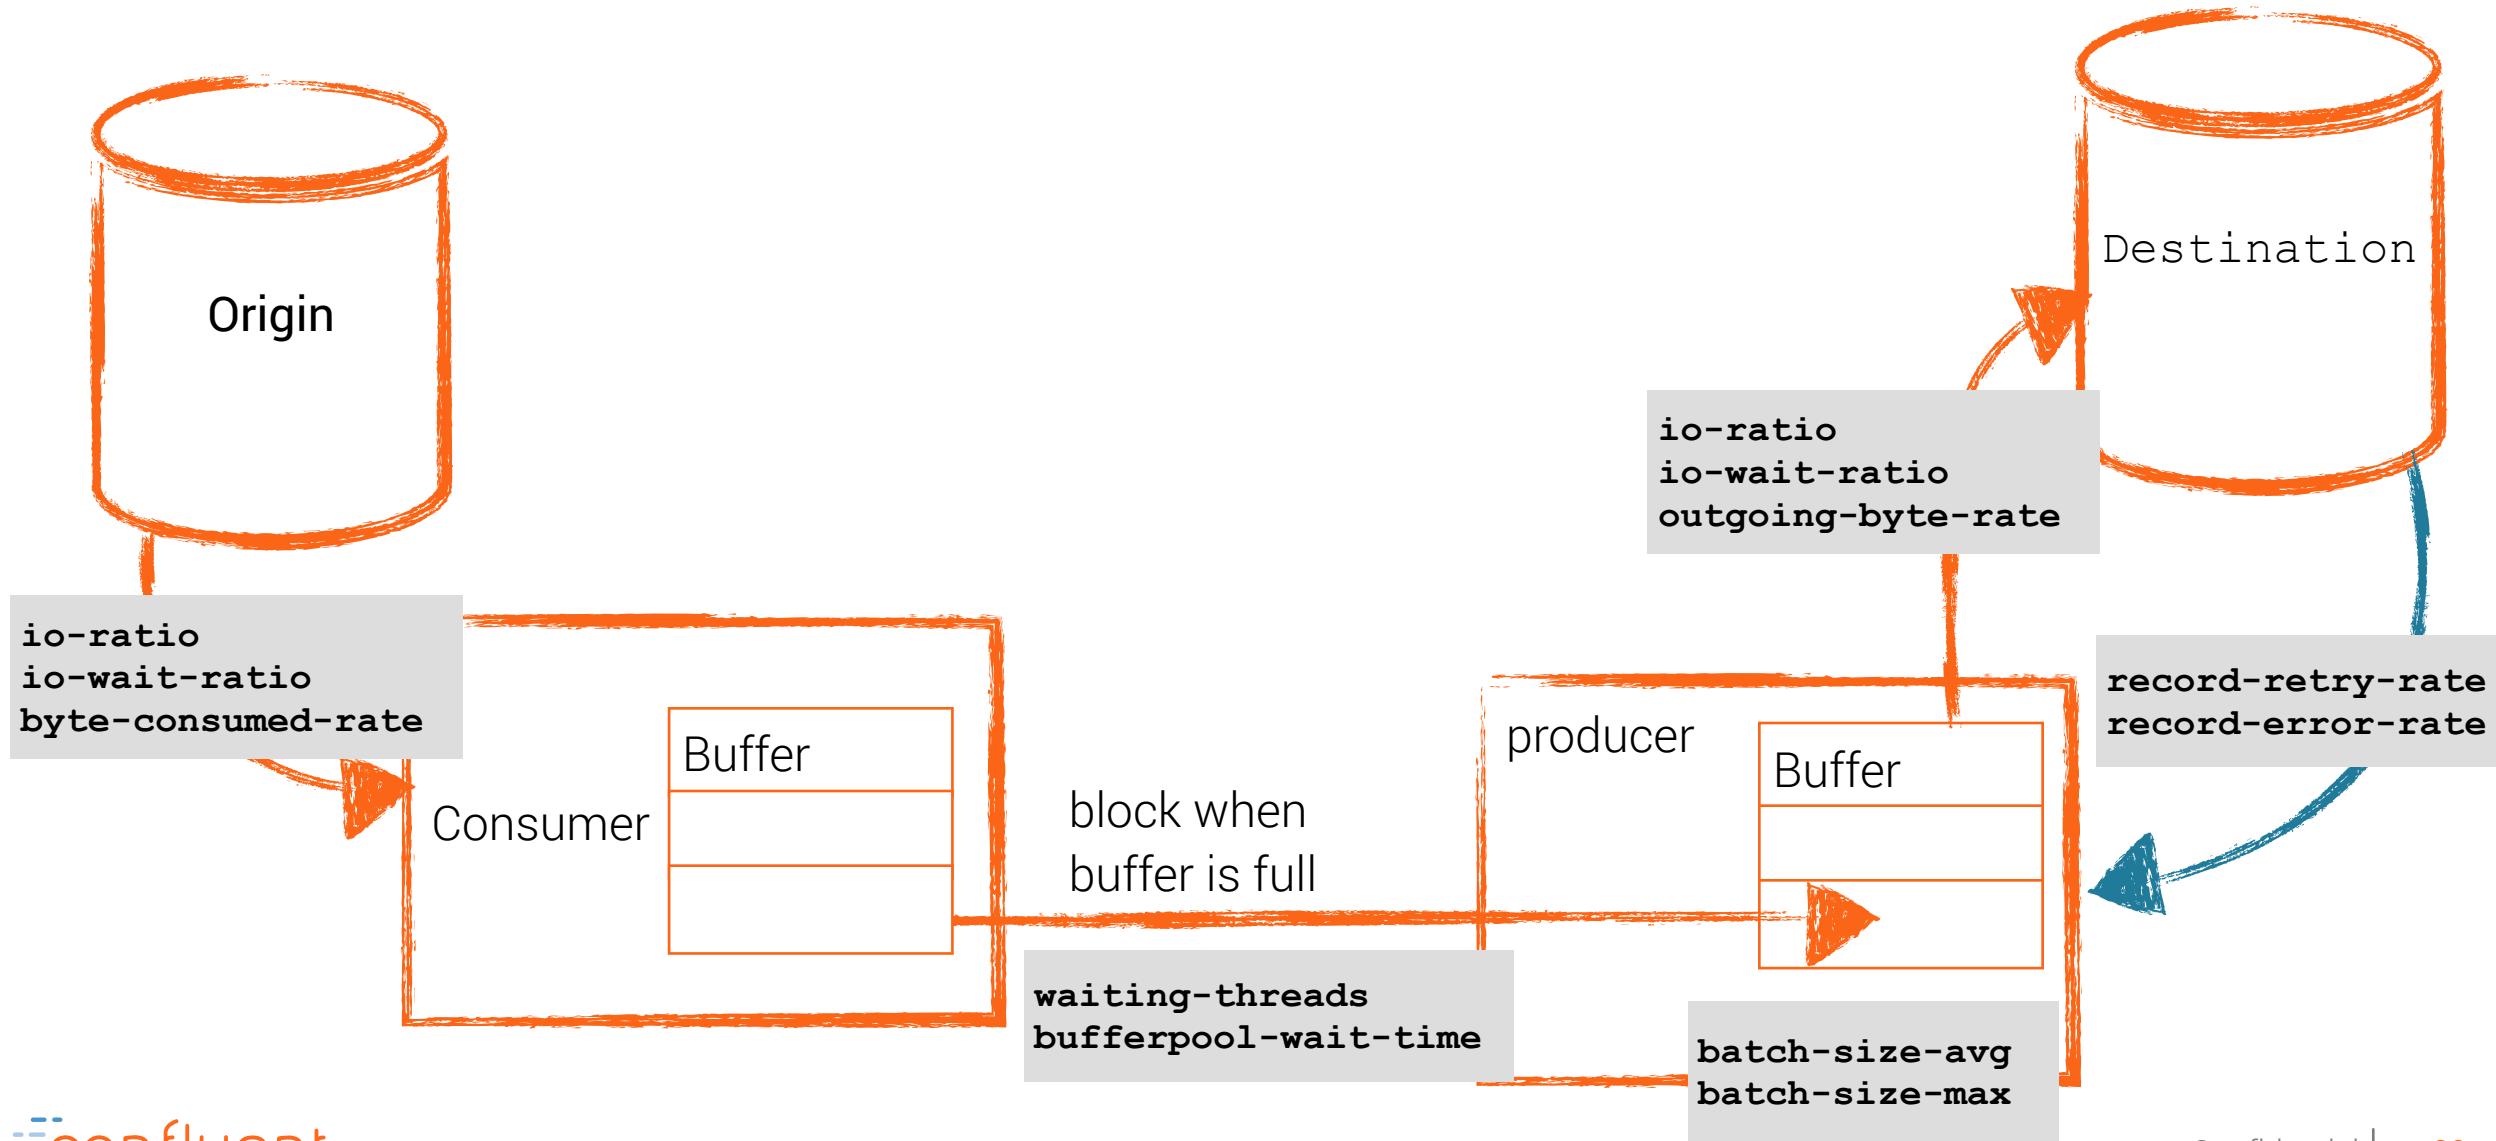
Simple and elegant design



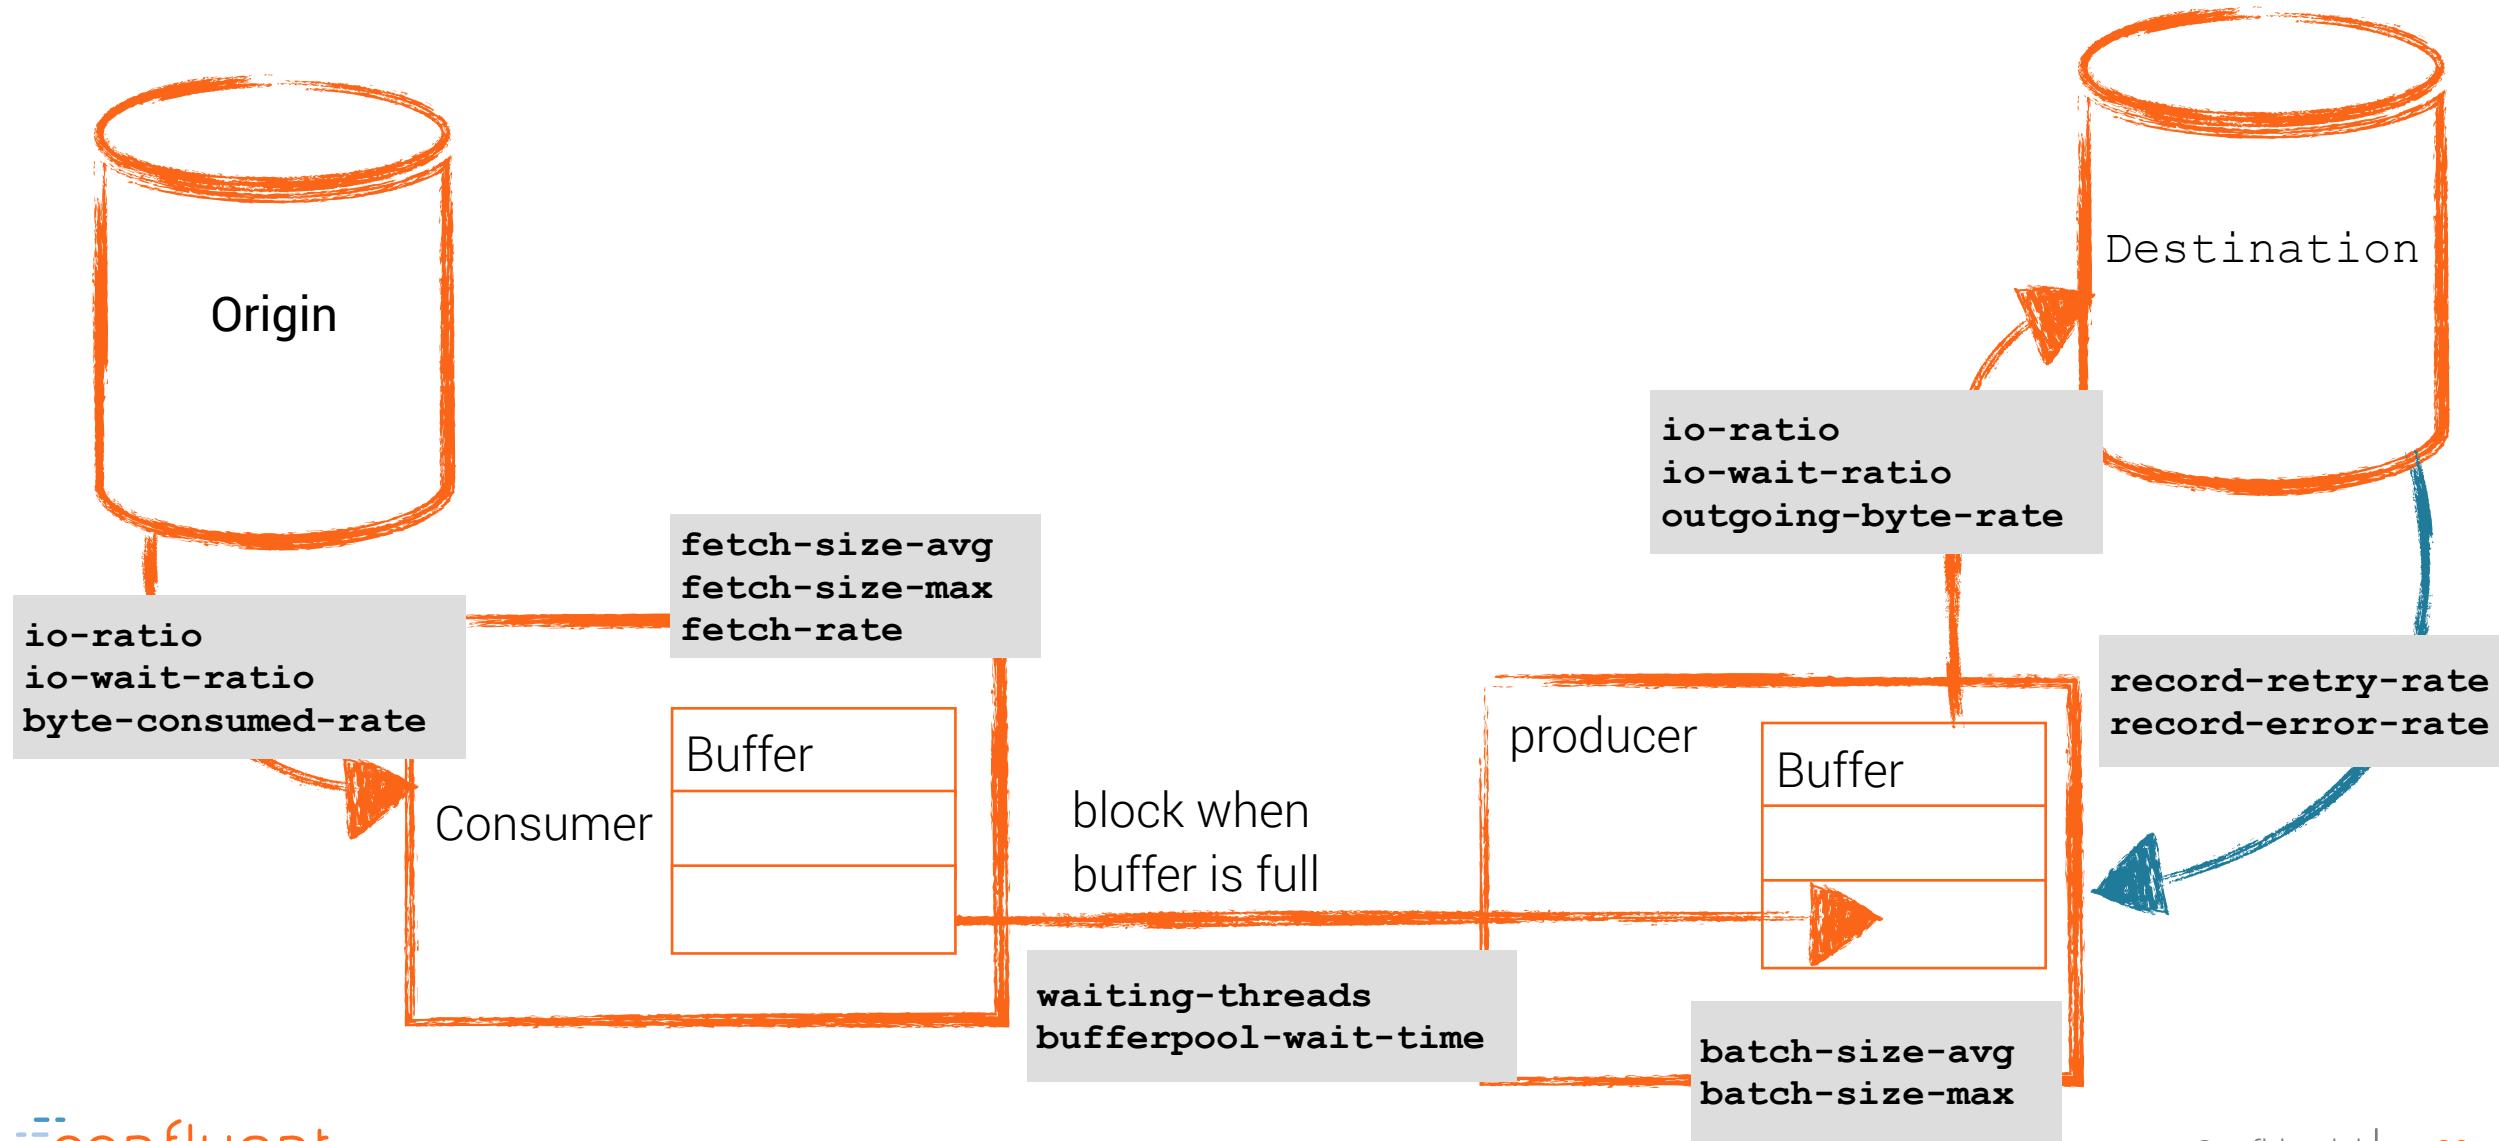
Simple and elegant design



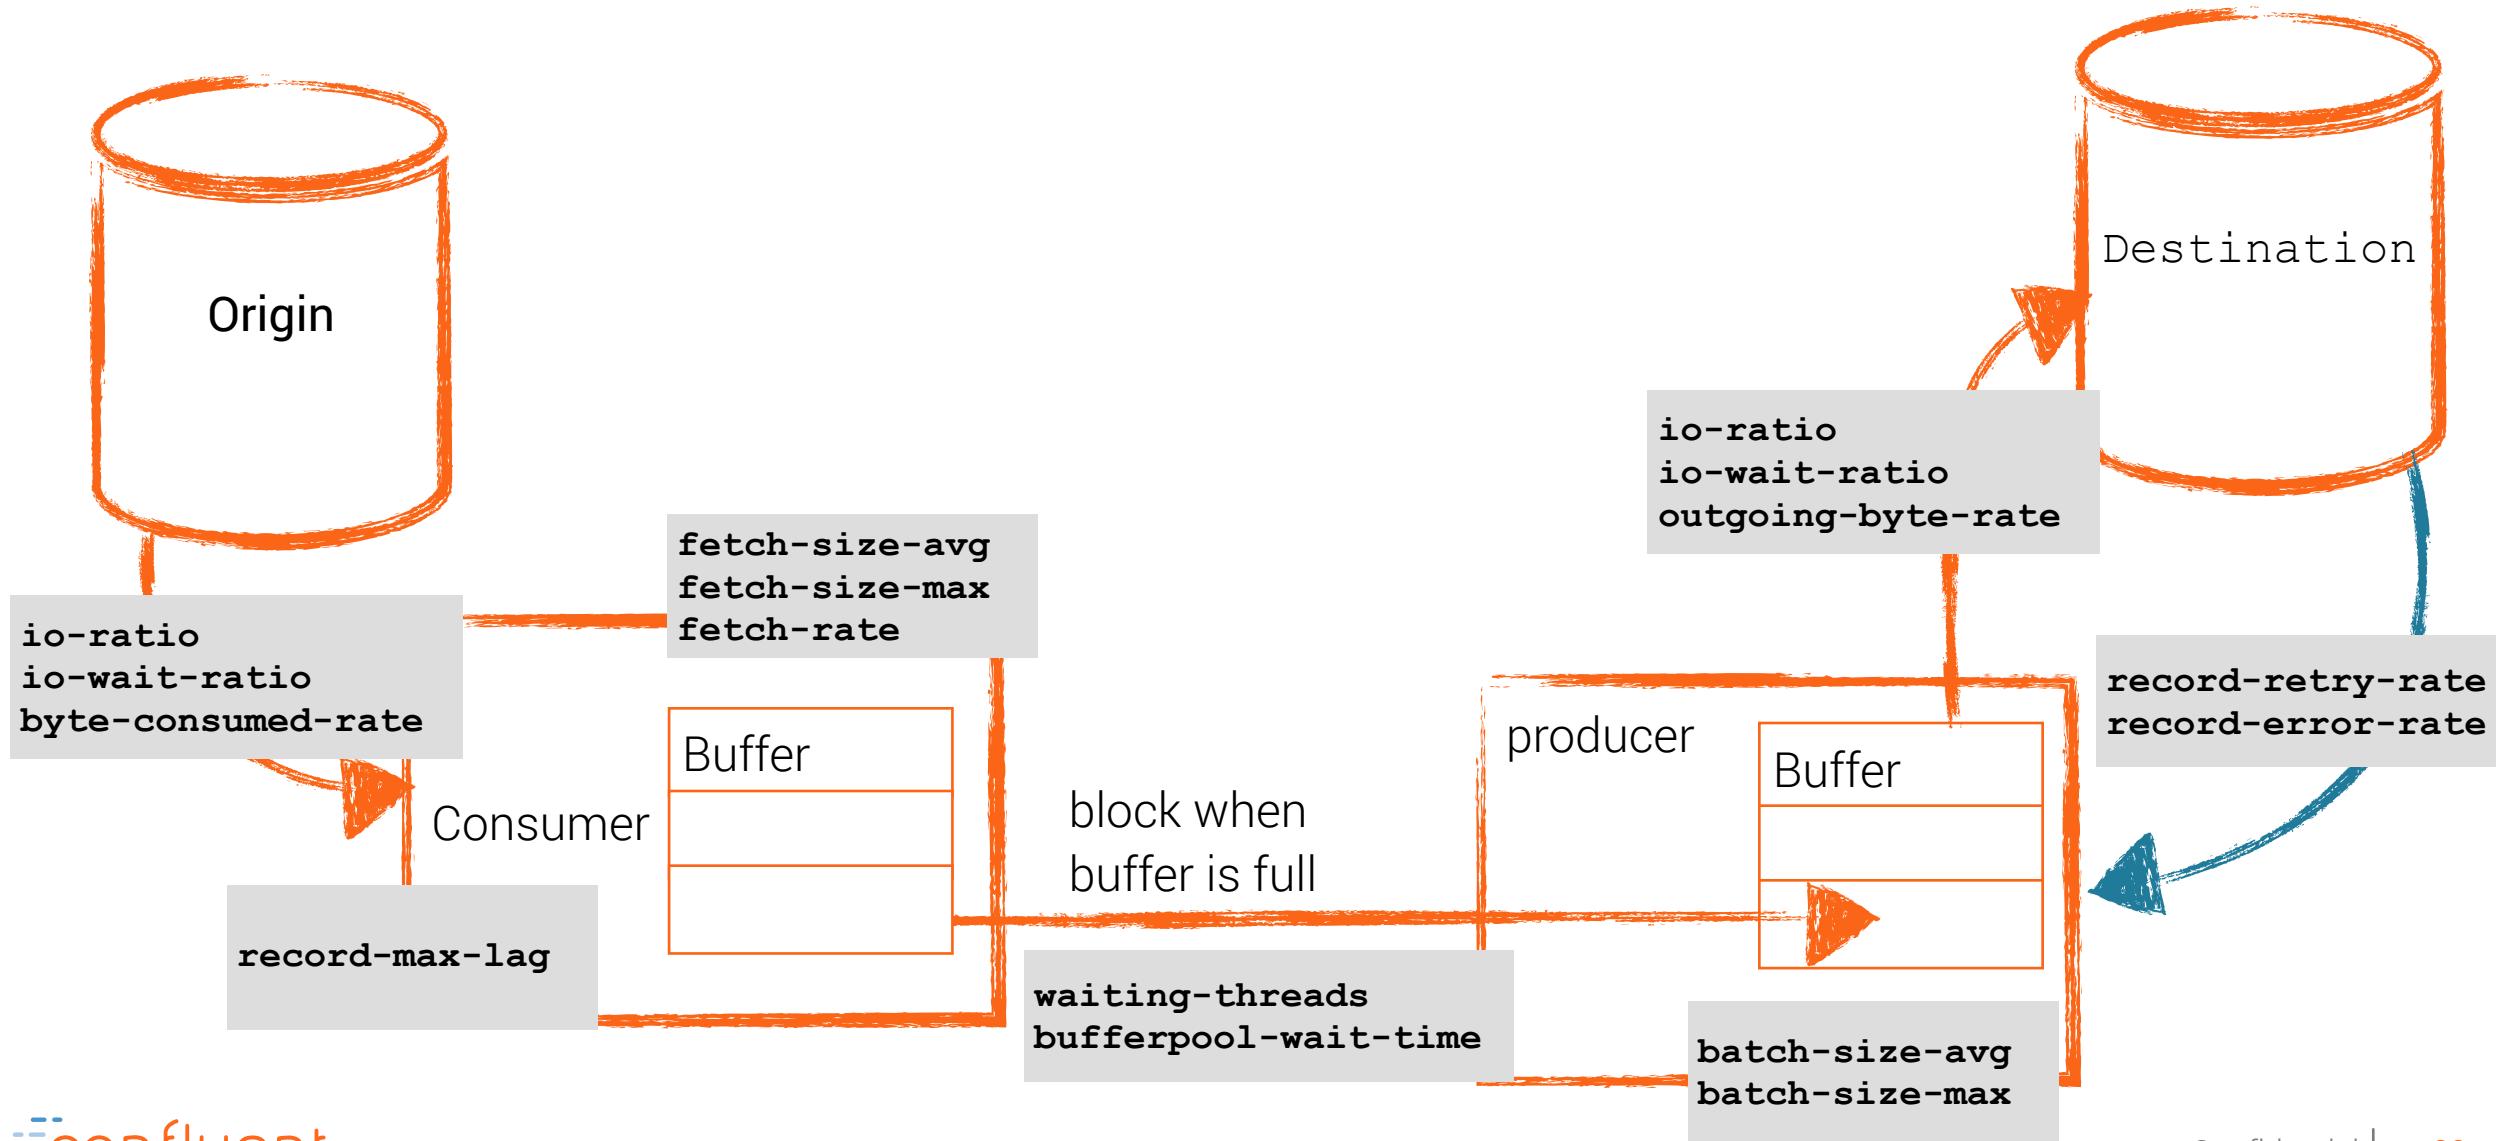
Simple and elegant design



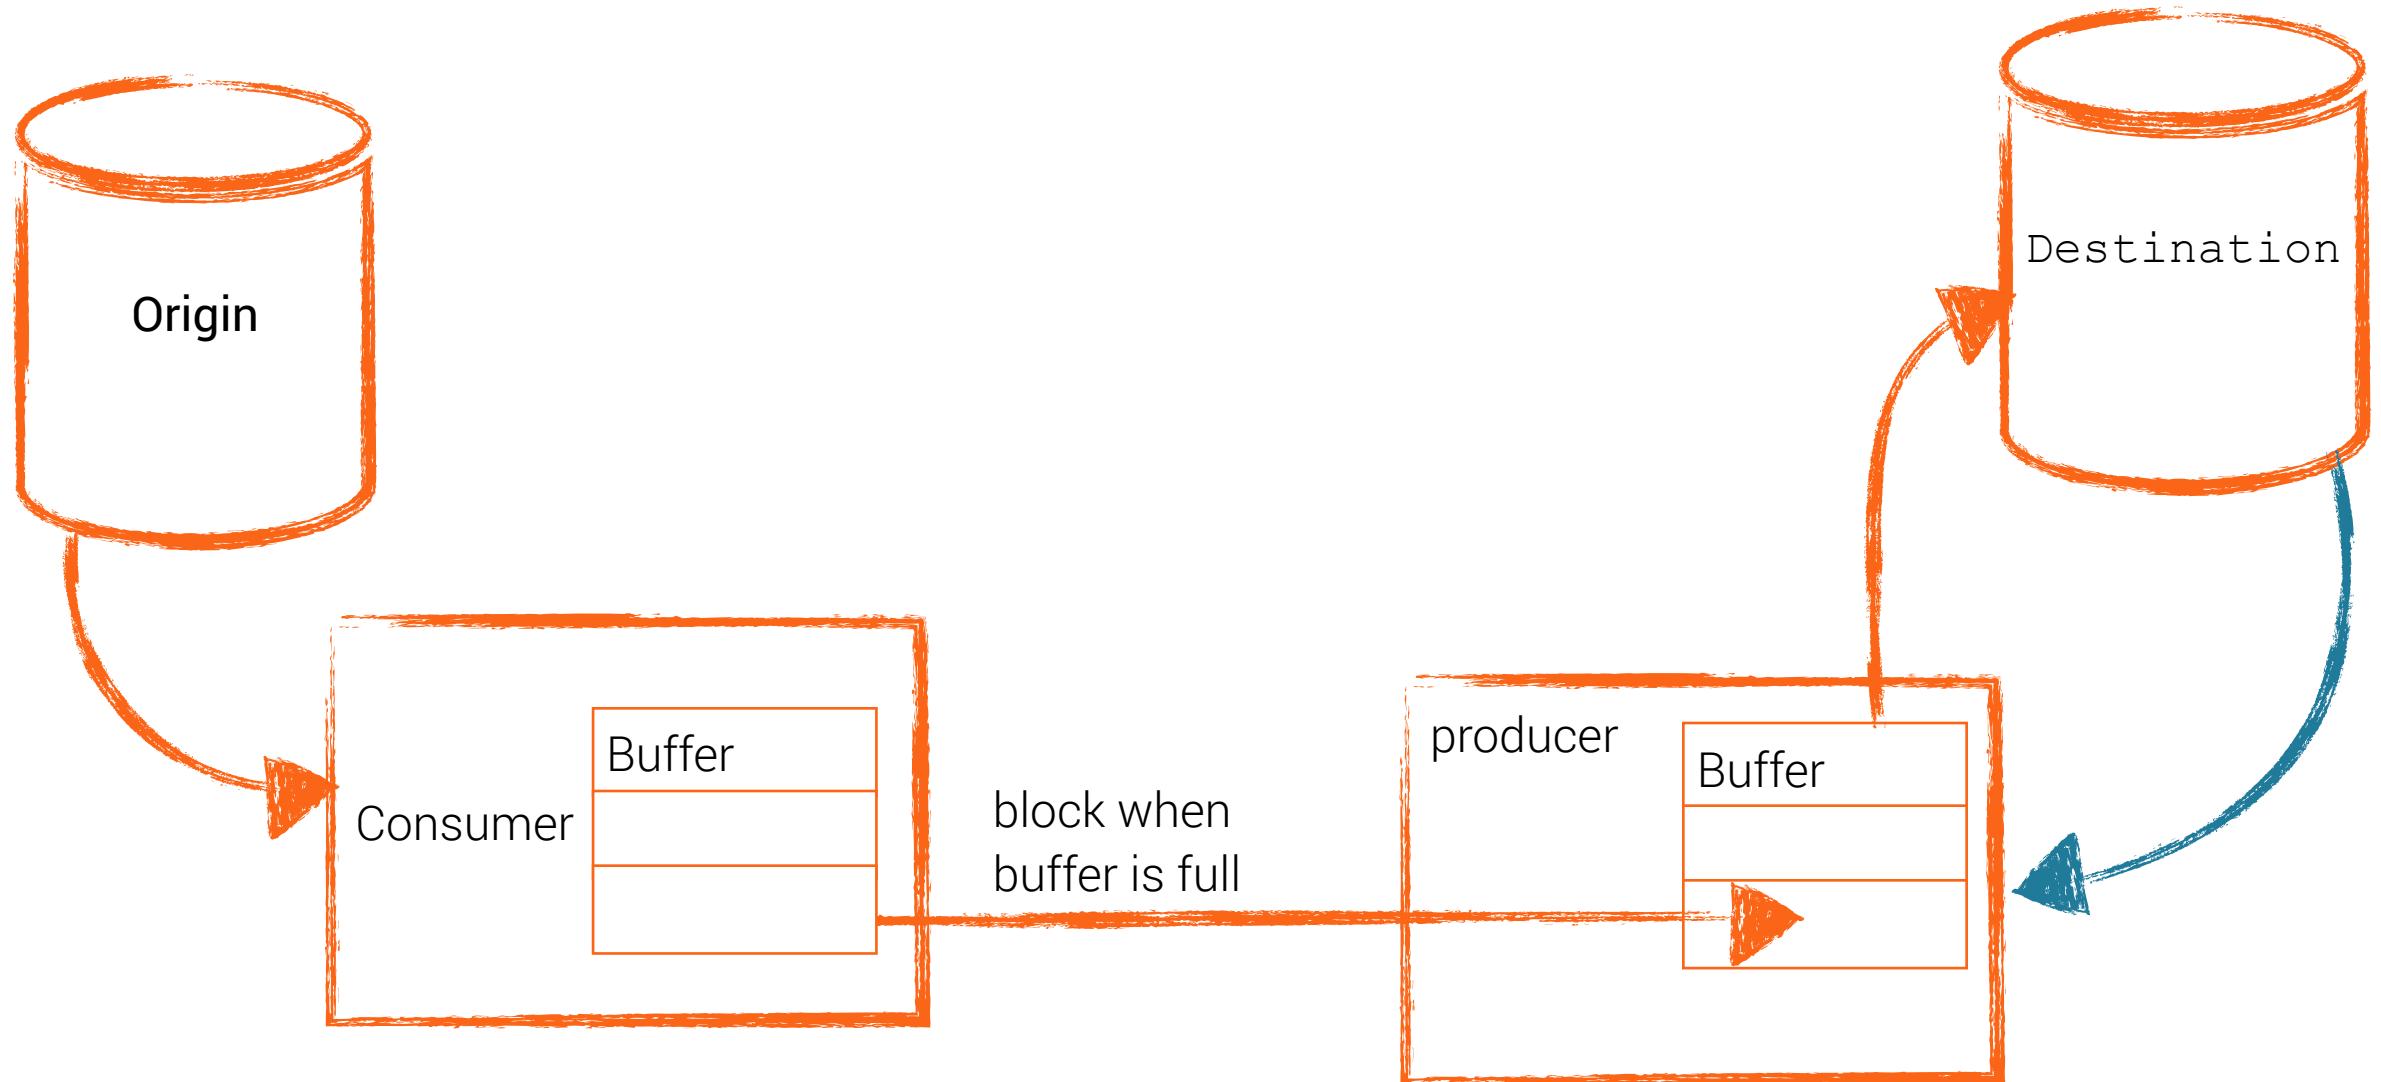
Simple and elegant design



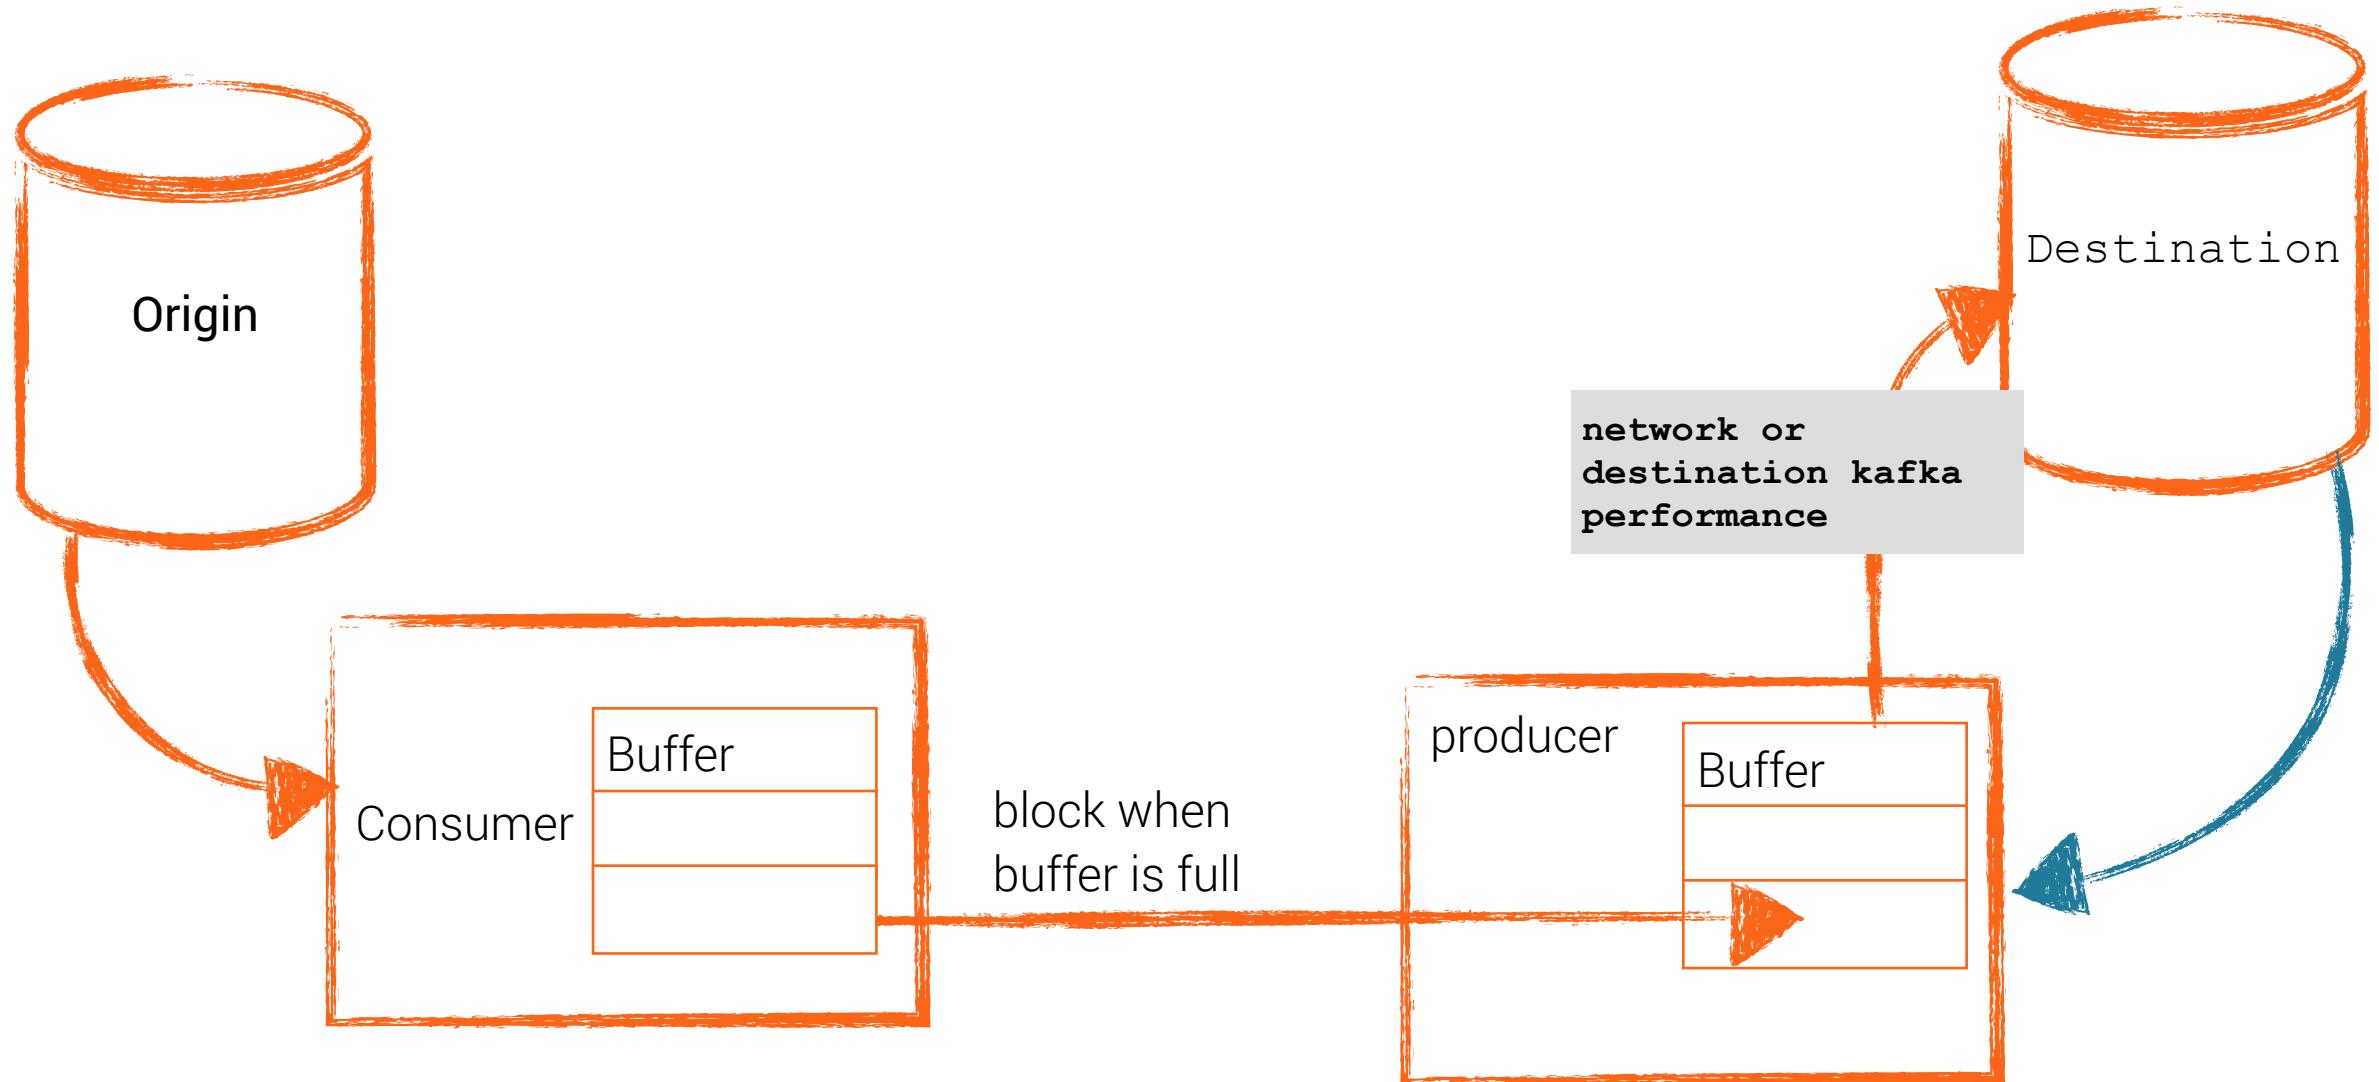
Simple and elegant design



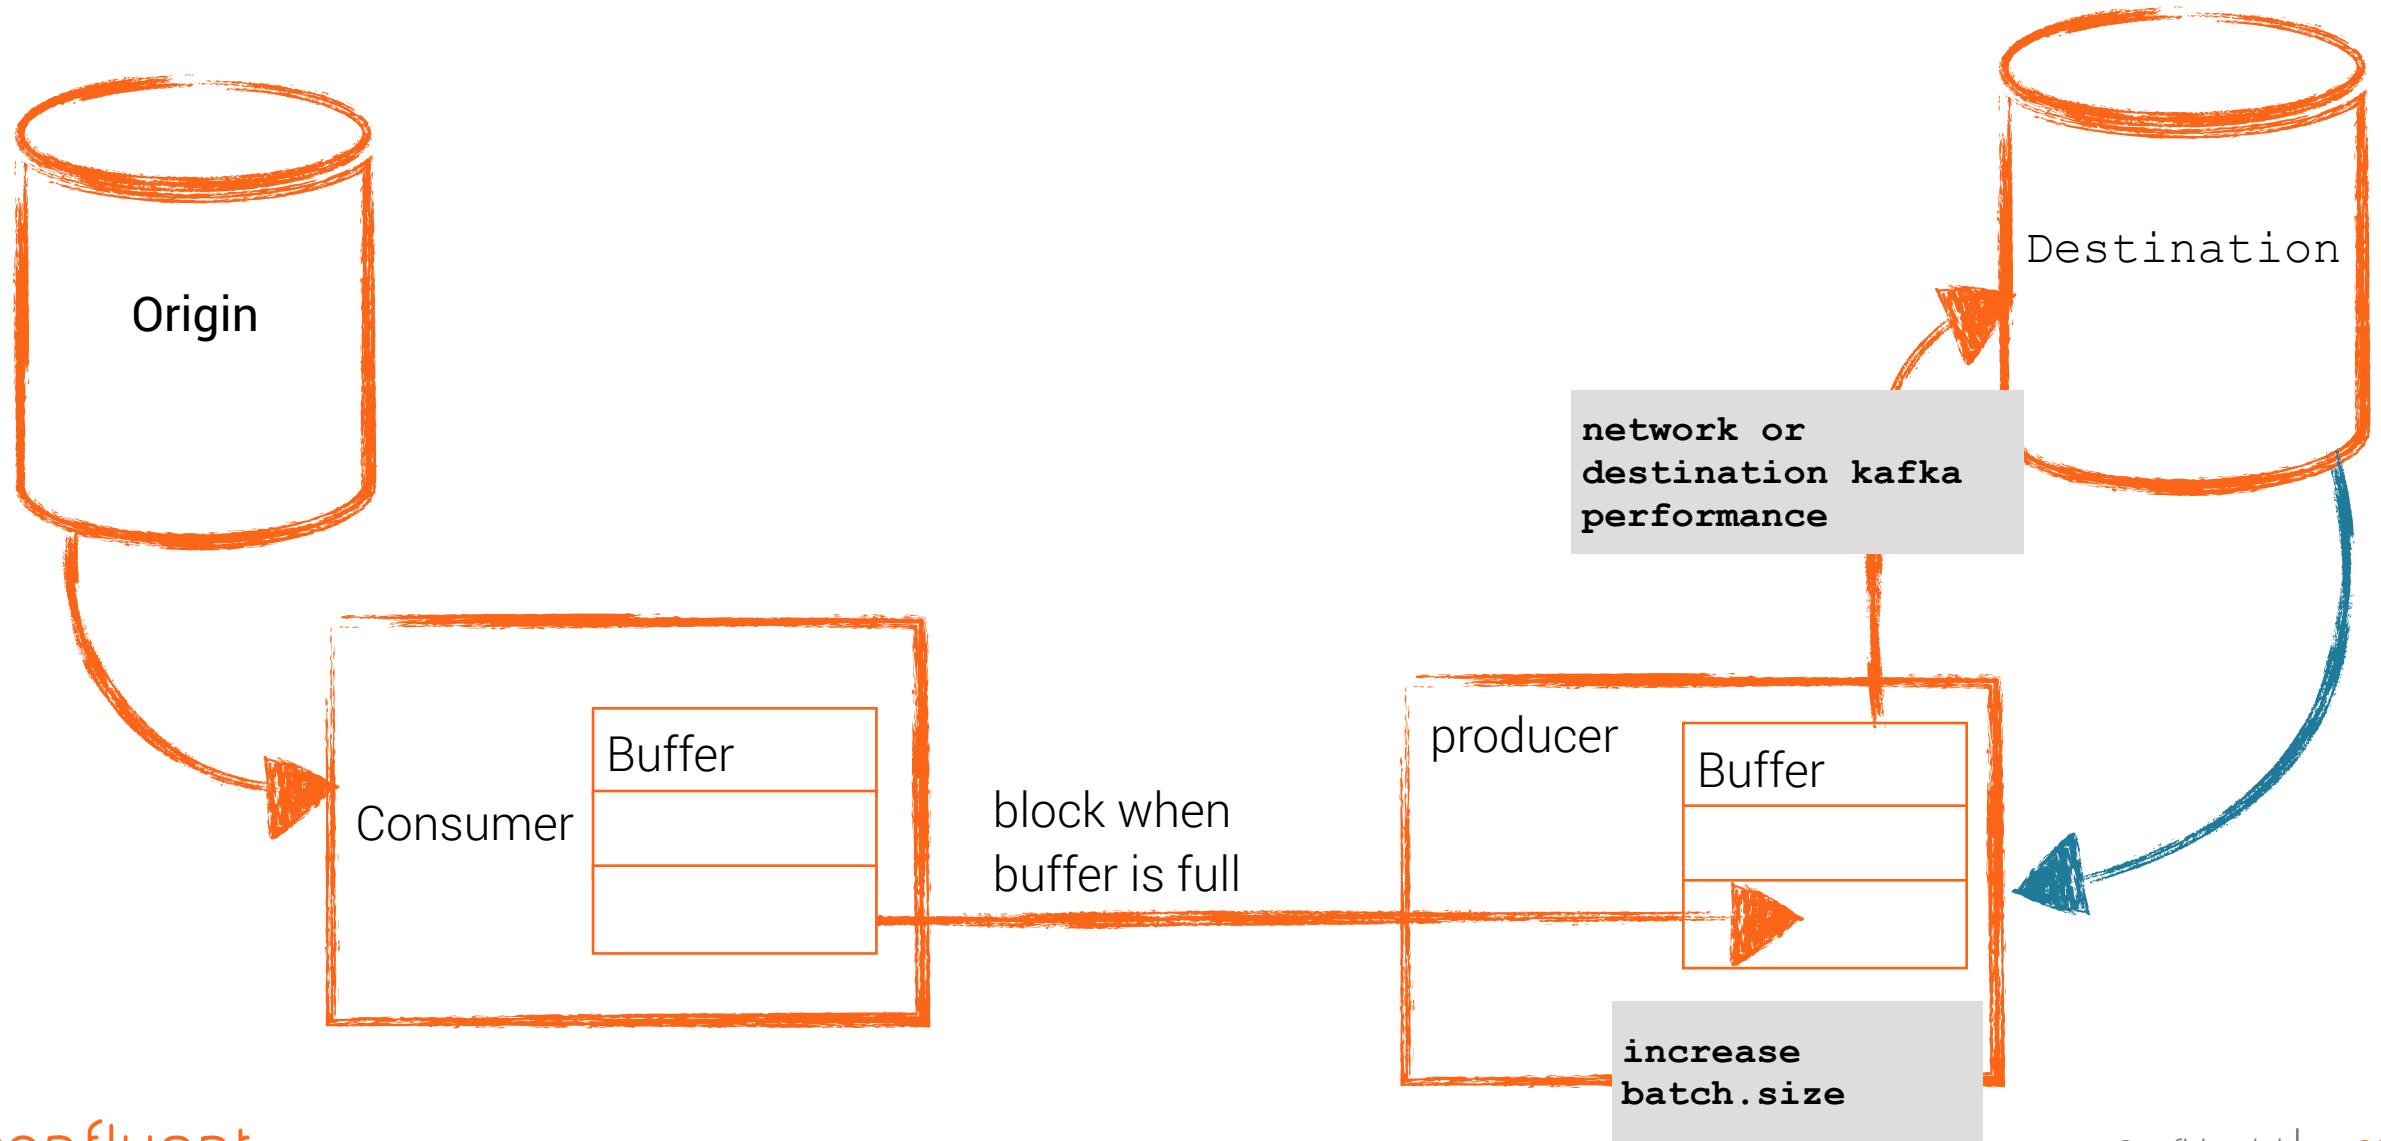
Simple and elegant design



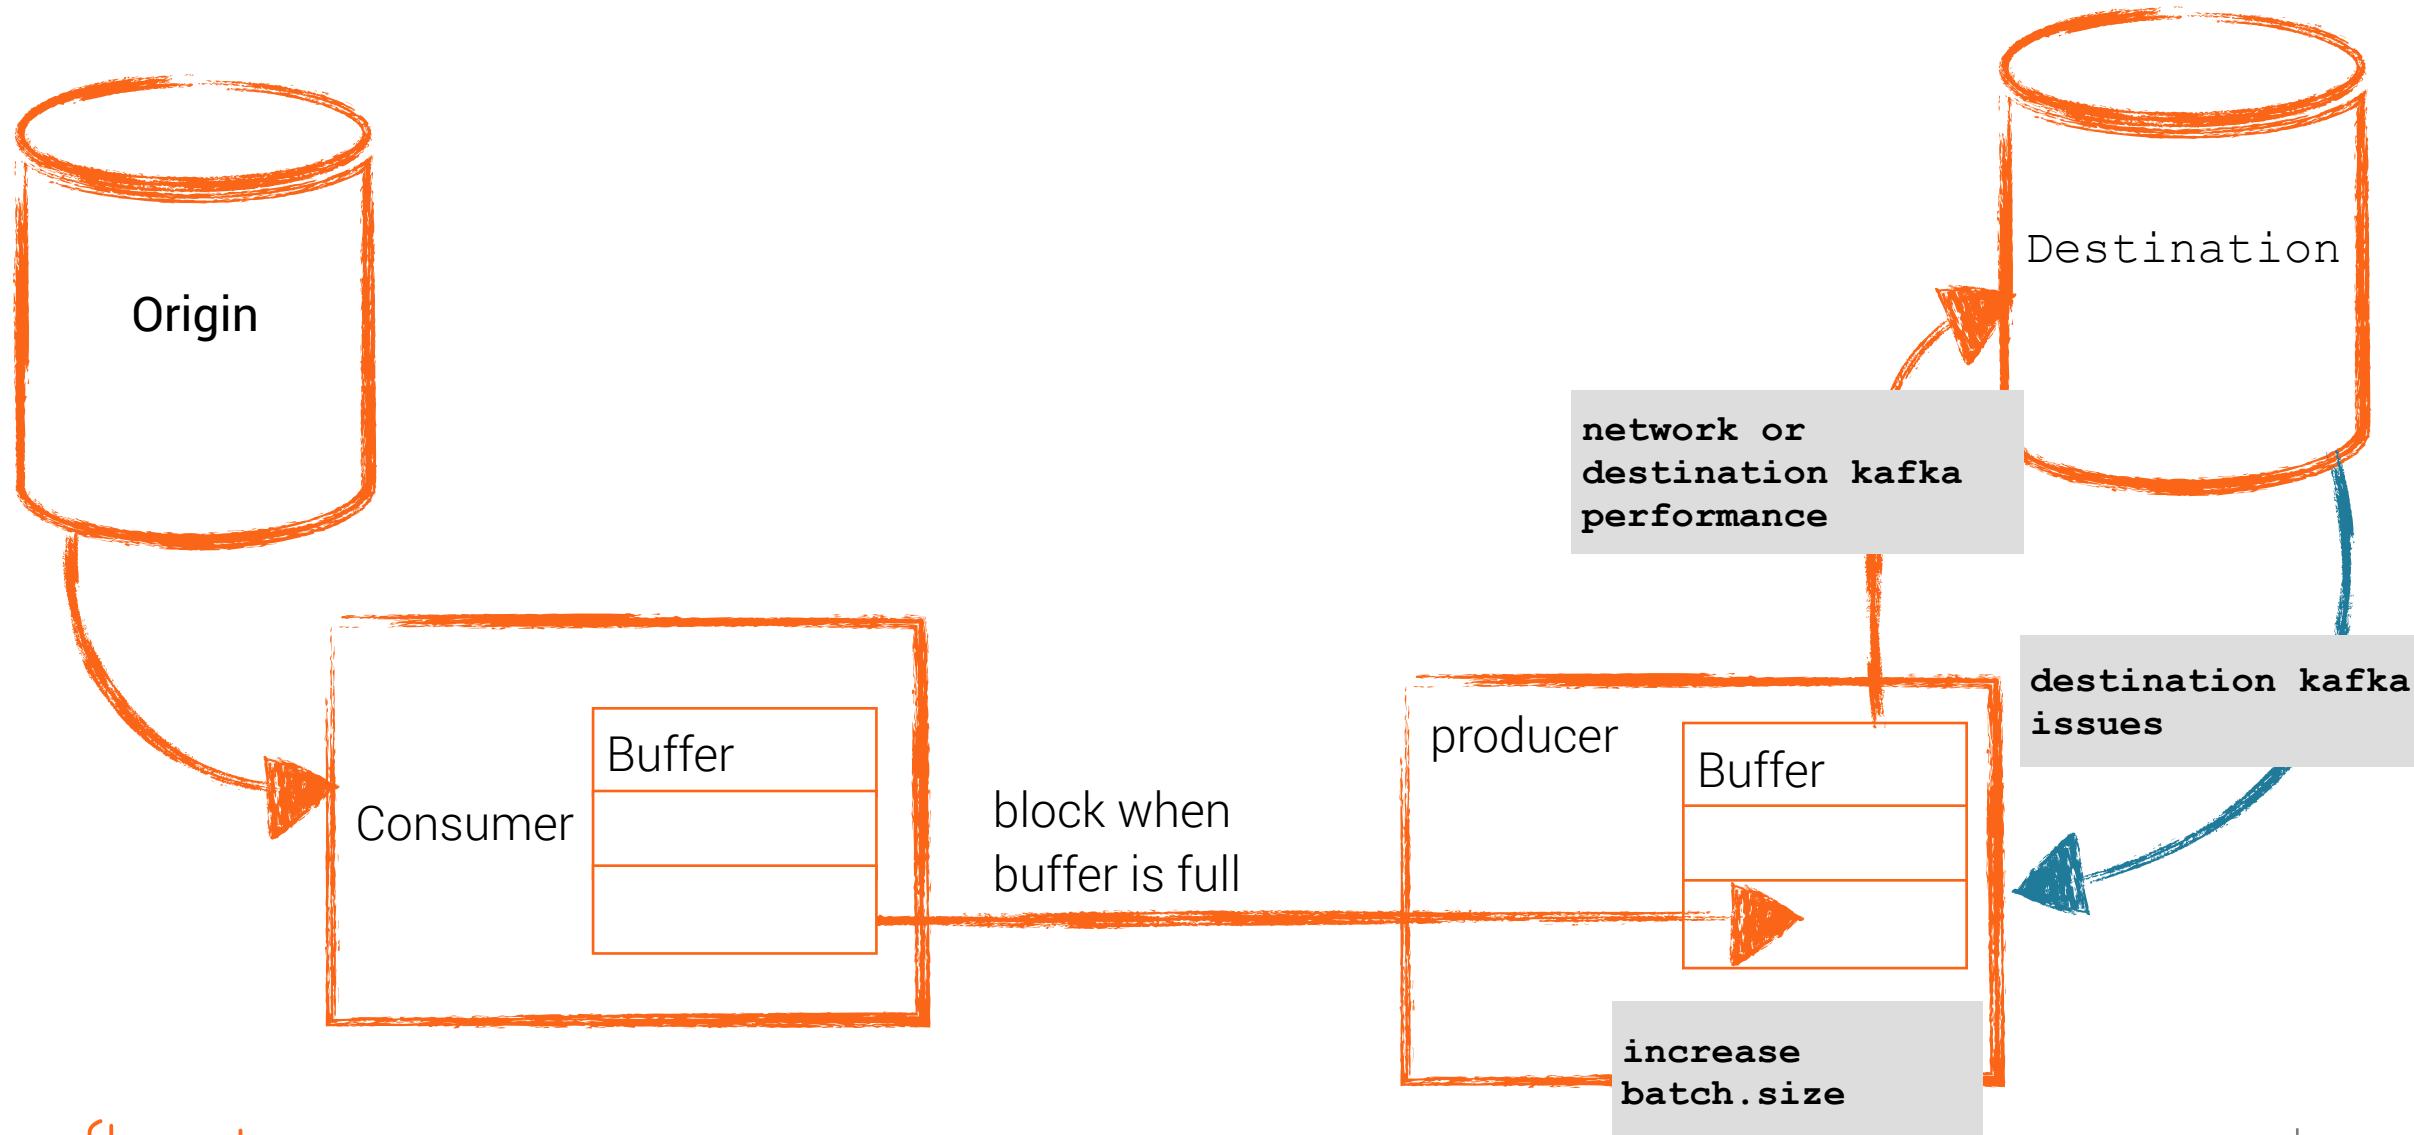
Simple and elegant design



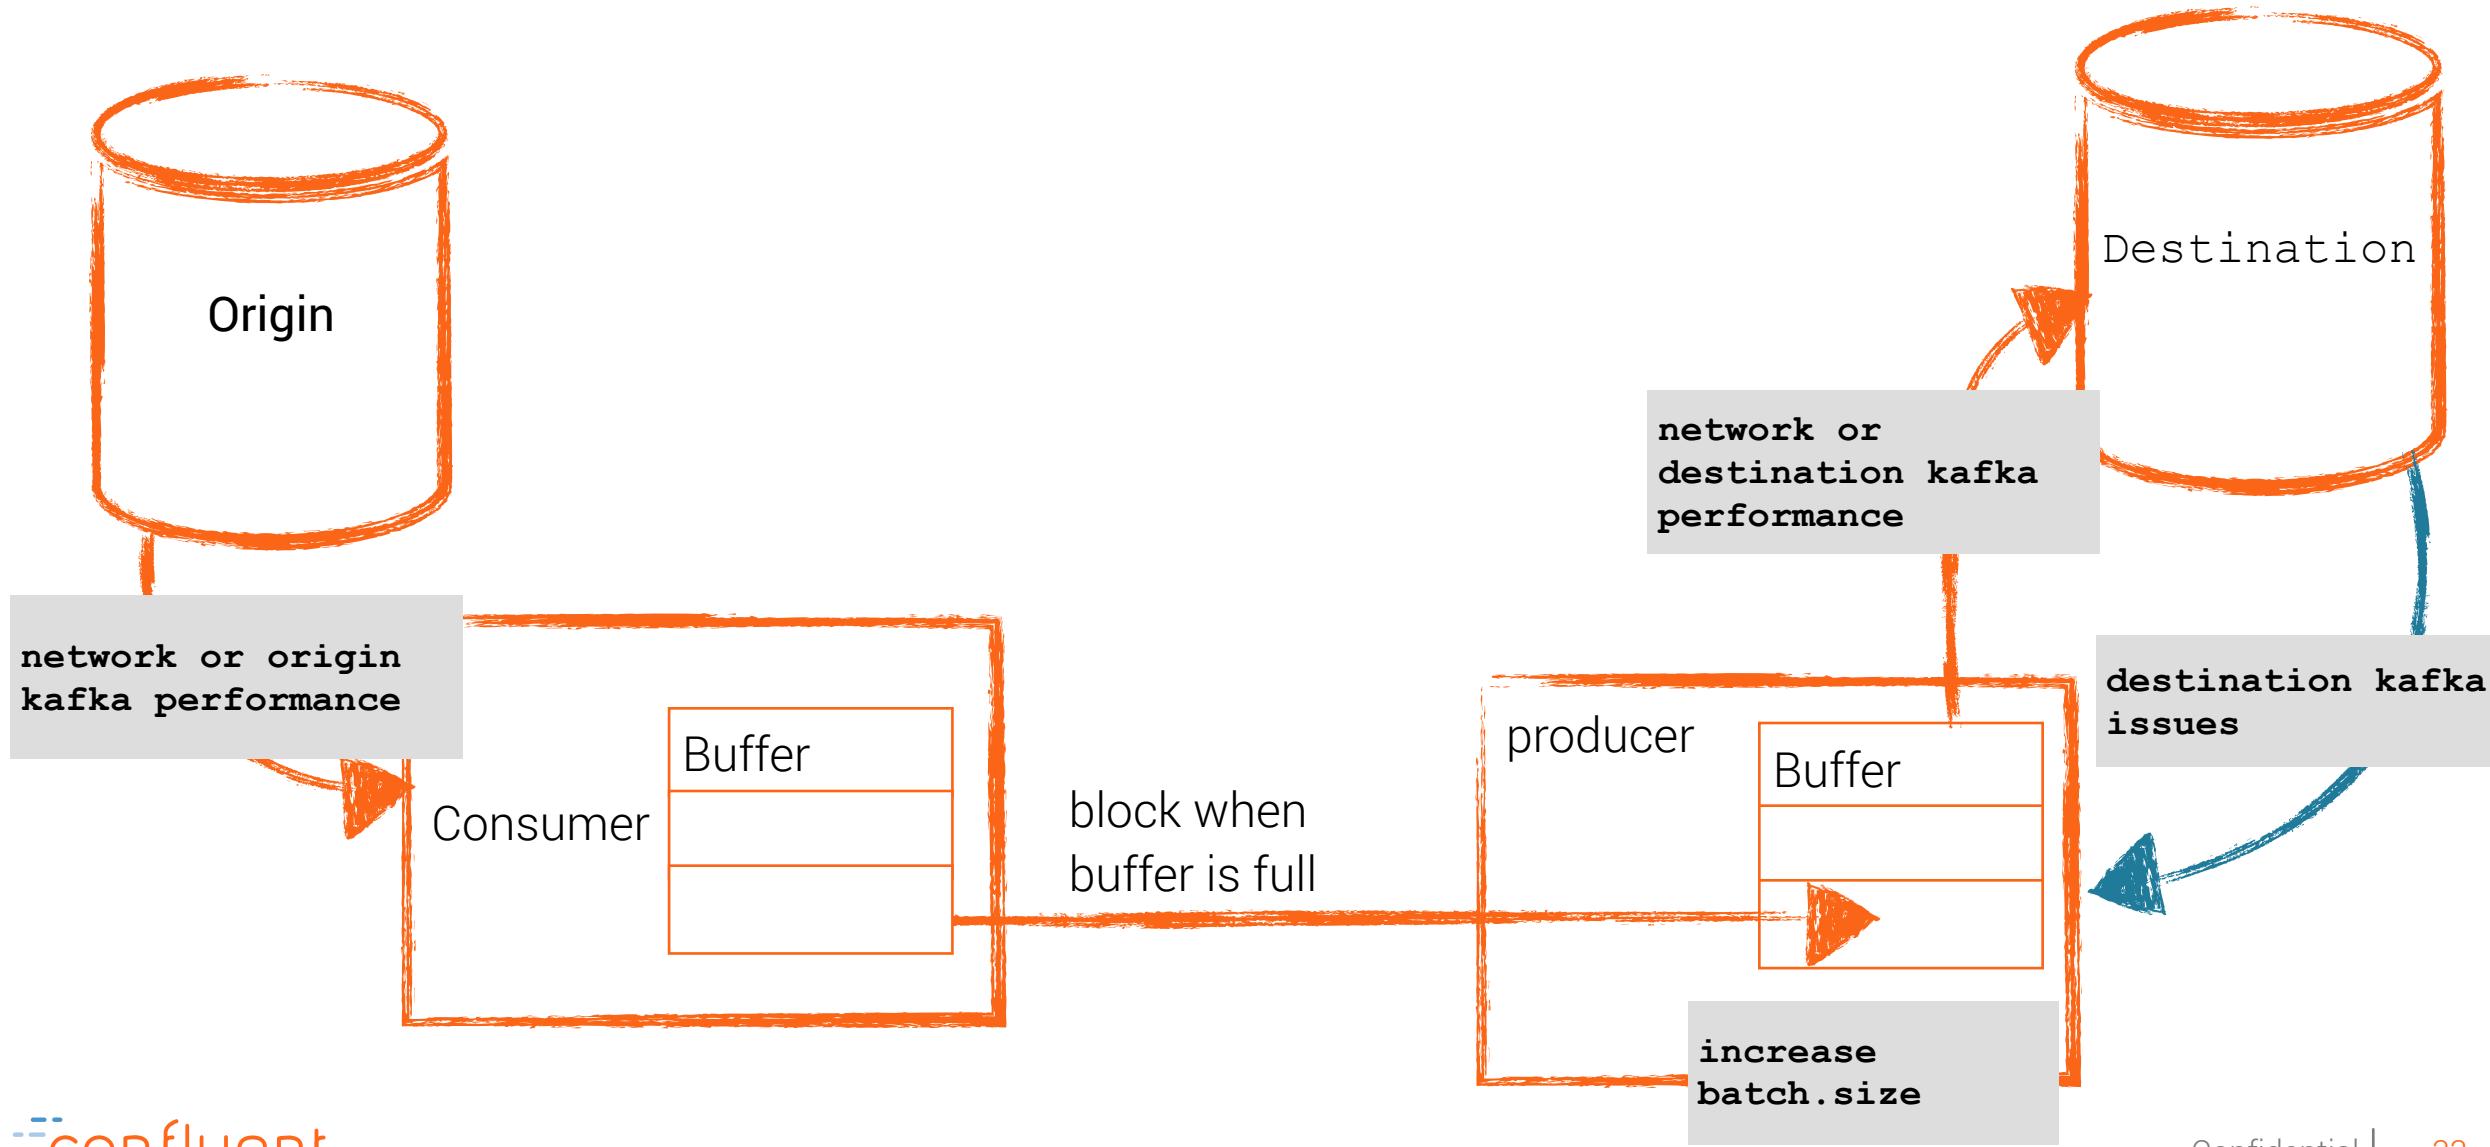
Simple and elegant design



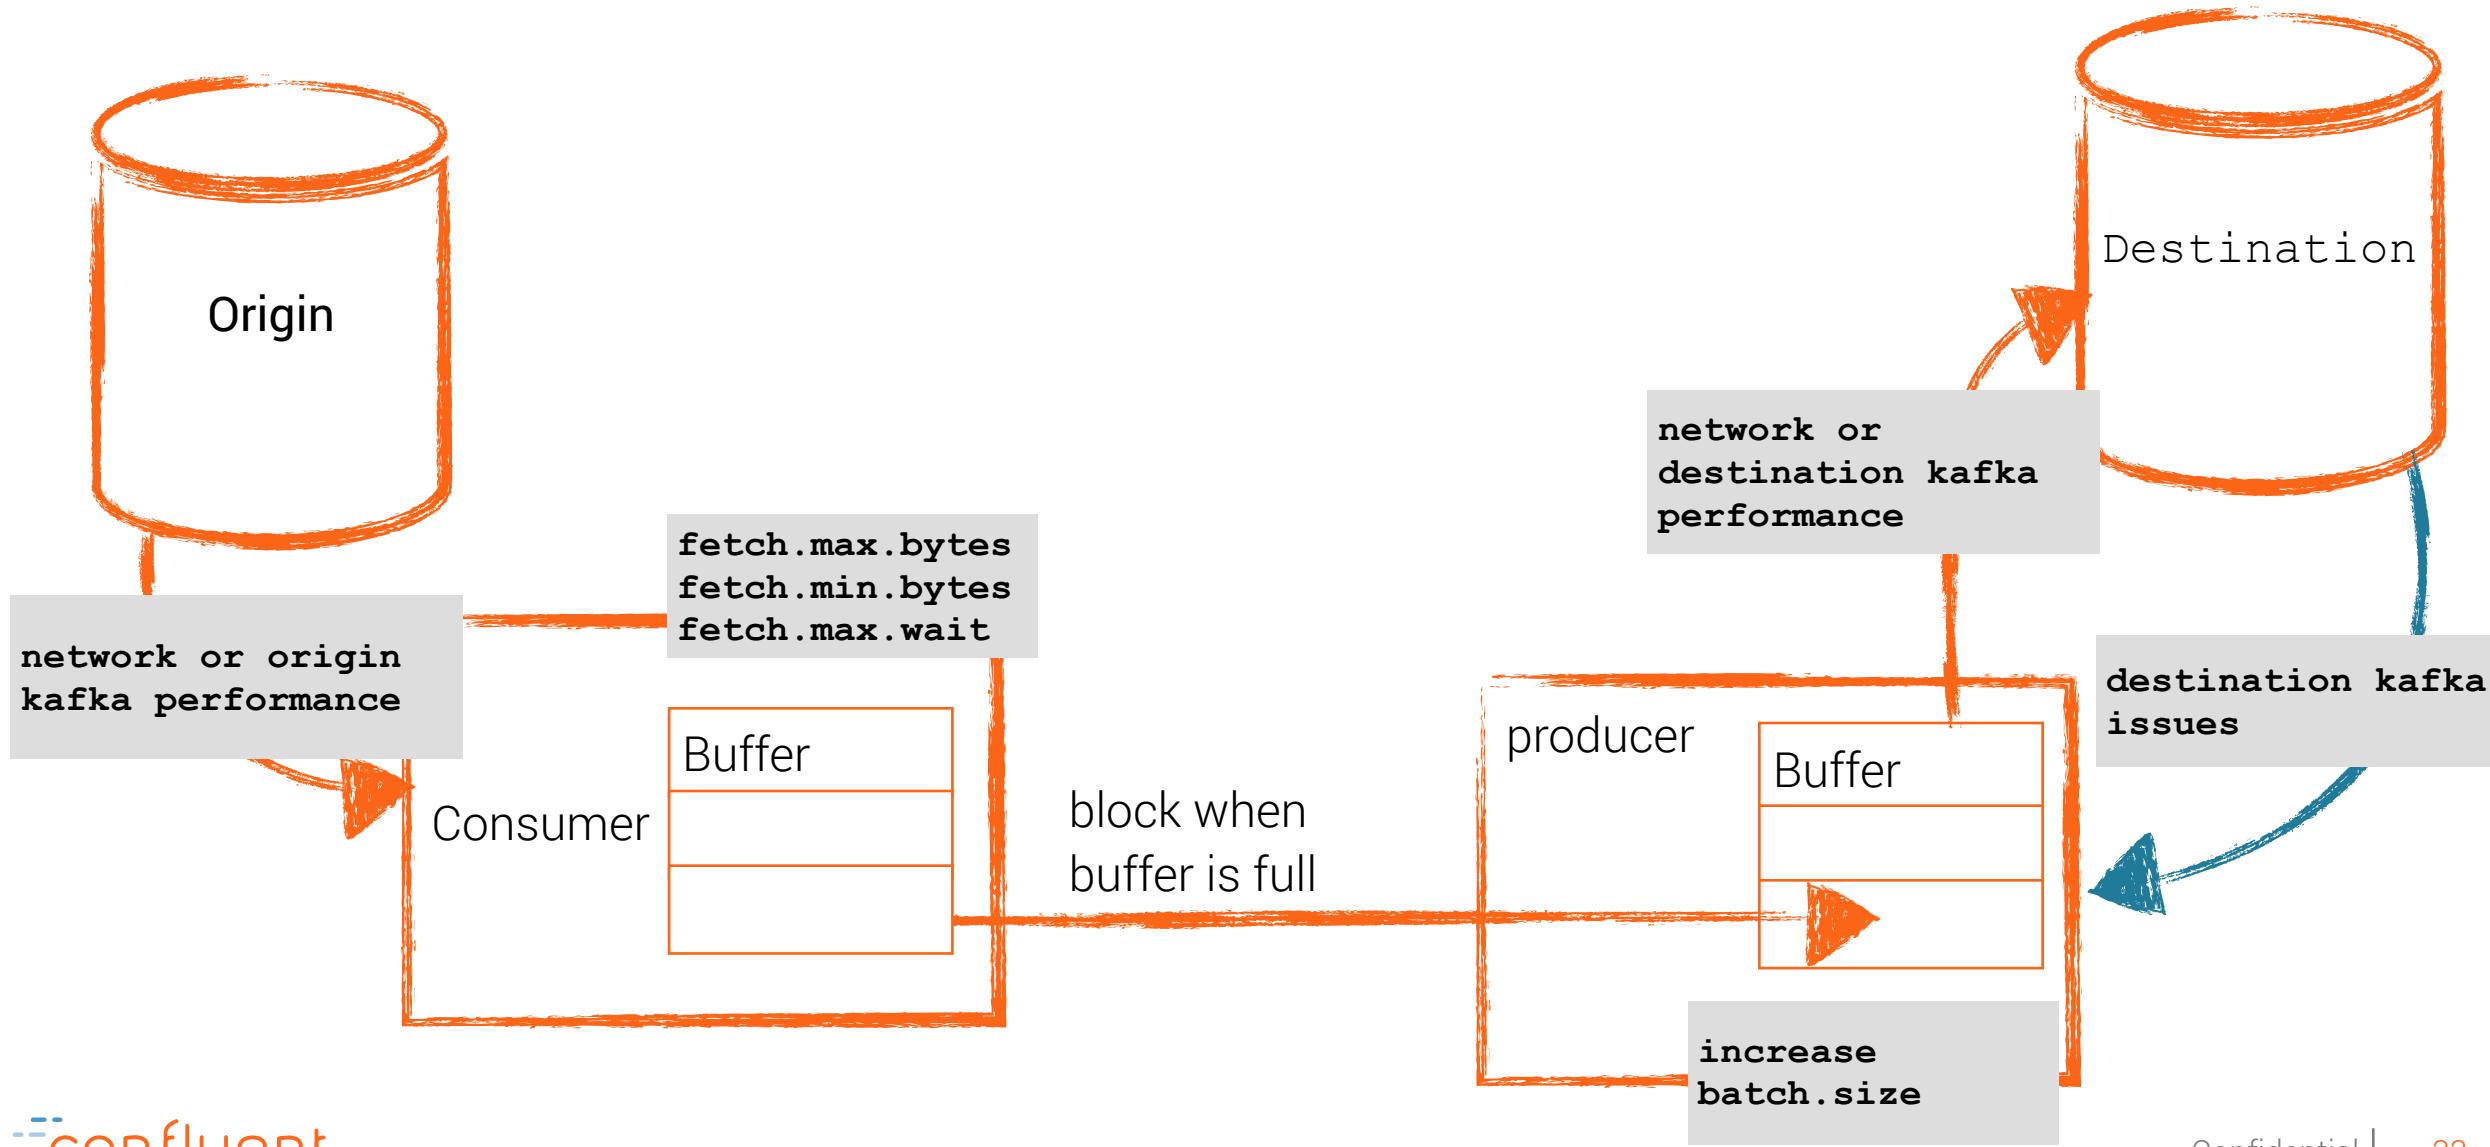
Simple and elegant design



Simple and elegant design



Simple and elegant design



Network Tuning

- WAN has high latency. We deal with it.
- Compute buffer size to match: https://www.switch.ch/network/tools/tcp_throughput/
- send.buffer.bytes and receive.buffer.bytes on producer, consumer, brokers
- OS tuning: https://wwwx.cs.unc.edu/~sparkst/howto/network_tuning.php
net.core.rmem_default, net.core.rmem_max, net.core.wmem_default,
net.core.wmem_max
- Enable logging to check if this had any effect:
`log4j.logger.org.apache.kafka.common.network.Selector=DEBUG`
- Additional tips in our docs

Competent users



- Monitor consumer lag
- Add processes when things are slow
- Automate deployment

Kafka Pros

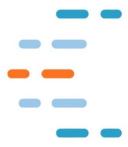


- Monitor time lag
- Collect client metrics
- Knows which side to blame
- Know which configs to tune
- Tunes the network over the WAN

Resources and Next Steps



<https://github.com/confluentinc/cp-demo>

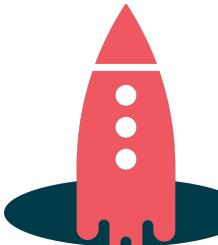


<https://www.confluent.io/download/>

<https://www.confluent.io/blog>



<https://slackpass.io/confluentcommunity>



Thank you!

@gwenshap
gwen@confluent.io

@xvrl
xavier@confluent.io

